

MAML: Model-Agnostic Meta-Learning for Fast Adaptation for Deep Networks

Yanjie Ze

July 2021

Contents

1. Meta-Learning Problem Setup (Definition & Goal)
2. Prior Work
3. Model Agnostic Meta Learning
 1. Characteristics
 2. Intuition
4. Approach
 1. Supervised Learning
 2. Gradient of Gradient
 3. Reinforcement Learning
5. Experiment

1. Problem Set-up

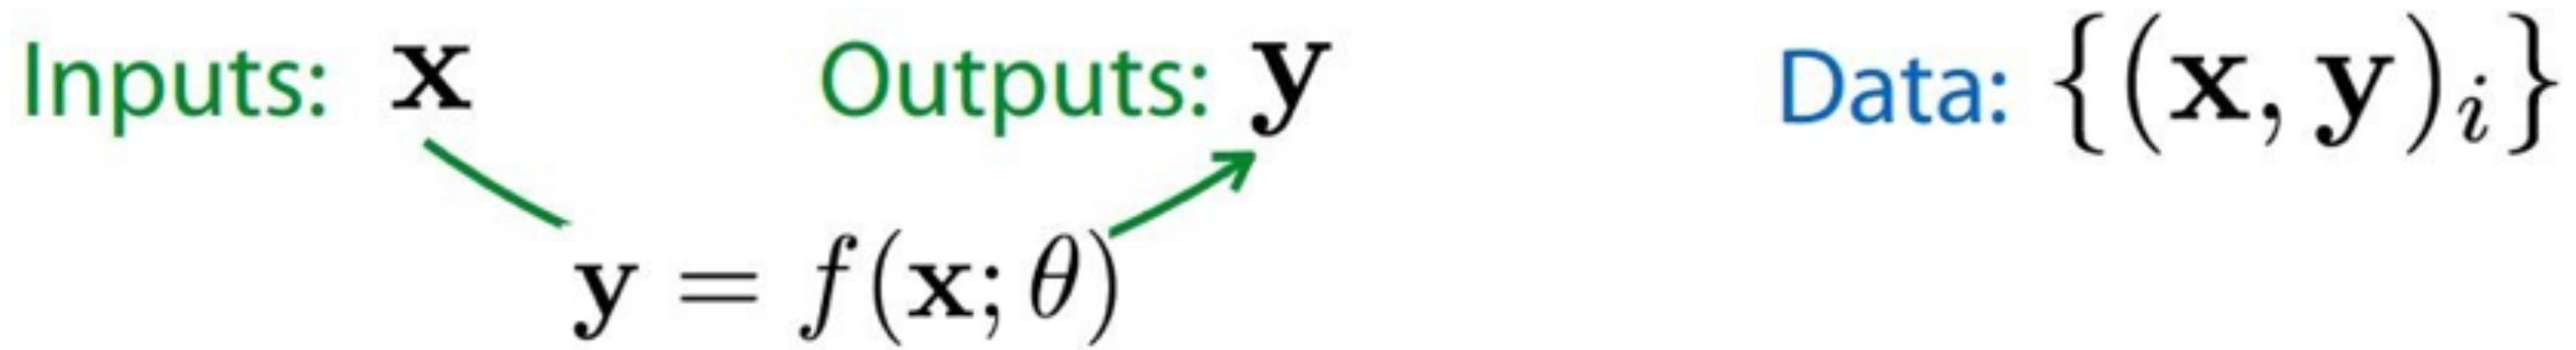
Meta-Learning Problem Set-up

- Definition of meta-learning:
 - Learner is trained by meta-learner to be able to learn on many different tasks.
- Goal:
 - Learner quickly learn new tasks from a small amount of new data.

1. Problem Set-up

Meta Supervised Learning

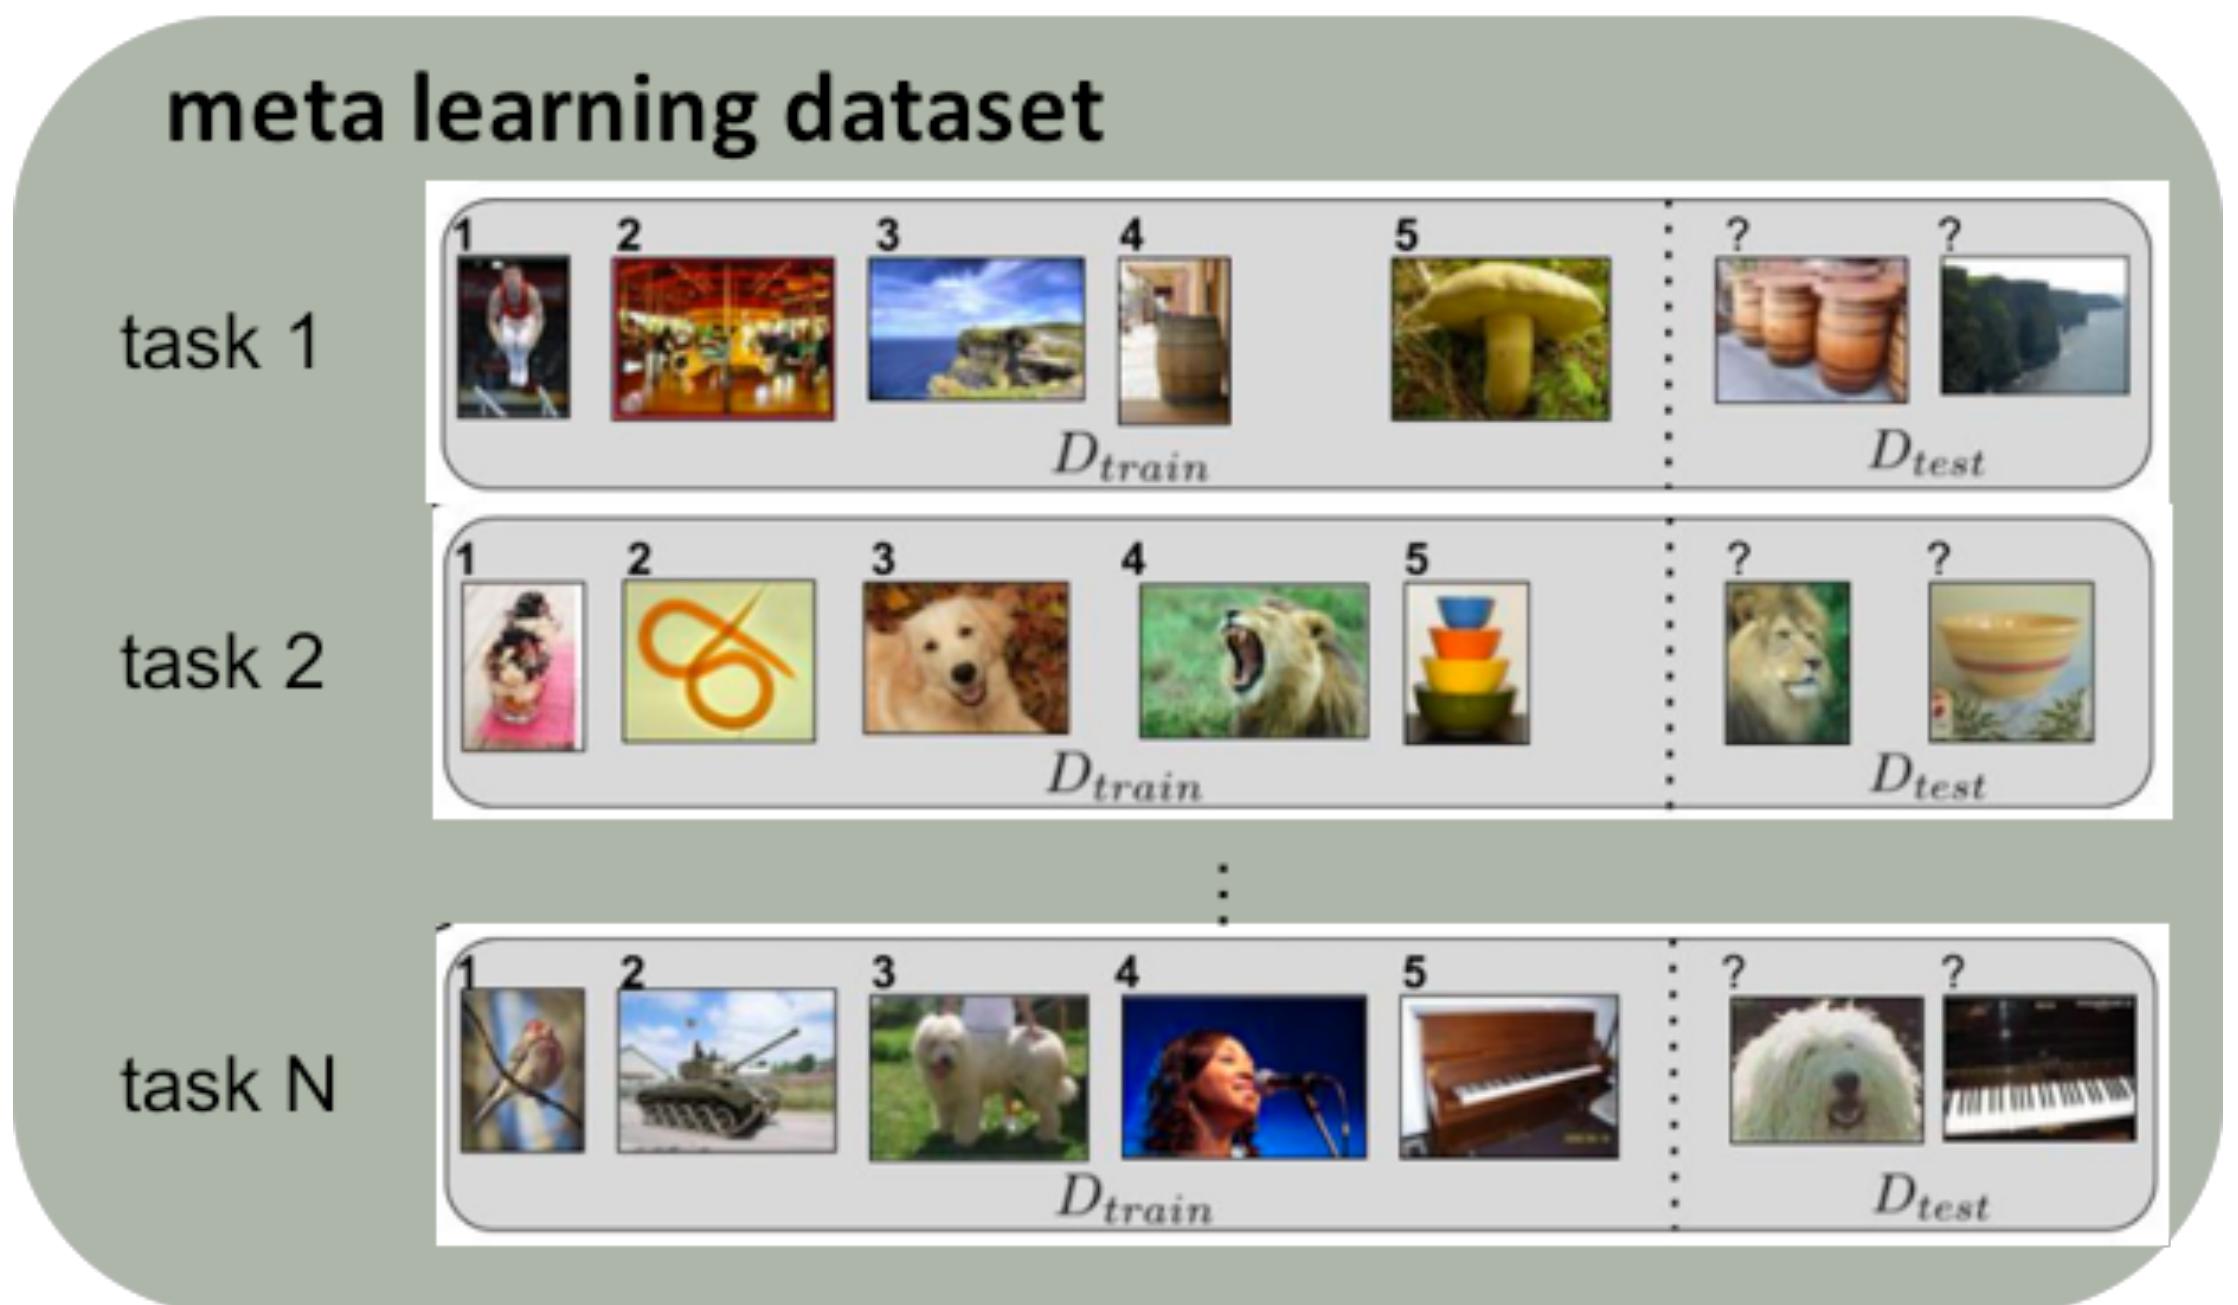
- Supervised learning:



- Meta supervised learning?

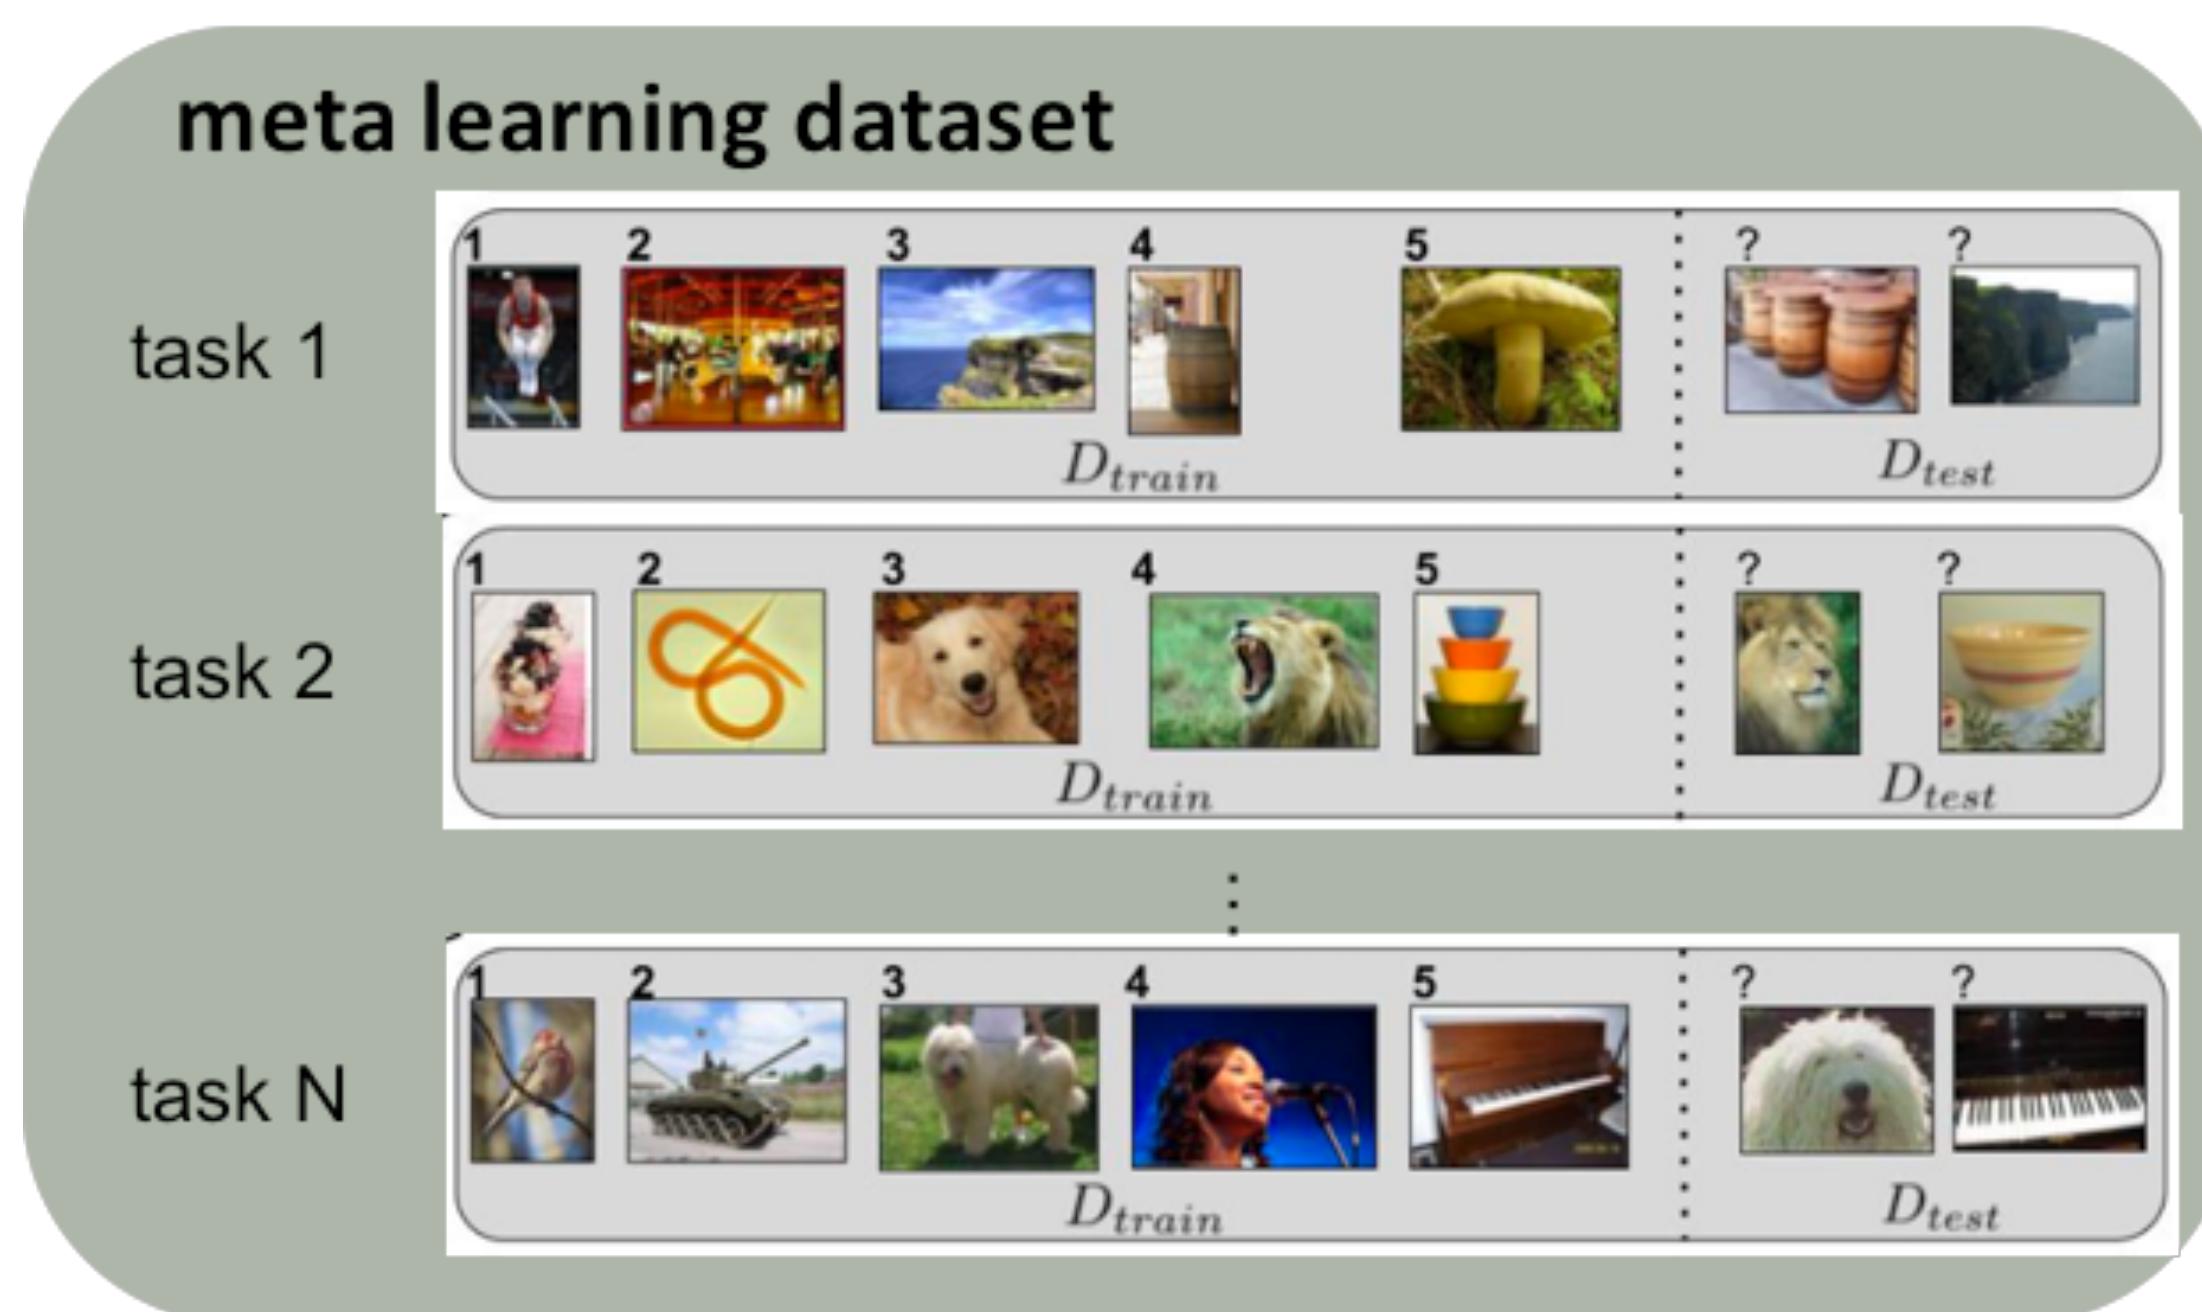
1. Problem Set-up

Meta Supervised Learning

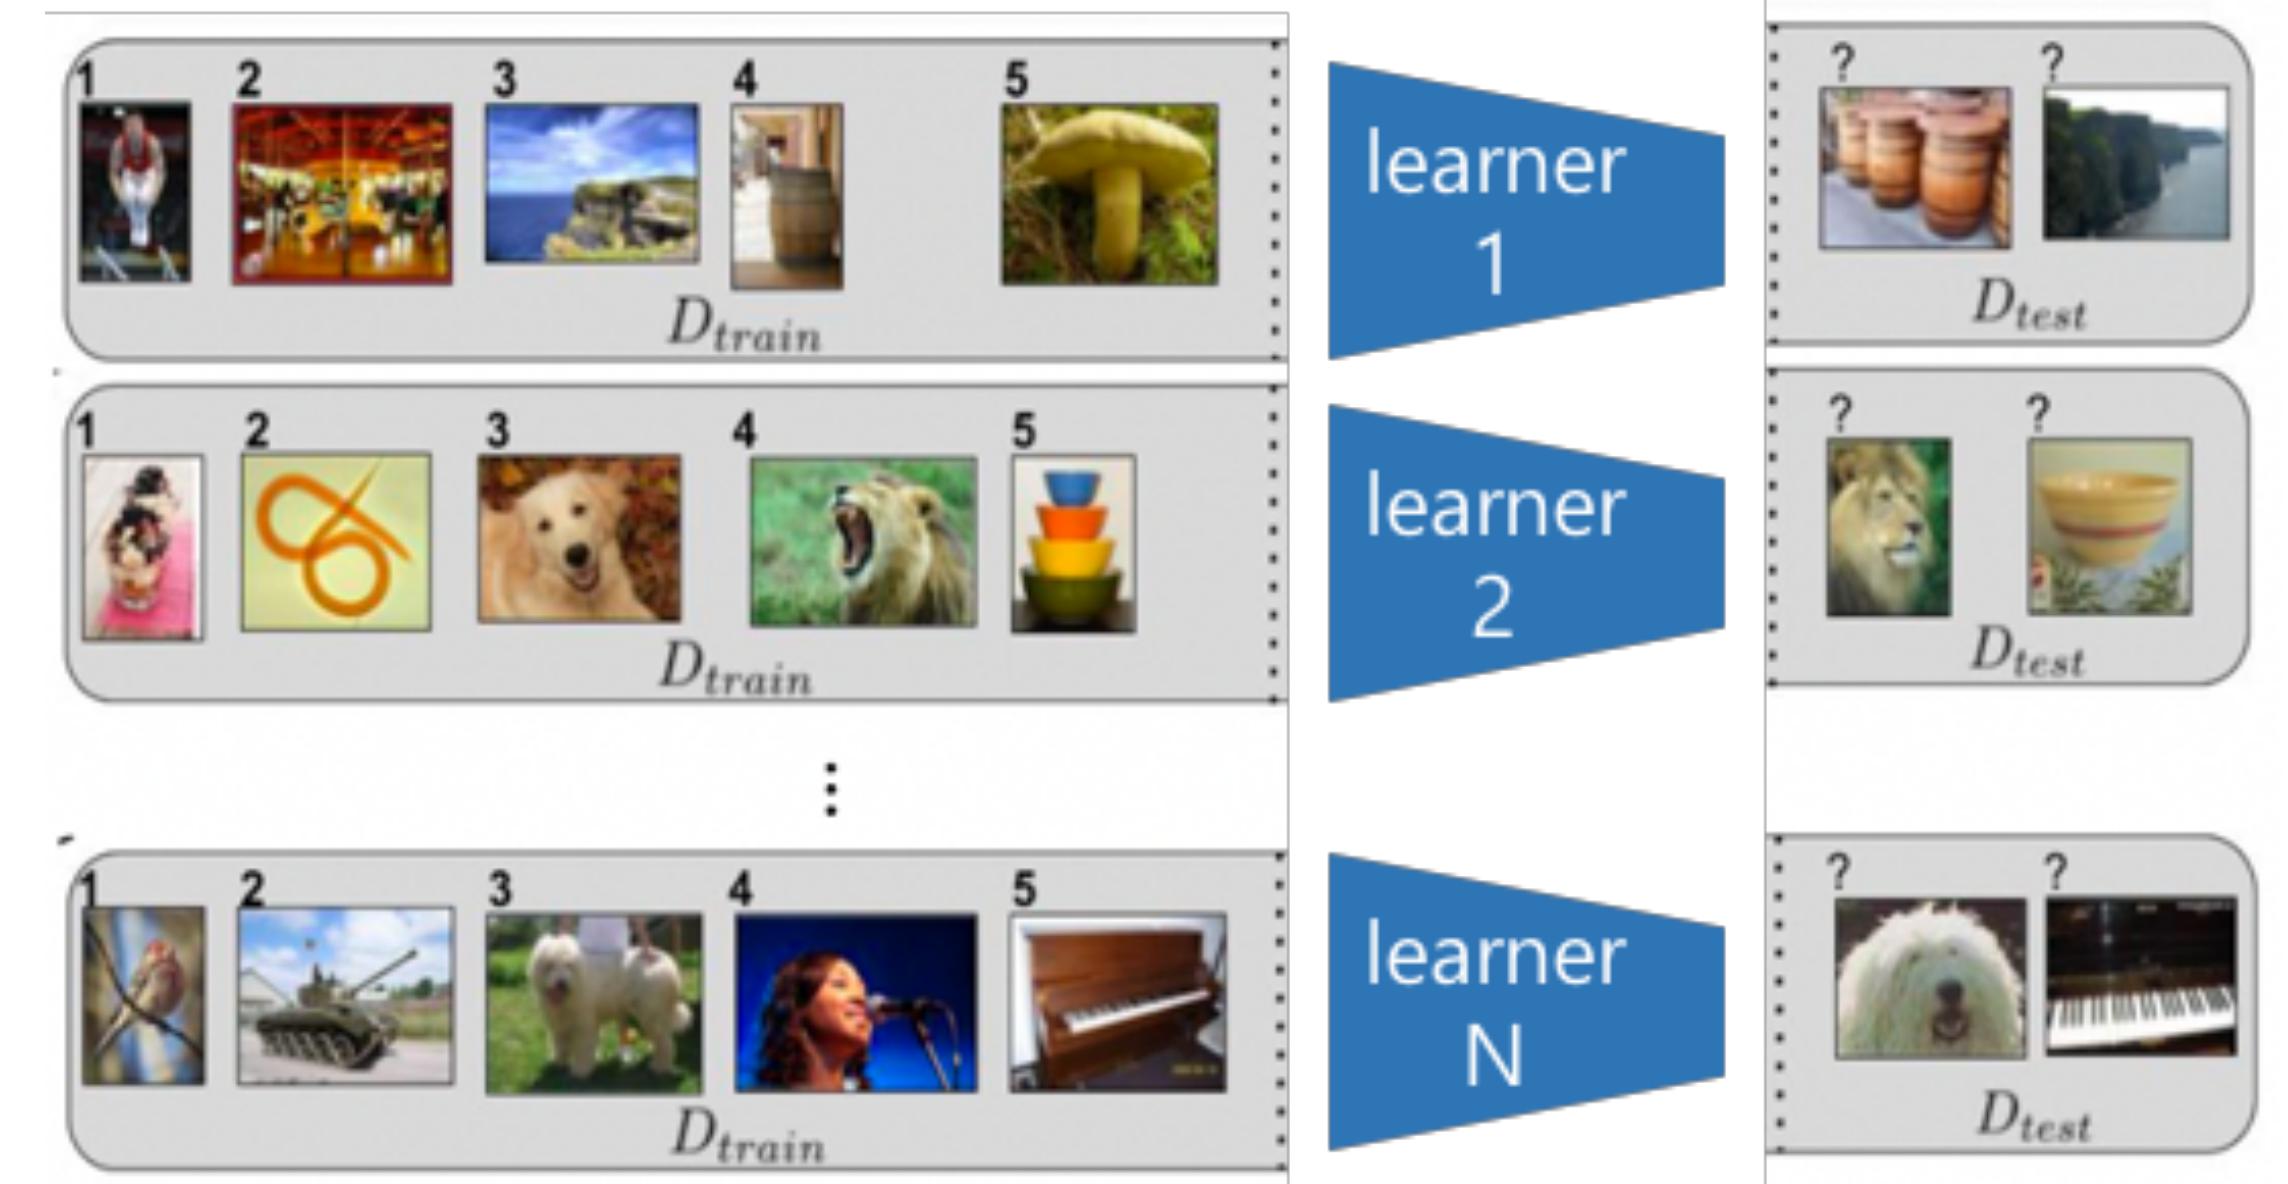


1. Problem Set-up

Meta Supervised Learning

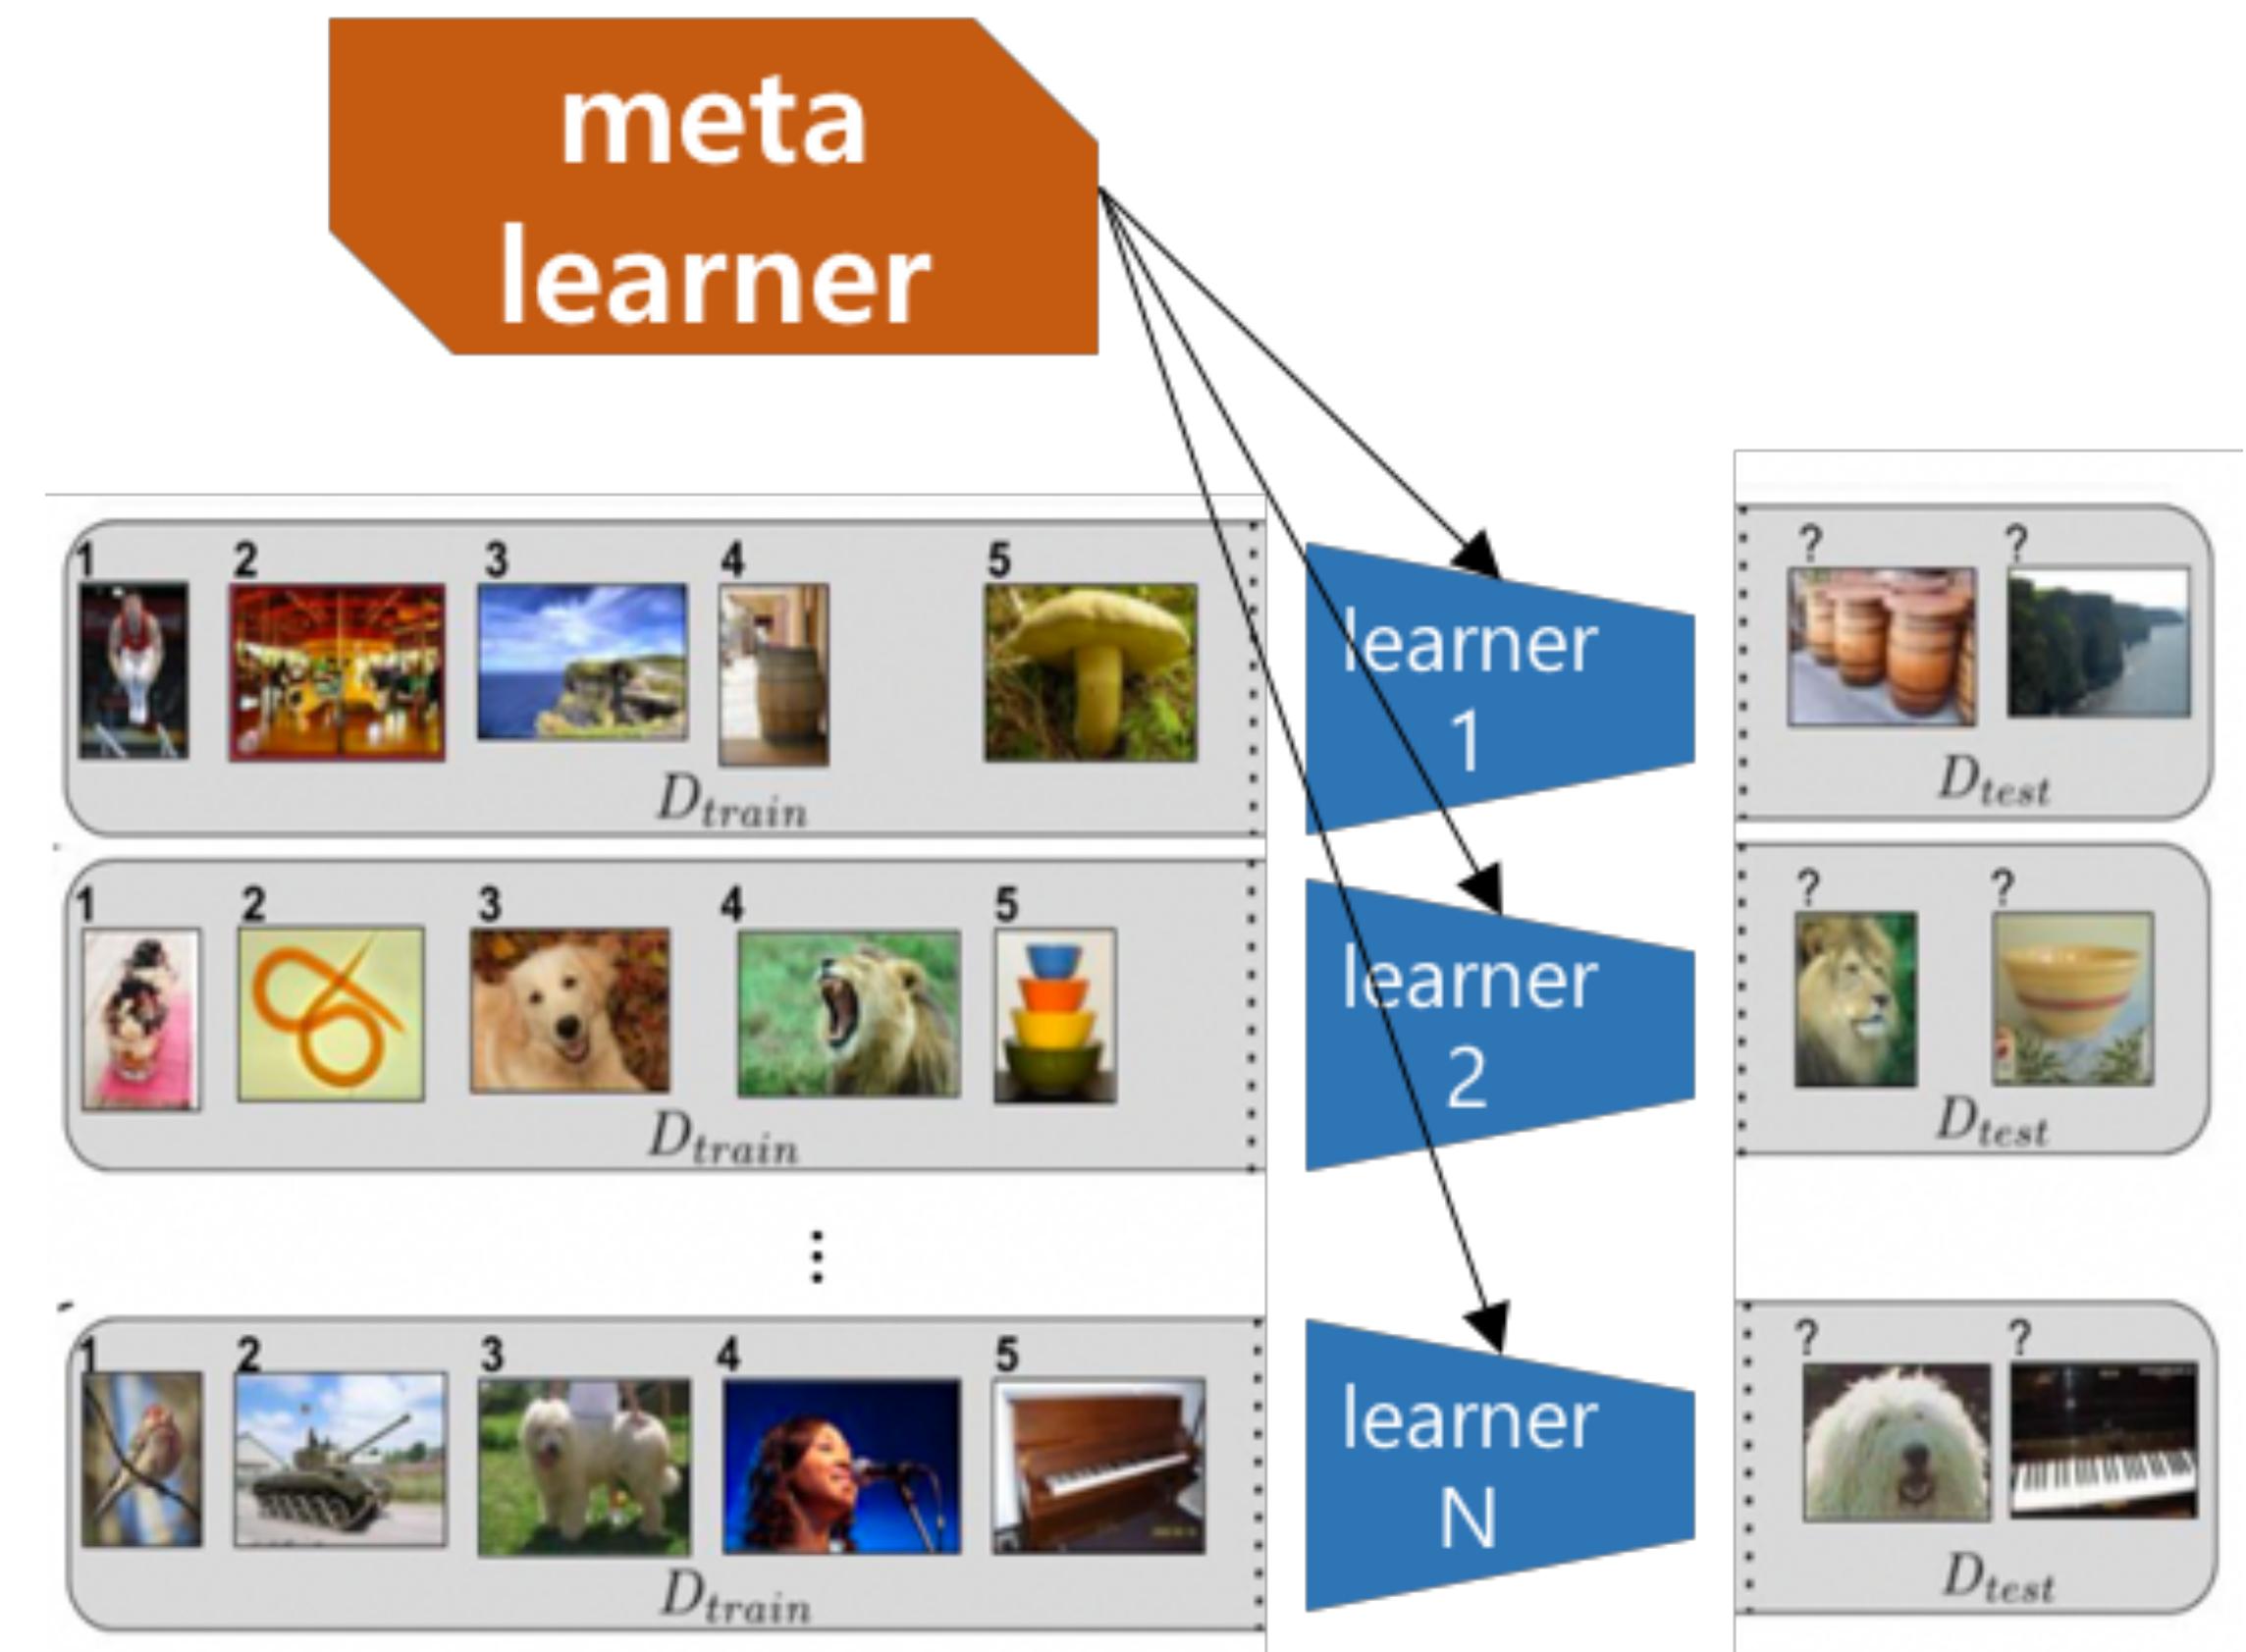
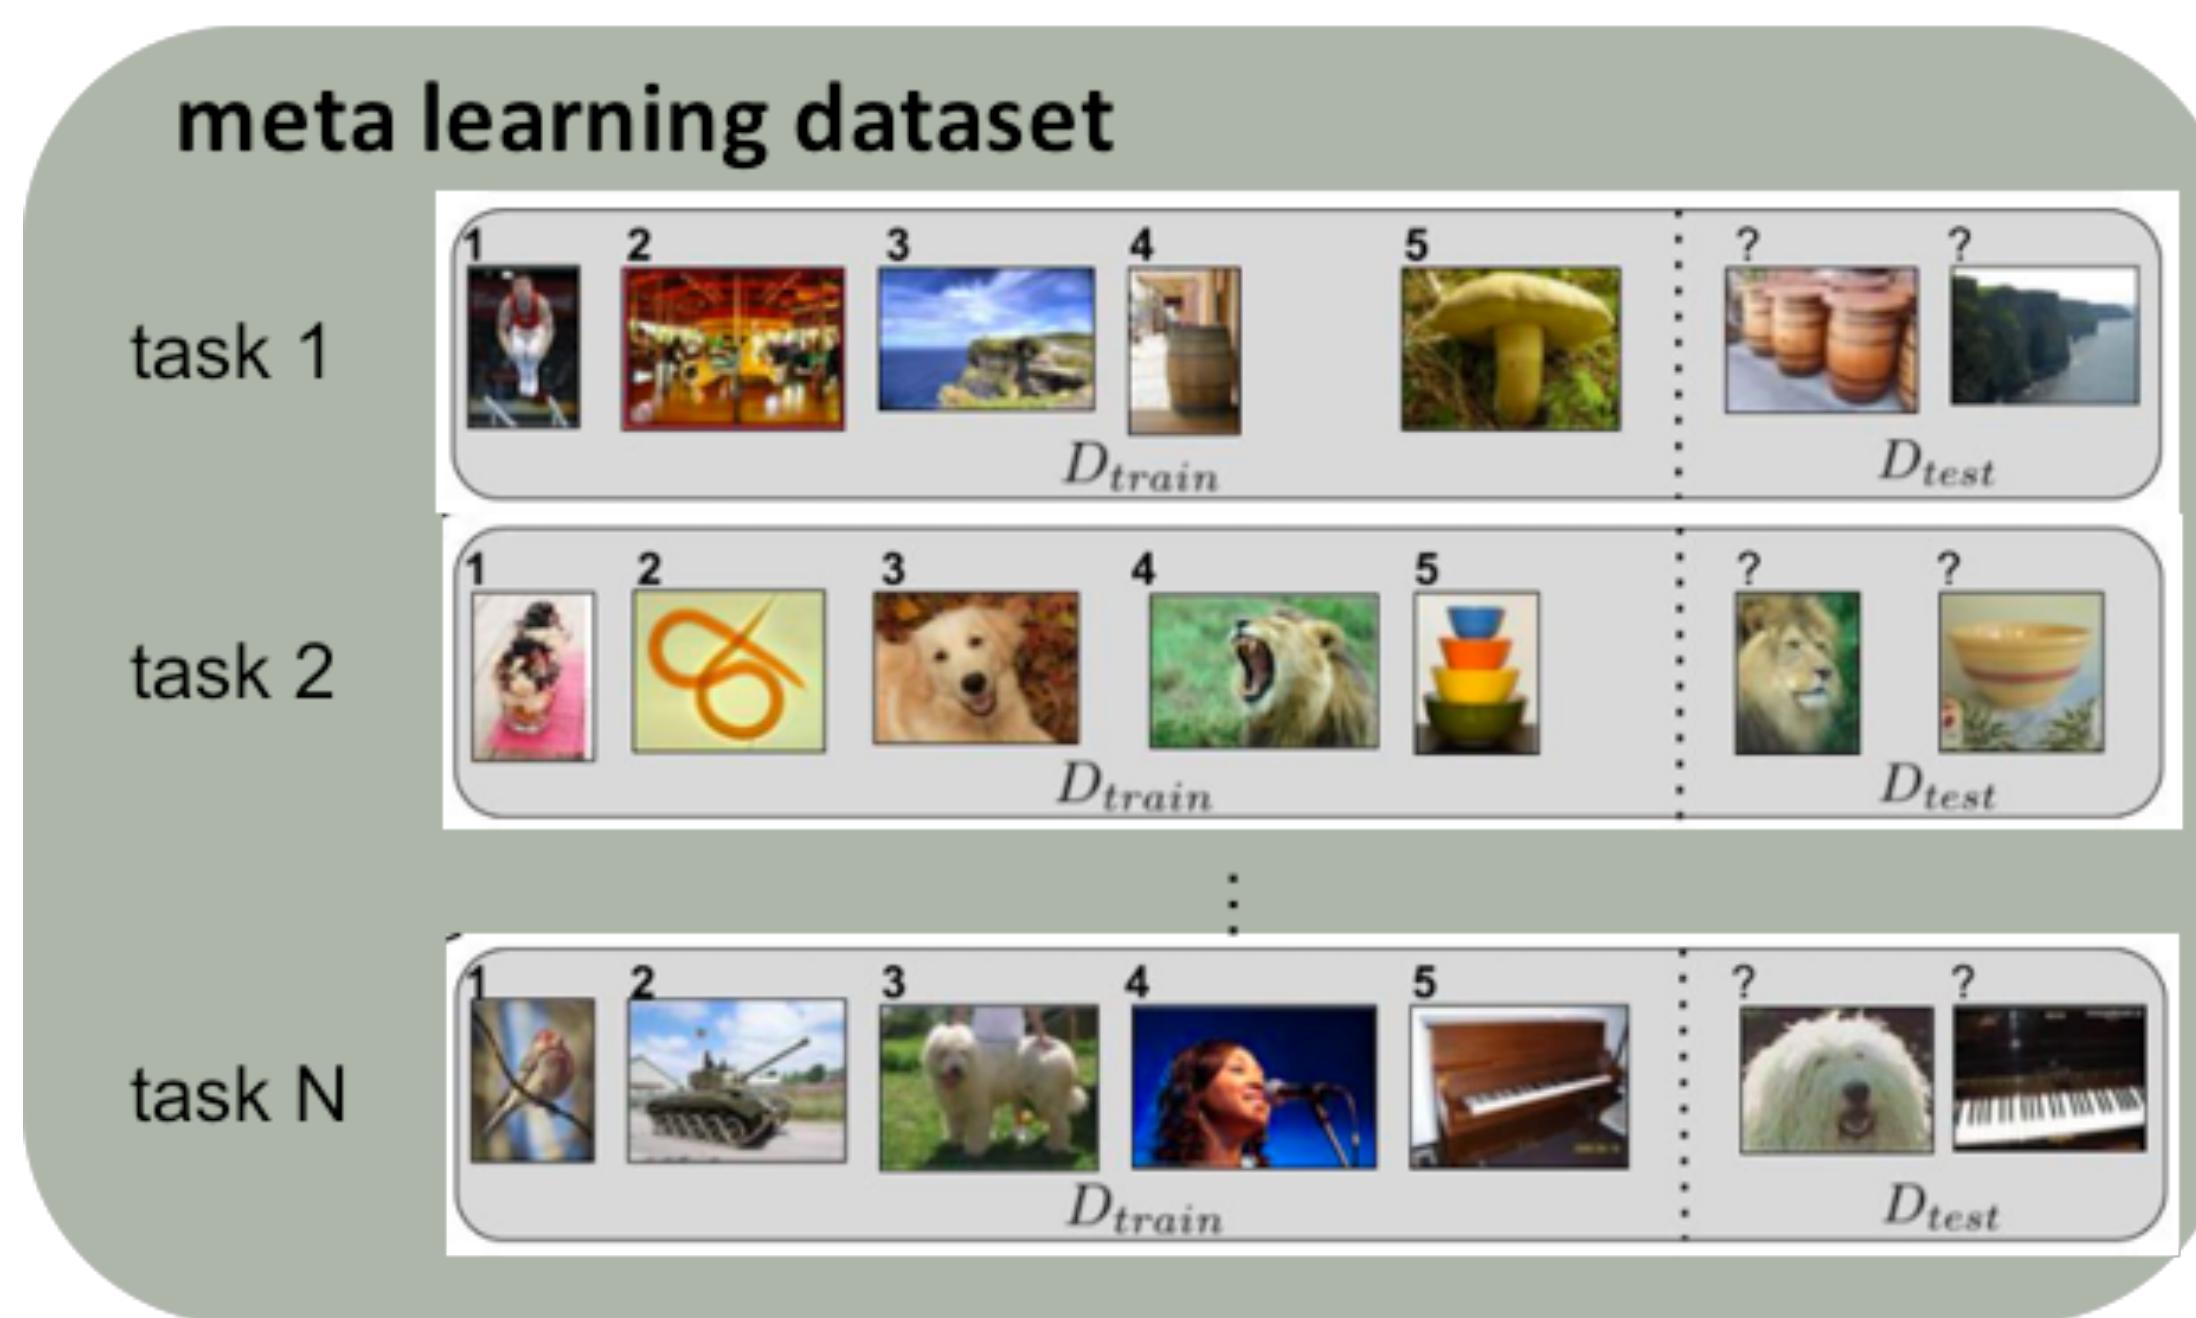


meta
learner



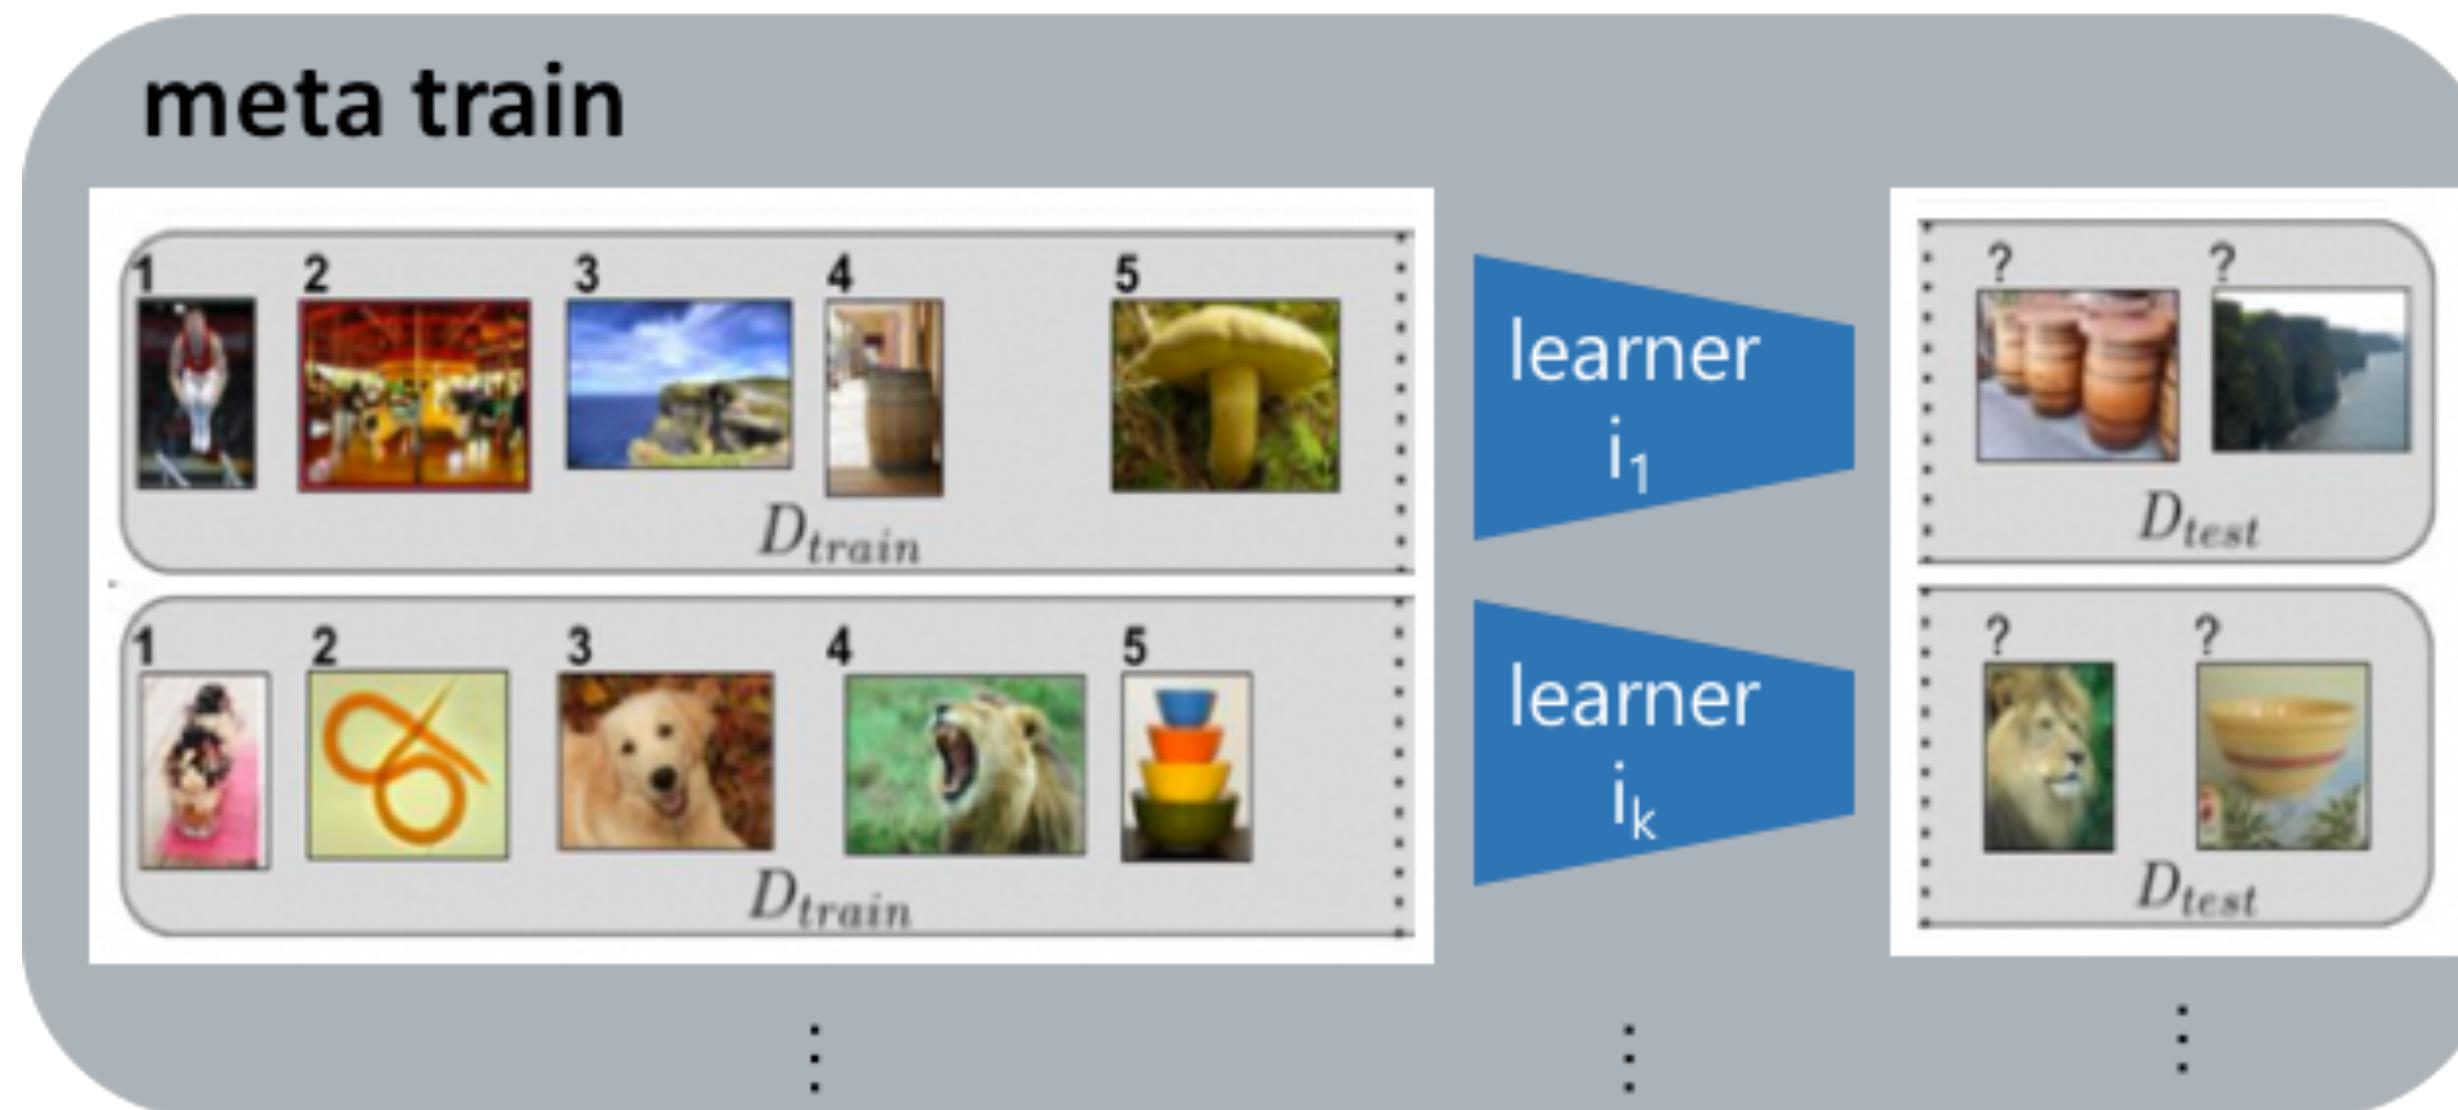
1. Problem Set-up

Meta Supervised Learning



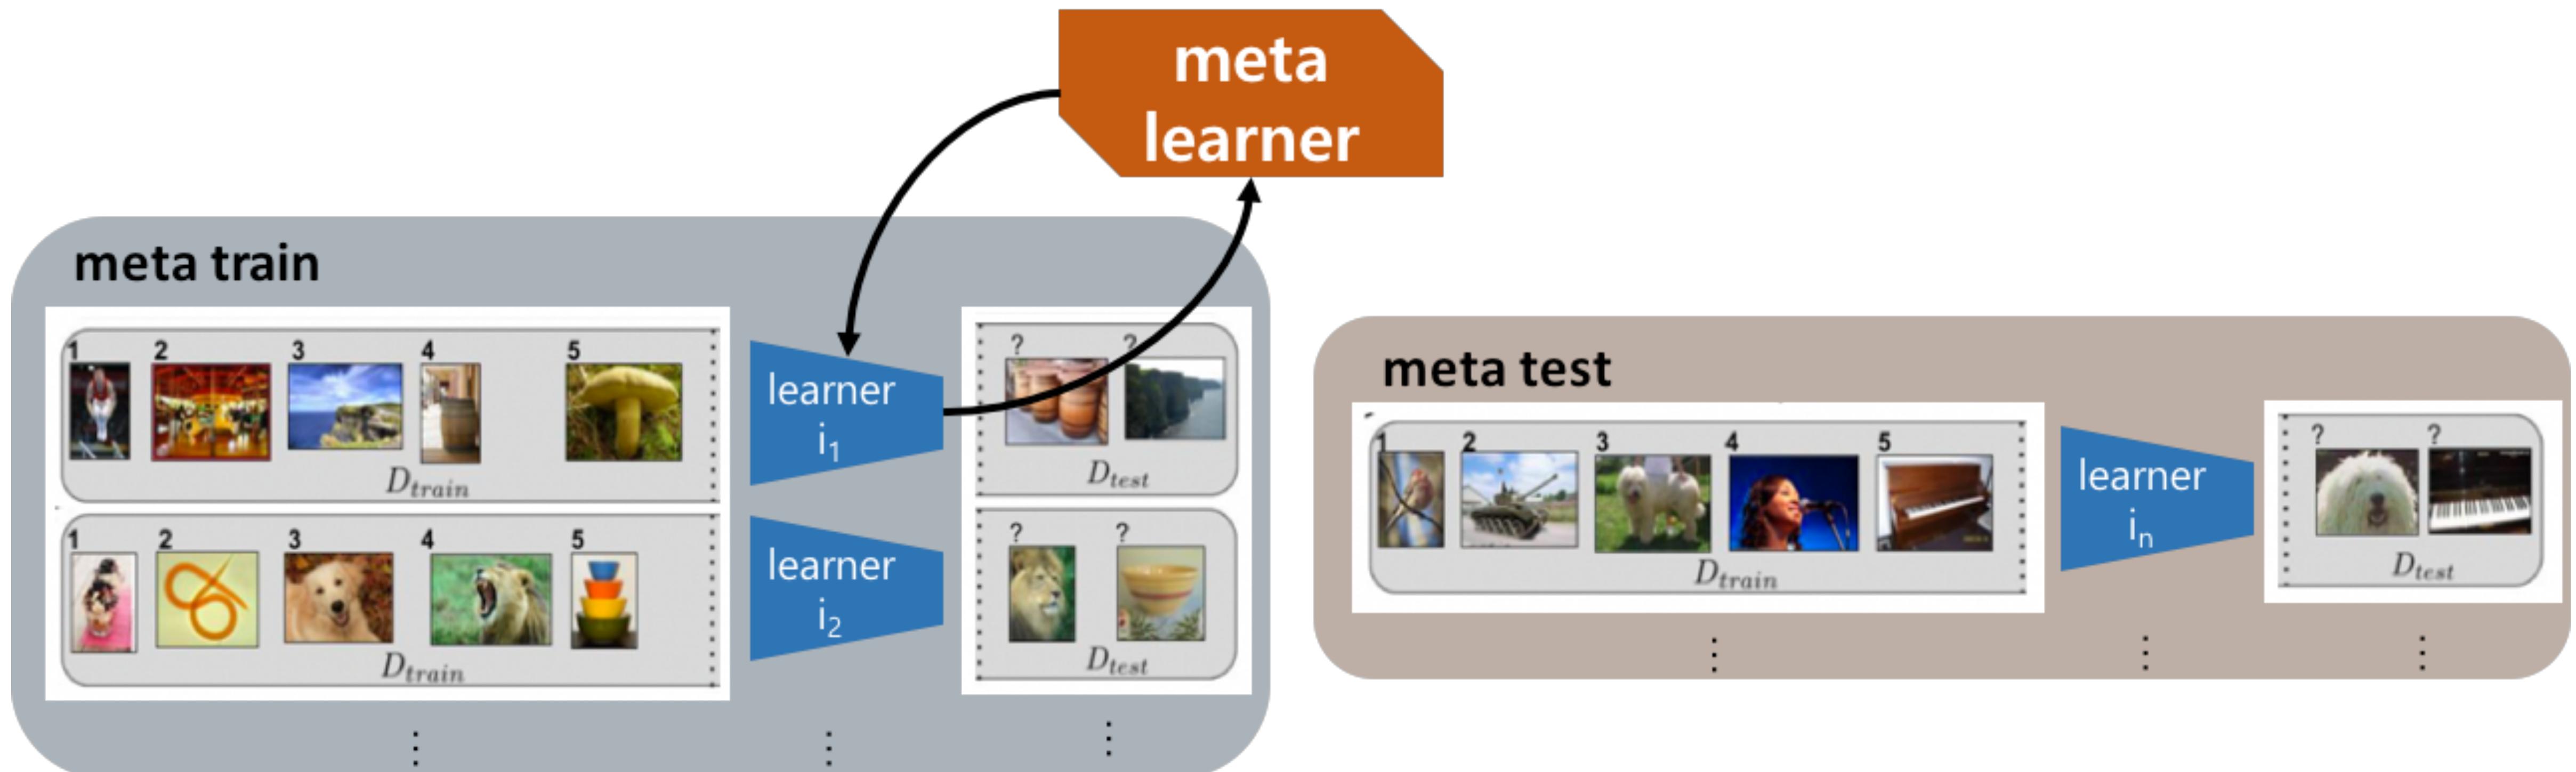
1. Problem Set-up

Meta Supervised Learning



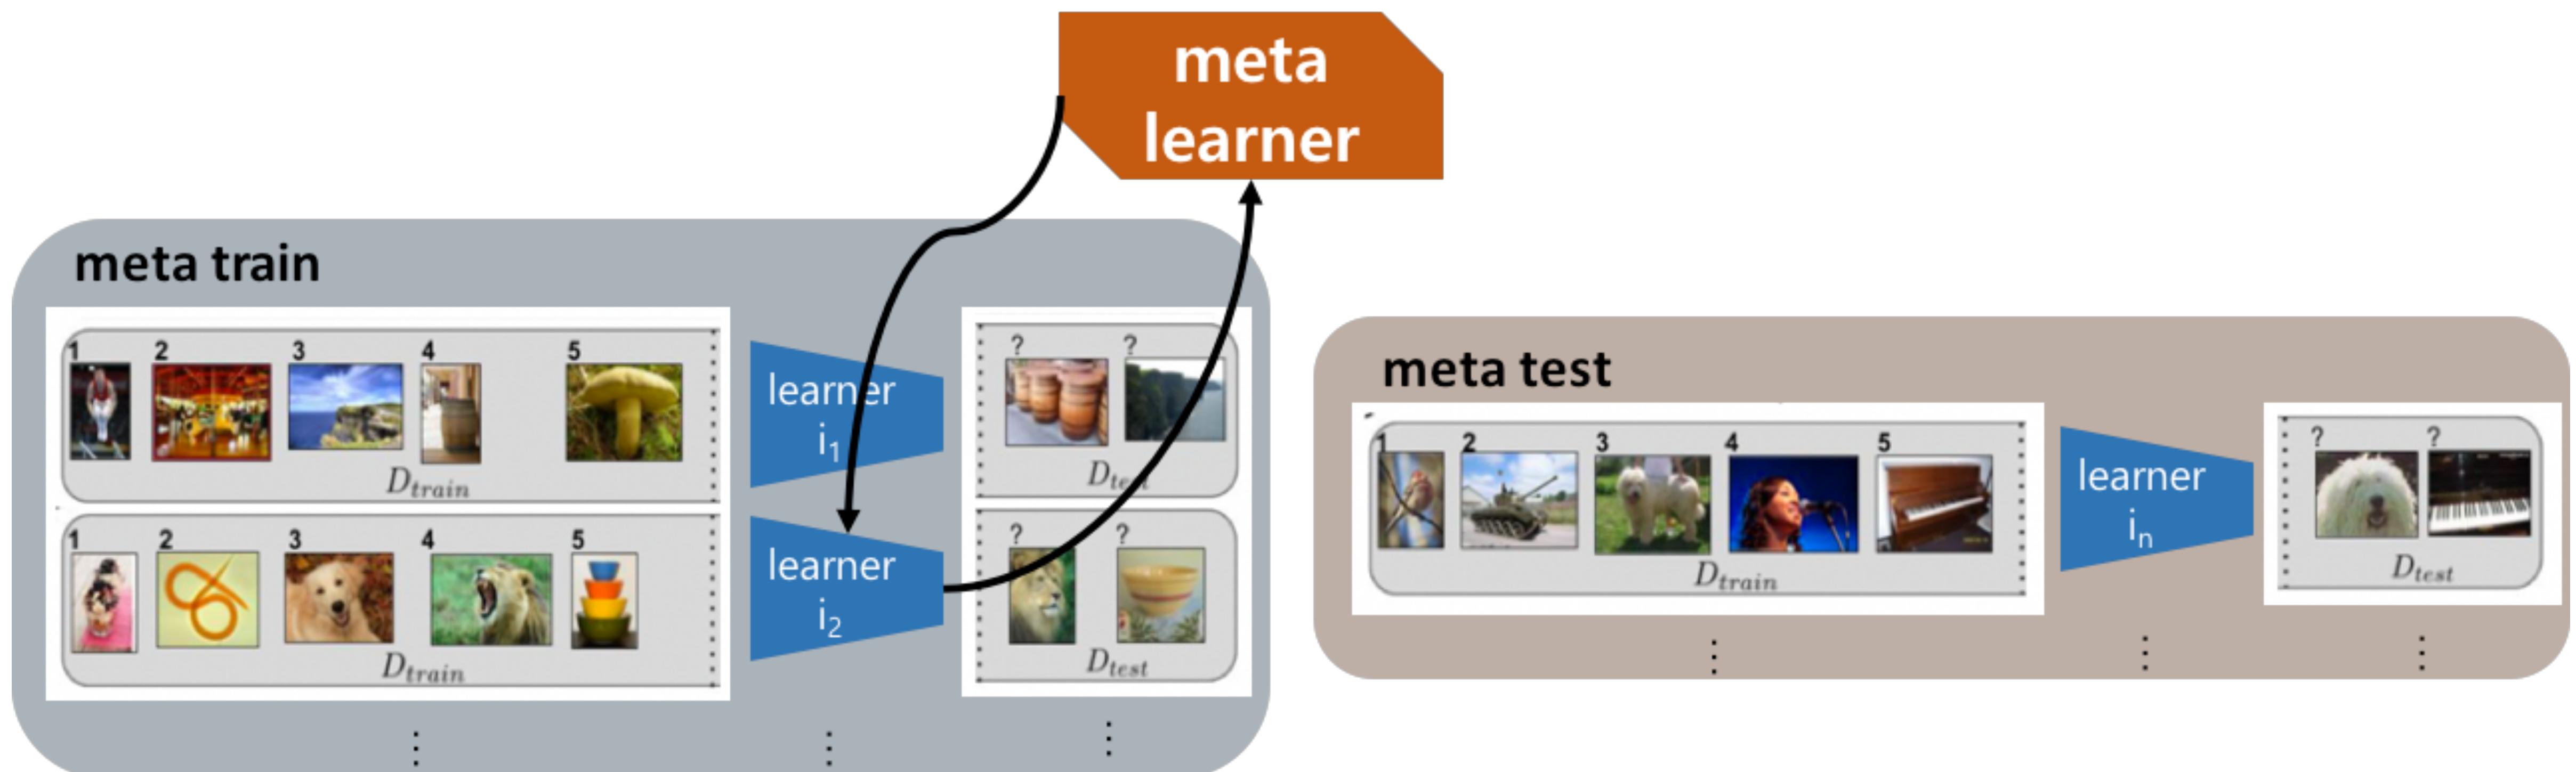
1. Problem Set-up

Meta Supervised Learning



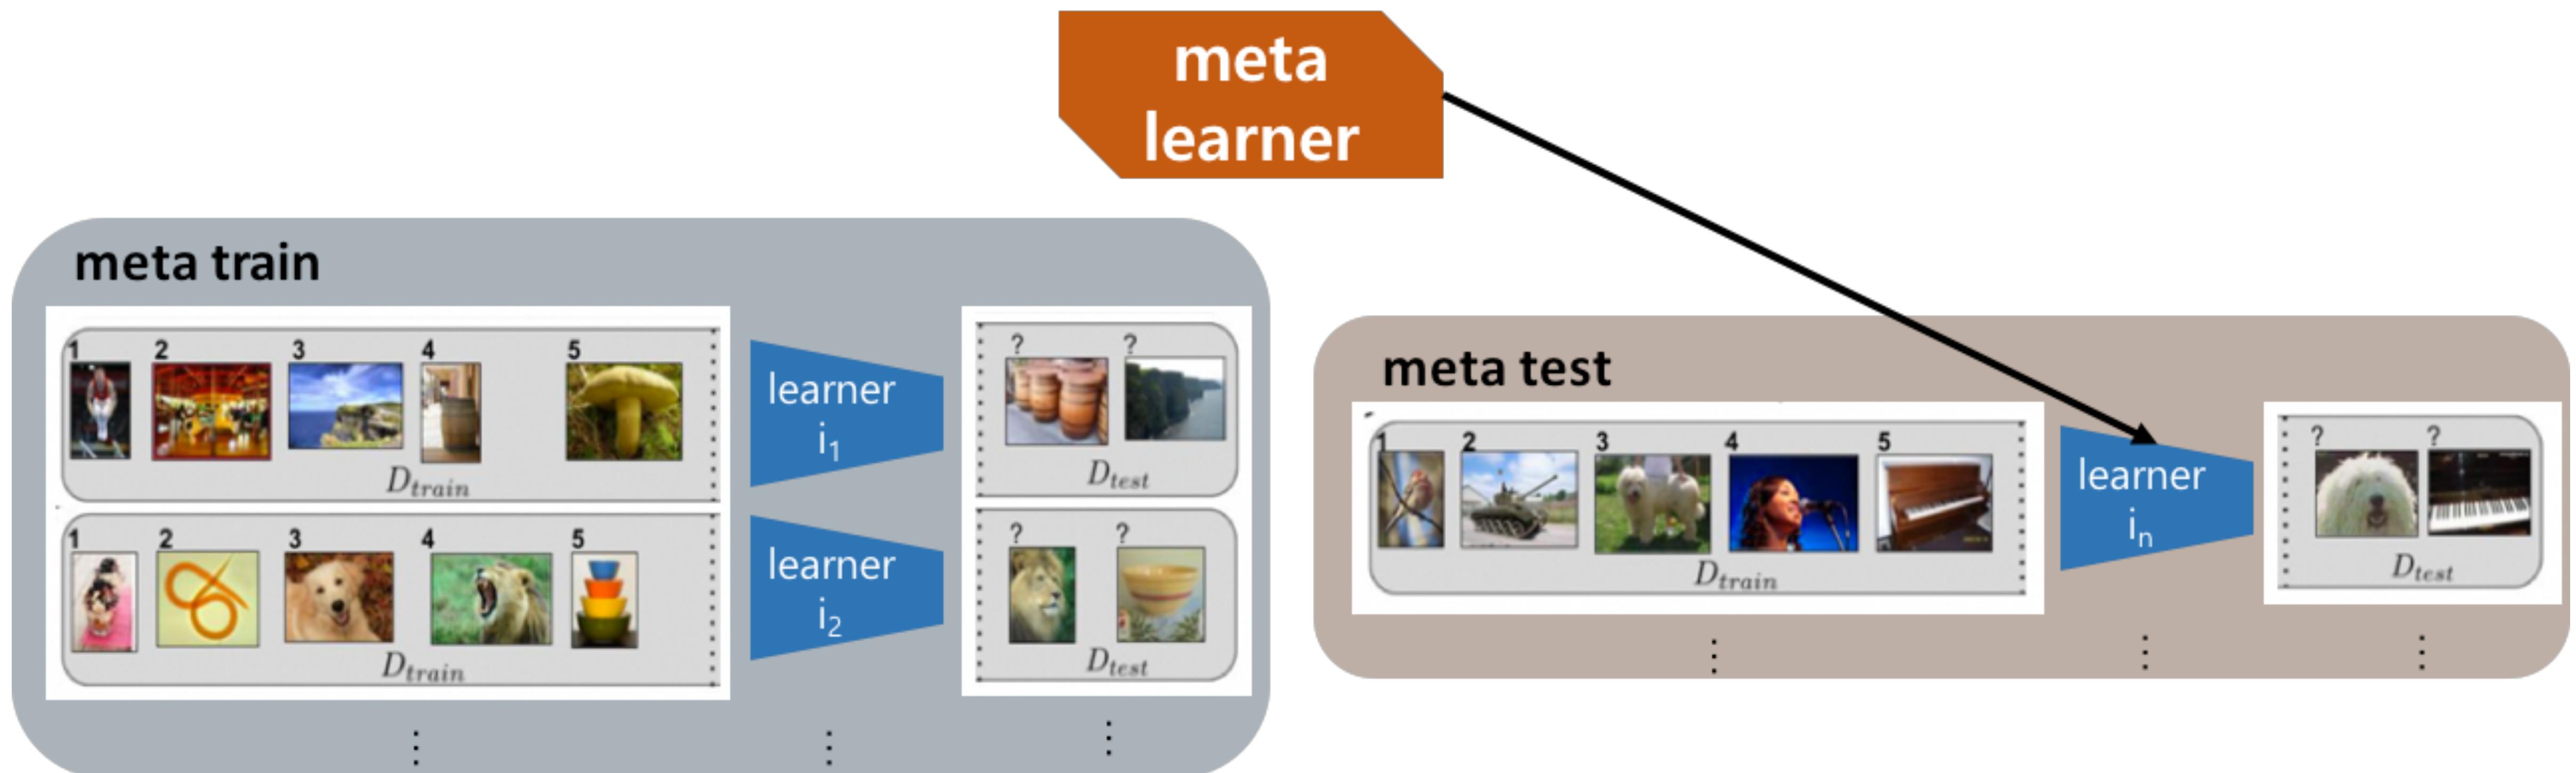
1. Problem Set-up

Meta Supervised Learning



1. Problem Set-up

Meta Supervised Learning



2. Prior Work

Prior Work

- Memory-augmented meta-learning
 - Memory-augmented models on many tasks
(Meta-learning with memory-augmented neural networks)
- Parameter initialization of deep networks
 - Exploring sensitivity while maintaining Internal representation
(Overcoming catastrophic forgetting in NN)
 - Data-dependent Initializer
(Data-Dependent Initialization of CNN)
 - learned initialization
(Gradient-based hyperparameter optimization through reversible learning)

3. Model-Agnostic Meta-Learning (MAML)

Model-Agnostic Meta-Learning (MAML)

- Goal:
 - Quickly adapt to new tasks on distribution with **only small amount of data** and with **only a few gradient steps,**
even one gradient step.
 - Learner:
 - Learn a new task by using a single gradient step.
 - Meta-learner:
 - Learn a **generalized parameter initialization** of model.
-
-

3. Model-Agnostic Meta-Learning (MAML)

Characteristics of MAML

- The MAML learner's **weights are updated using the gradient**, rather than a learned update.
 - Not need additional parameters nor require a particular learner architecture.
- **Fast adaptability** through good parameter initialization
 - Explicitly optimizes to learn Internal representation (i.e. suitable for many tasks).
 - Maximizes sensitivity of new task losses to the model parameters.

3. Model-Agnostic Meta-Learning (MAML)

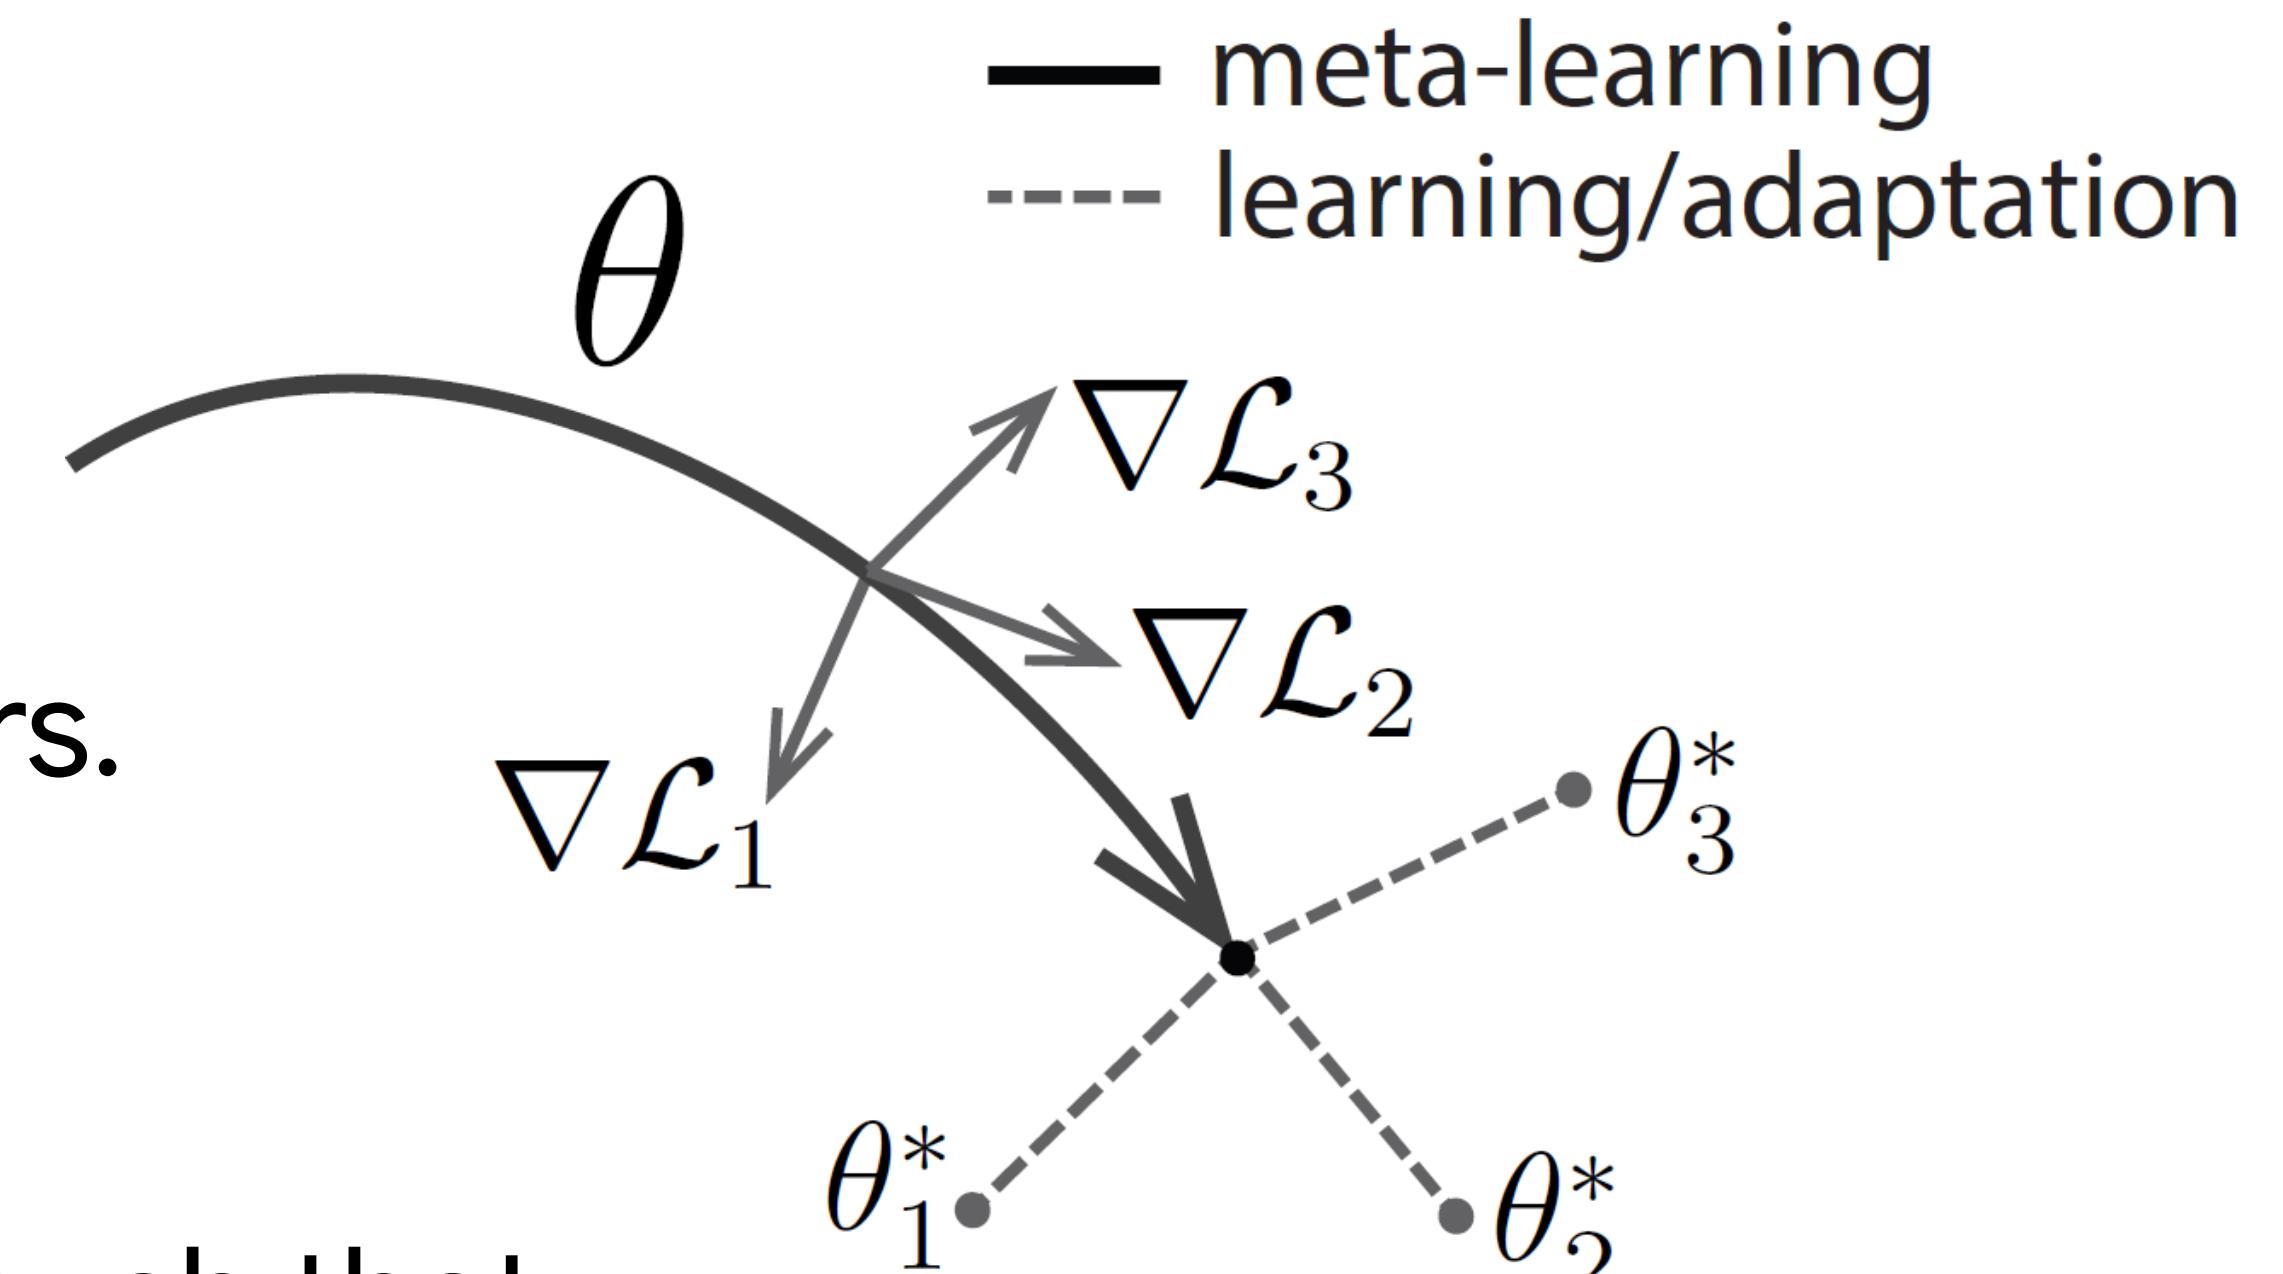
Characteristics of MAML

- **Model-agnostic** (No matter whatever model is)
 - Classification & regression with differentiable losses, Policy gradient RL.
 - The model should be parameterized.
 - No other assumptions on the form of the model.
- **Task-agnostic** (No matter whatever task is)
 - Adopted all knowledge-transferable tasks.
 - No other assumption is required.

3. Model-Agnostic Meta-Learning (MAML)

Intuition of MAML

- Some internal representations are **more transferrable** than others.
- Desired model parameter set is Θ such that:
Applying one (or a small # of) gradient step to Θ on a new task will produce maximally effective behavior.
→ Find Θ that commonly decreases loss of each task **after adaptation.**



4. Approach

Supervised Learning

- Notations
 - Model: f_θ (Model-agnostic)
 - Task distribution: $p(\mathcal{T})$
- For each Task: $T = \{L(x_1, y_1, \dots, x_K, y_K), (x, y) \sim q\}$
 - Data distribution: q (K samples are drawn from q)
 - Loss function: \mathcal{L}
 - MSE for Regression
 - Cross entropy for classification
 - Any other **differentiable** loss functions can be used.

4. Approach

Supervised Learning

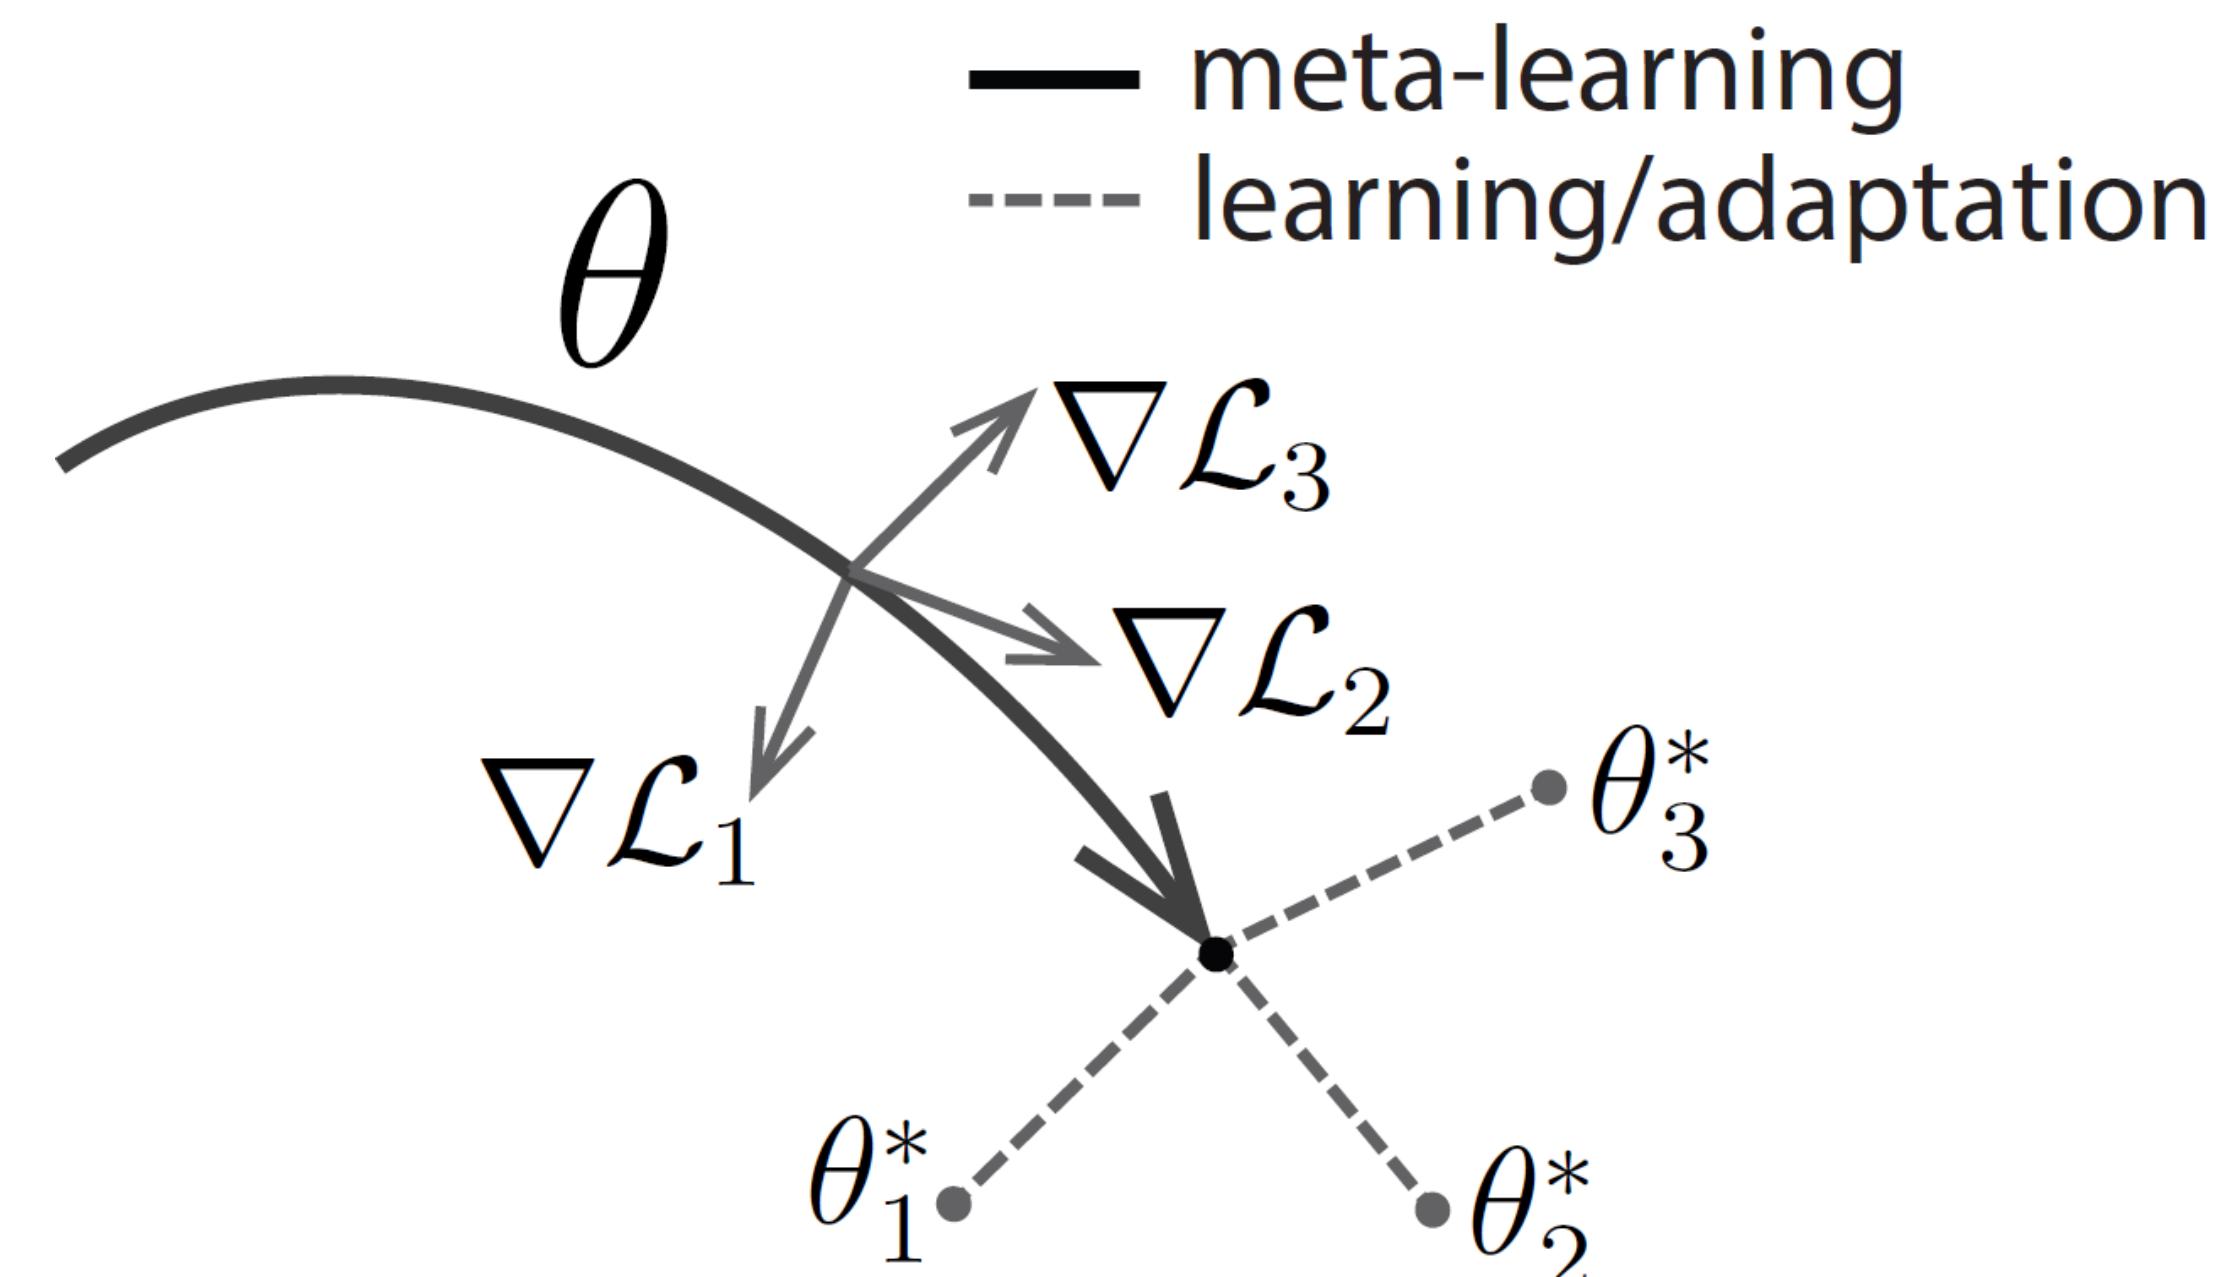
Algorithm 2 MAML for Few-Shot Supervised Learning

Require: $p(\mathcal{T})$: distribution over tasks

Require: α, β : step size hyperparameters

- 1: randomly initialize θ
 - 2: **while** not done **do**
 - 3: Sample batch of tasks $\mathcal{T}_i \sim p(\mathcal{T})$
 - 4: **for all** \mathcal{T}_i **do**
 - 5: Sample K datapoints $\mathcal{D} = \{\mathbf{x}^{(j)}, \mathbf{y}^{(j)}\}$ from \mathcal{T}_i
 - 6: Evaluate $\nabla_{\theta} \mathcal{L}_{\mathcal{T}_i}(f_{\theta})$ using \mathcal{D} and $\mathcal{L}_{\mathcal{T}_i}$
 - 7: Compute adapted parameters with gradient descent:

$$\theta'_i = \theta - \alpha \nabla_{\theta} \mathcal{L}_{\mathcal{T}_i}(f_{\theta})$$
 - 8: Sample datapoints $\mathcal{D}'_i = \{\mathbf{x}^{(j)}, \mathbf{y}^{(j)}\}$ from \mathcal{T}_i for the meta-update
 - 9: **end for**
 - 10: Update $\theta \leftarrow \theta - \beta \nabla_{\theta} \sum_{\mathcal{T}_i \sim p(\mathcal{T})} \mathcal{L}_{\mathcal{T}_i}(f_{\theta'_i})$ using each \mathcal{D}'_i and $\mathcal{L}_{\mathcal{T}_i}$
 - 11: **end while**
-



4. Approach

Gradient of Gradient

- From line 10 in Algorithm 2,

$$\theta \leftarrow \theta - \beta \nabla_{\theta} \sum_{\mathcal{T}_i \sim p(\mathcal{T})} \mathcal{L}_{\mathcal{T}_i}(f_{\theta'_i})$$

(Recall: $\theta'_i = \theta - \alpha \nabla_{\theta} \mathcal{L}_{\mathcal{T}_i}(f_{\theta})$)

4. Approach

Gradient of Gradient

- From line 10 in Algorithm 2,

$$\theta \leftarrow \theta - \beta \nabla_{\theta} \sum_{\mathcal{T}_i \sim p(\mathcal{T})} \mathcal{L}_{\mathcal{T}_i}(f_{\theta'_i})$$

$$= \theta - \beta \sum_{\mathcal{T}_i \sim p(\mathcal{T})} \nabla_{\theta} \mathcal{L}_{\mathcal{T}_i}(f_{\theta'_i})$$

(Recall: $\theta'_i = \theta - \alpha \nabla_{\theta} \mathcal{L}_{\mathcal{T}_i}(f_{\theta})$)

(\mathcal{L} is differentiable)

4. Approach

Gradient of Gradient

- From line 10 in Algorithm 2,

$$\theta \leftarrow \theta - \beta \nabla_{\theta} \sum_{\mathcal{T}_i \sim p(\mathcal{T})} \mathcal{L}_{\mathcal{T}_i}(f_{\theta'_i})$$

(Recall: $\theta'_i = \theta - \alpha \nabla_{\theta} \mathcal{L}_{\mathcal{T}_i}(f_{\theta})$)

$$= \theta - \beta \sum_{\mathcal{T}_i \sim p(\mathcal{T})} \nabla_{\theta} \mathcal{L}_{\mathcal{T}_i}(f_{\theta'_i})$$

(\mathcal{L} is differentiable)

$$= \theta - \beta \sum_{\mathcal{T}_i \sim p(\mathcal{T})} (\nabla_{\theta} \theta'_i) \nabla_{\theta'_i} \mathcal{L}_{\mathcal{T}_i}(f_{\theta'_i})$$

4. Approach

Gradient of Gradient

- From line 10 in Algorithm 2,

$$\begin{aligned}
 \theta &\leftarrow \theta - \beta \nabla_{\theta} \sum_{\mathcal{T}_i \sim p(\mathcal{T})} \mathcal{L}_{\mathcal{T}_i}(f_{\theta'_i}) && (\text{Recall: } \theta'_i = \theta - \alpha \nabla_{\theta} \mathcal{L}_{\mathcal{T}_i}(f_{\theta})) \\
 &= \theta - \beta \sum_{\mathcal{T}_i \sim p(\mathcal{T})} \nabla_{\theta} \mathcal{L}_{\mathcal{T}_i}(f_{\theta'_i}) && (\mathcal{L} \text{ is differentiable}) \\
 &= \theta - \beta \sum_{\mathcal{T}_i \sim p(\mathcal{T})} (\nabla_{\theta} \theta'_i) \nabla_{\theta'_i} \mathcal{L}_{\mathcal{T}_i}(f_{\theta'_i}) \\
 &= \theta - \beta \sum_{\mathcal{T}_i \sim p(\mathcal{T})} (I - \alpha \nabla_{\theta}^2 \mathcal{L}_{\mathcal{T}_i}(f_{\theta})) \nabla_{\theta'_i} \mathcal{L}_{\mathcal{T}_i}(f_{\theta'_i})
 \end{aligned}$$

4. Approach

Gradient of Gradient

- From line 10 in Algorithm 2,

$$\begin{aligned}
 \theta &\leftarrow \theta - \beta \nabla_{\theta} \sum_{\mathcal{T}_i \sim p(\mathcal{T})} \mathcal{L}_{\mathcal{T}_i}(f_{\theta'_i}) && (\text{Recall: } \theta'_i = \theta - \alpha \nabla_{\theta} \mathcal{L}_{\mathcal{T}_i}(f_{\theta})) \\
 &= \theta - \beta \sum_{\mathcal{T}_i \sim p(\mathcal{T})} \nabla_{\theta} \mathcal{L}_{\mathcal{T}_i}(f_{\theta'_i}) && (\mathcal{L} \text{ is differentiable}) \\
 &= \theta - \beta \sum_{\mathcal{T}_i \sim p(\mathcal{T})} (\nabla_{\theta} \theta'_i) \nabla_{\theta'_i} \mathcal{L}_{\mathcal{T}_i}(f_{\theta'_i}) \\
 &= \theta - \beta \sum_{\mathcal{T}_i \sim p(\mathcal{T})} \boxed{(I - \alpha \nabla_{\theta}^2 \mathcal{L}_{\mathcal{T}_i}(f_{\theta}))} \nabla_{\theta'_i} \mathcal{L}_{\mathcal{T}_i}(f_{\theta'_i})
 \end{aligned}$$

Calculation of Hessian matrix is required.

4. Approach

Gradient of Gradient

- From line 10 in Algorithm 2,

$$\begin{aligned}
 \theta &\leftarrow \theta - \beta \nabla_{\theta} \sum_{\mathcal{T}_i \sim p(\mathcal{T})} \mathcal{L}_{\mathcal{T}_i}(f_{\theta'_i}) && (\text{Recall: } \theta'_i = \theta - \alpha \nabla_{\theta} \mathcal{L}_{\mathcal{T}_i}(f_{\theta})) \\
 &= \theta - \beta \sum_{\mathcal{T}_i \sim p(\mathcal{T})} \nabla_{\theta} \mathcal{L}_{\mathcal{T}_i}(f_{\theta'_i}) && (\mathcal{L} \text{ is differentiable}) \\
 &= \theta - \beta \sum_{\mathcal{T}_i \sim p(\mathcal{T})} (\nabla_{\theta} \theta'_i) \nabla_{\theta'_i} \mathcal{L}_{\mathcal{T}_i}(f_{\theta'_i}) \\
 &= \theta - \beta \sum_{\mathcal{T}_i \sim p(\mathcal{T})} \boxed{(I - \alpha \nabla_{\theta}^2 \mathcal{L}_{\mathcal{T}_i}(f_{\theta}))} \nabla_{\theta'_i} \mathcal{L}_{\mathcal{T}_i}(f_{\theta'_i})
 \end{aligned}$$

Calculation of Hessian matrix is required.
 → MAML suggest 1st order approximation.

4. Approach

1st Order Approximation

- Update rule of MAML:

$$\theta \leftarrow \theta - \beta \nabla_{\theta} \sum_{\mathcal{T} \sim p(\mathcal{T})} \mathcal{L}_{\mathcal{T}_i}(f_{\theta - \alpha \nabla \mathcal{L}_{\mathcal{T}_i}(f_{\theta})})$$

4. Approach

1st Order Approximation

- Update rule of MAML:

$$\theta \leftarrow \theta - \beta \nabla_{\theta} \sum_{\mathcal{T} \sim p(\mathcal{T})} \mathcal{L}_{\mathcal{T}_i}(f_{\theta - \boxed{\alpha \nabla \mathcal{L}_{\mathcal{T}_i}(f_{\theta})}})$$

This part needs calculating Hessian.
Hessian makes MAML be slow.

4. Approach

1st Order Approximation

- Update rule of MAML:

$$\theta \leftarrow \theta - \beta \nabla_{\theta} \sum_{\mathcal{T} \sim p(\mathcal{T})} \mathcal{L}_{\mathcal{T}_i}(f_{\theta - \alpha \nabla \mathcal{L}_{\mathcal{T}_i}(f_{\theta})})$$

- Update rule of MAML with 1st order approximation:

$$\delta \leftarrow \nabla \mathcal{L}_{\mathcal{T}_i}(f_{\theta})$$

$$\theta \leftarrow \theta - \beta \nabla_{\theta} \sum_{\mathcal{T} \sim p(\mathcal{T})} \mathcal{L}_{\mathcal{T}_i}(f_{\theta - \alpha \delta})$$

4. Approach

1st Order Approximation

- Update rule of MAML:

$$\theta \leftarrow \theta - \beta \nabla_{\theta} \sum_{\mathcal{T} \sim p(\mathcal{T})} \mathcal{L}_{\mathcal{T}_i}(f_{\theta - \alpha \nabla \mathcal{L}_{\mathcal{T}_i}(f_{\theta})})$$

- Update rule of MAML with 1st order approximation:

$$\delta \leftarrow \nabla \mathcal{L}_{\mathcal{T}_i}(f_{\theta}) \quad (\text{Regard } \delta \text{ as constant})$$

$$\theta \leftarrow \theta - \beta \nabla_{\theta} \sum_{\mathcal{T} \sim p(\mathcal{T})} \mathcal{L}_{\mathcal{T}_i}(f_{\theta - \alpha \delta})$$

4. Approach

Recall: Gradient of Gradient

- From line 10 in Algorithm 2,

$$\theta \leftarrow \theta - \beta \nabla_{\theta} \sum_{\mathcal{T}_i \sim p(\mathcal{T})} \mathcal{L}_{\mathcal{T}_i}(f_{\theta'_i}) \quad (\text{Recall: } \theta'_i = \theta - \alpha \nabla_{\theta} \mathcal{L}_{\mathcal{T}_i}(f_{\theta}))$$

$$= \theta - \beta \sum_{\mathcal{T}_i \sim p(\mathcal{T})} \nabla_{\theta} \mathcal{L}_{\mathcal{T}_i}(f_{\theta'_i}) \quad (\mathcal{L} \text{ is differentiable})$$

$$= \theta - \beta \sum_{\mathcal{T}_i \sim p(\mathcal{T})} (\nabla_{\theta} \theta'_i) \nabla_{\theta'_i} \mathcal{L}_{\mathcal{T}_i}(f_{\theta'_i})$$

$$= \theta - \beta \sum_{\mathcal{T}_i \sim p(\mathcal{T})} \boxed{(I - \alpha \nabla_{\theta}^2 \mathcal{L}_{\mathcal{T}_i}(f_{\theta}))} \nabla_{\theta'_i} \mathcal{L}_{\mathcal{T}_i}(f_{\theta'_i})$$

In 1st order approximation,
we regard this as identity matrix I .

4. Approach

Reinforcement Learning

- Notations
 - Policy of an agent: f_θ such that $a_t \sim f_\theta(x_t)$
 - Task distribution: $p(\mathcal{T})$
- For each Task: $T = \{L(e_1, e_2, \dots, e_K), x_1 \sim q(x), x_{t+1} \sim q(x|x_t, a_t), H\}$
 - Episode: $e = (x_1, a_1, \dots, x_H, a_H)$
 - Transition distribution: q (K trajectories are drawn from q and f)
 - Episode length: H
 - Loss function: \mathcal{L} = expectation of sum of rewards

4. Approach

Reinforcement Learning

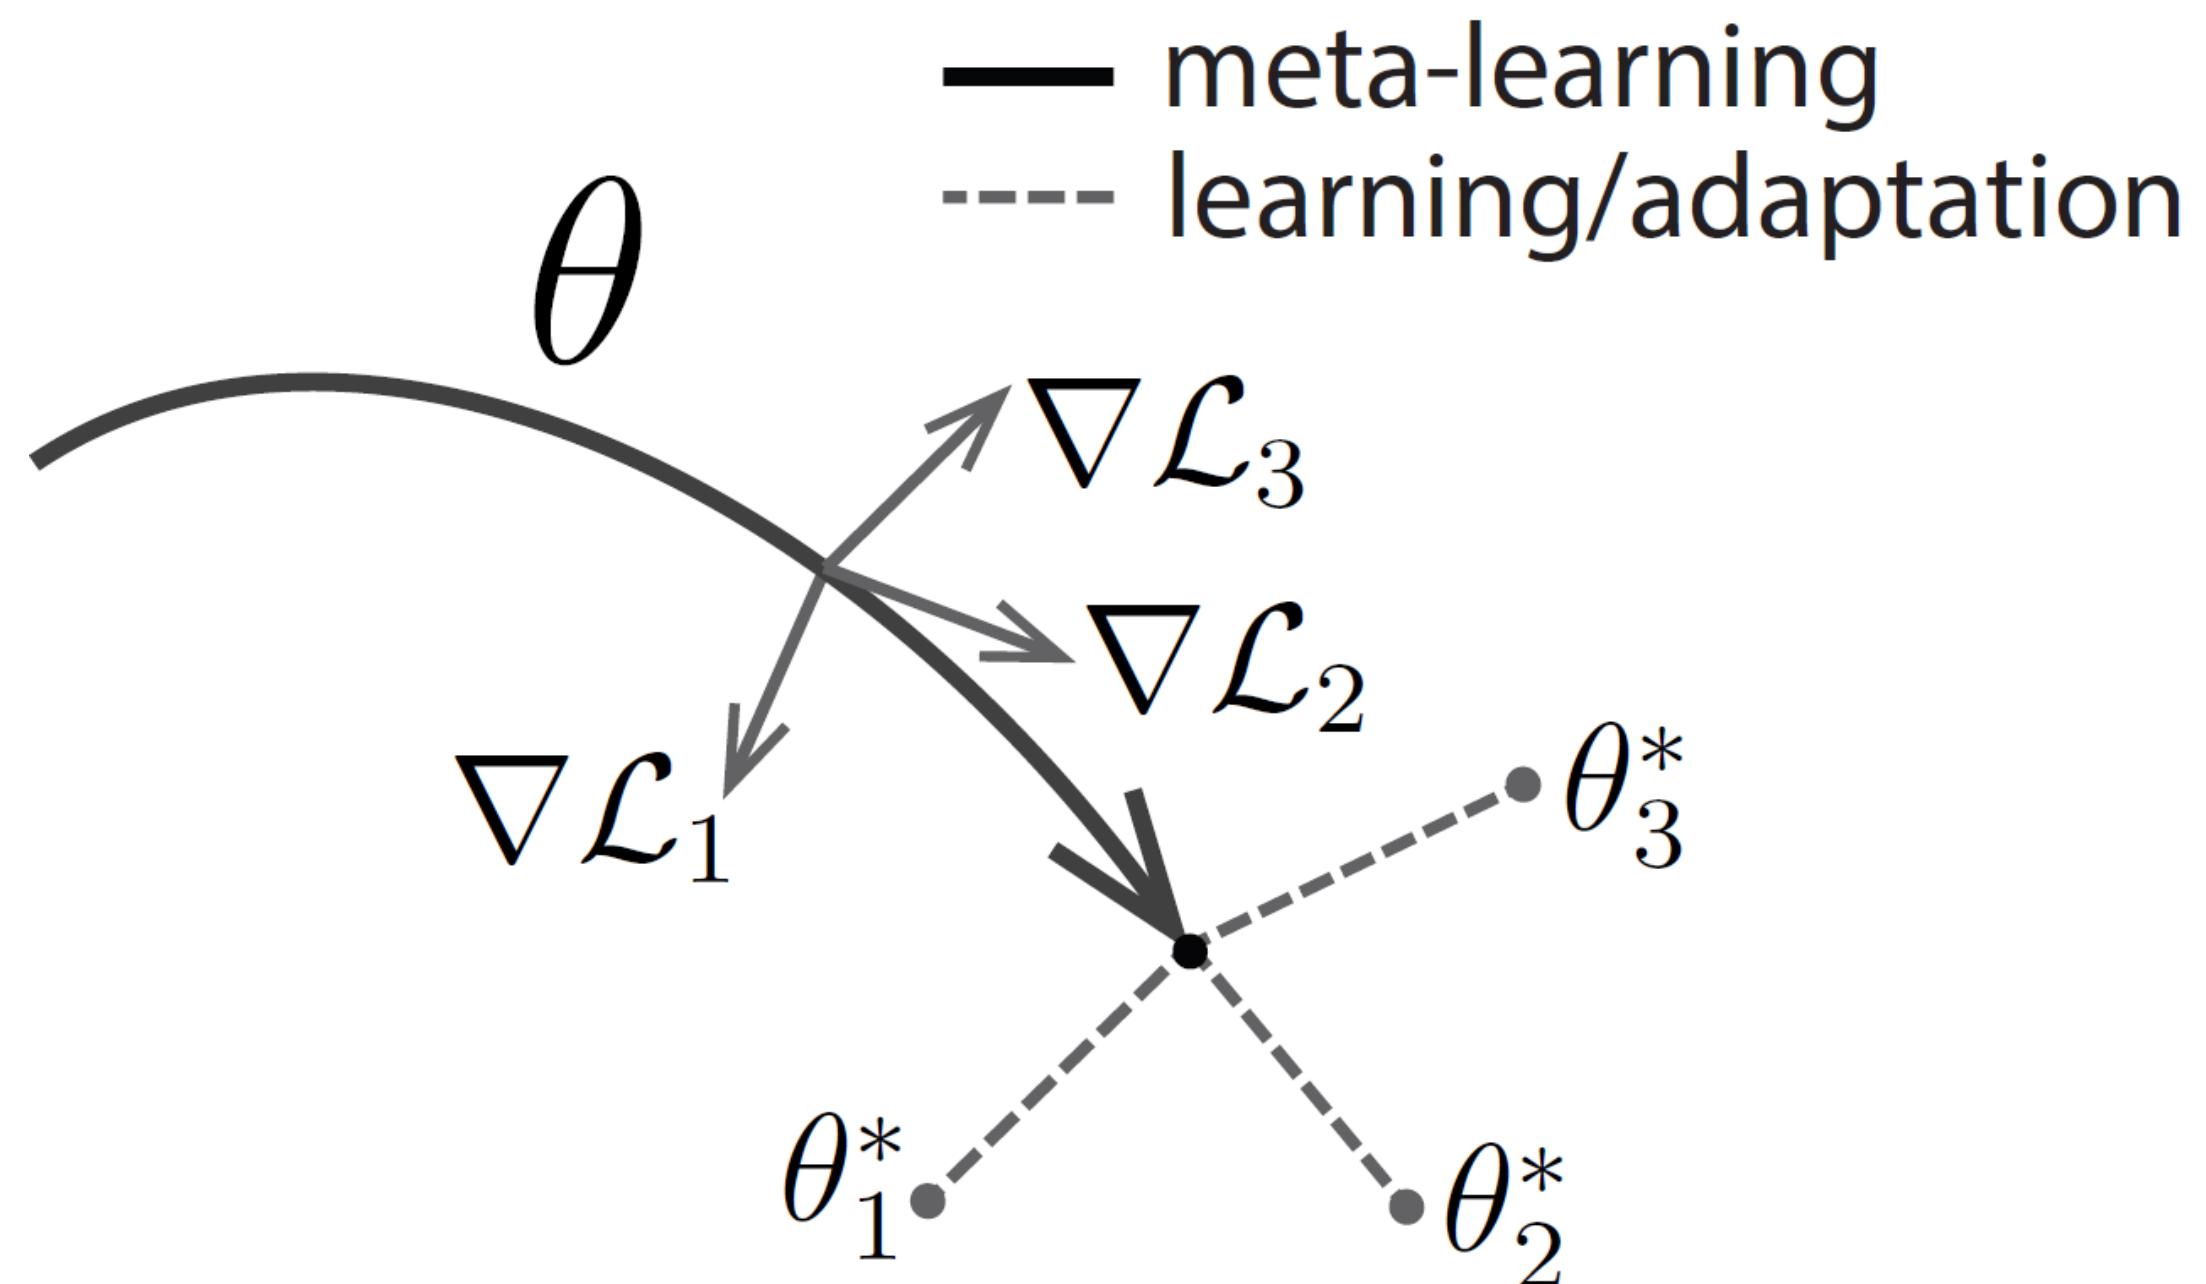
Algorithm 3 MAML for Reinforcement Learning

Require: $p(\mathcal{T})$: distribution over tasks

Require: α, β : step size hyperparameters

- 1: randomly initialize θ
 - 2: **while** not done **do**
 - 3: Sample batch of tasks $\mathcal{T}_i \sim p(\mathcal{T})$
 - 4: **for all** \mathcal{T}_i **do**
 - 5: Sample K trajectories $\mathcal{D} = \{(\mathbf{x}_1, \mathbf{a}_1, \dots, \mathbf{x}_H)\}$ using f_θ in \mathcal{T}_i
 - 6: Evaluate $\nabla_\theta \mathcal{L}_{\mathcal{T}_i}(f_\theta)$ using \mathcal{D} and $\mathcal{L}_{\mathcal{T}_i}$
 - 7: Compute adapted parameters with gradient descent:

$$\theta'_i = \theta - \alpha \nabla_\theta \mathcal{L}_{\mathcal{T}_i}(f_\theta)$$
 - 8: Sample trajectories $\mathcal{D}'_i = \{(\mathbf{x}_1, \mathbf{a}_1, \dots, \mathbf{x}_H)\}$ using $f_{\theta'_i}$ in \mathcal{T}_i
 - 9: **end for**
 - 10: Update $\theta \leftarrow \theta - \beta \nabla_\theta \sum_{\mathcal{T}_i \sim p(\mathcal{T})} \mathcal{L}_{\mathcal{T}_i}(f_{\theta'_i})$ using each \mathcal{D}'_i and $\mathcal{L}_{\mathcal{T}_i}$
 - 11: **end while**
-



5. Experiment

Experiment on Few-Shot Classification

- Omniglot (Lake et al., 2012)
 - 50 different alphabets, 1623 characters.
 - 20 instances for each characters were drawn by 20 different people.
 - 1200 for training, 423 for test.
- Mini-Imagenet (Ravi & Larochelle, 2017)
 - Classes for each set: train=64, validation=12, test=24.

5. Experiment

Results of Few-Shot Classification

- MAML outperforms state-of-arts algorithms.

	5-way Accuracy		20-way Accuracy	
	1-shot	5-shot	1-shot	5-shot
Omniglot (Lake et al., 2011)				
MANN, no conv (Santoro et al., 2016)	82.8%	94.9%	–	–
MAML, no conv (ours)	$89.7 \pm 1.1\%$	$97.5 \pm 0.6\%$	–	–
Siamese nets (Koch, 2015)	97.3%	98.4%	88.2%	97.0%
matching nets (Vinyals et al., 2016)	98.1%	98.9%	93.8%	98.5%
neural statistician (Edwards & Storkey, 2017)	98.1%	99.5%	93.2%	98.1%
memory mod. (Kaiser et al., 2017)	98.4%	99.6%	95.0%	98.6%
MAML (ours)	$98.7 \pm 0.4\%$	$99.9 \pm 0.1\%$	$95.8 \pm 0.3\%$	$98.9 \pm 0.2\%$

5. Experiment

Effect of 1st Order Approximation

- In 1st order approximation, computation is roughly 33% better.
- Performances are similar.

MiniImagenet (Ravi & Larochelle, 2017)	5-way Accuracy	
	1-shot	5-shot
fine-tuning baseline	28.86 ± 0.54%	49.79 ± 0.79%
nearest neighbor baseline	41.08 ± 0.70%	51.04 ± 0.65%
matching nets (Vinyals et al., 2016)	43.56 ± 0.84%	55.31 ± 0.73%
meta-learner LSTM (Ravi & Larochelle, 2017)	43.44 ± 0.77%	60.60 ± 0.71%
MAML, first order approx. (ours)	48.07 ± 1.75%	63.15 ± 0.91%
MAML (ours)	48.70 ± 1.84%	63.11 ± 0.92%

5. Experiment

Effect of 1st Order Approximation

- In 1st order approximation, computation is roughly 33% better.
- Performances are similar.

MiniImagenet (Ravi & Larochelle, 2017)	5-way Accuracy	
	1-shot	5-shot
fine-tuning baseline	28.86 ± 0.54%	49.79 ± 0.79%
nearest neighbor baseline	41.08 ± 0.70%	51.04 ± 0.65%
matching nets (Vinyals et al., 2016)	43.56 ± 0.84%	55.31 ± 0.73%
meta-learner LSTM (Ravi & Larochelle, 2017)	43.44 ± 0.77%	60.60 ± 0.71%
MAML, first order approx. (ours)	48.07 ± 1.75%	63.15 ± 0.91%
MAML (ours)	48.70 ± 1.84%	63.11 ± 0.92%

Performances are not that different.

5. Experiment

Effect of 1st Order Approximation

- In 1st order approximation, computation is roughly 33% better.
- Performances are similar.

	5-way Accuracy	
	1-shot	5-shot
MiniImagenet (Ravi & Larochelle, 2017)		
fine-tuning baseline	28.86 ± 0.54%	49.79 ± 0.79%
nearest neighbor baseline	41.08 ± 0.70%	51.04 ± 0.65%
matching nets (Vinyals et al., 2016)	43.56 ± 0.84%	55.31 ± 0.73%
meta-learner LSTM (Ravi & Larochelle, 2017)	43.44 ± 0.77%	60.60 ± 0.71%
MAML, first order approx. (ours)	48.07 ± 1.75%	63.15 ± 0.91%
MAML (ours)	48.70 ± 1.84%	63.11 ± 0.92%

Some performance is even better.

5. Experiment

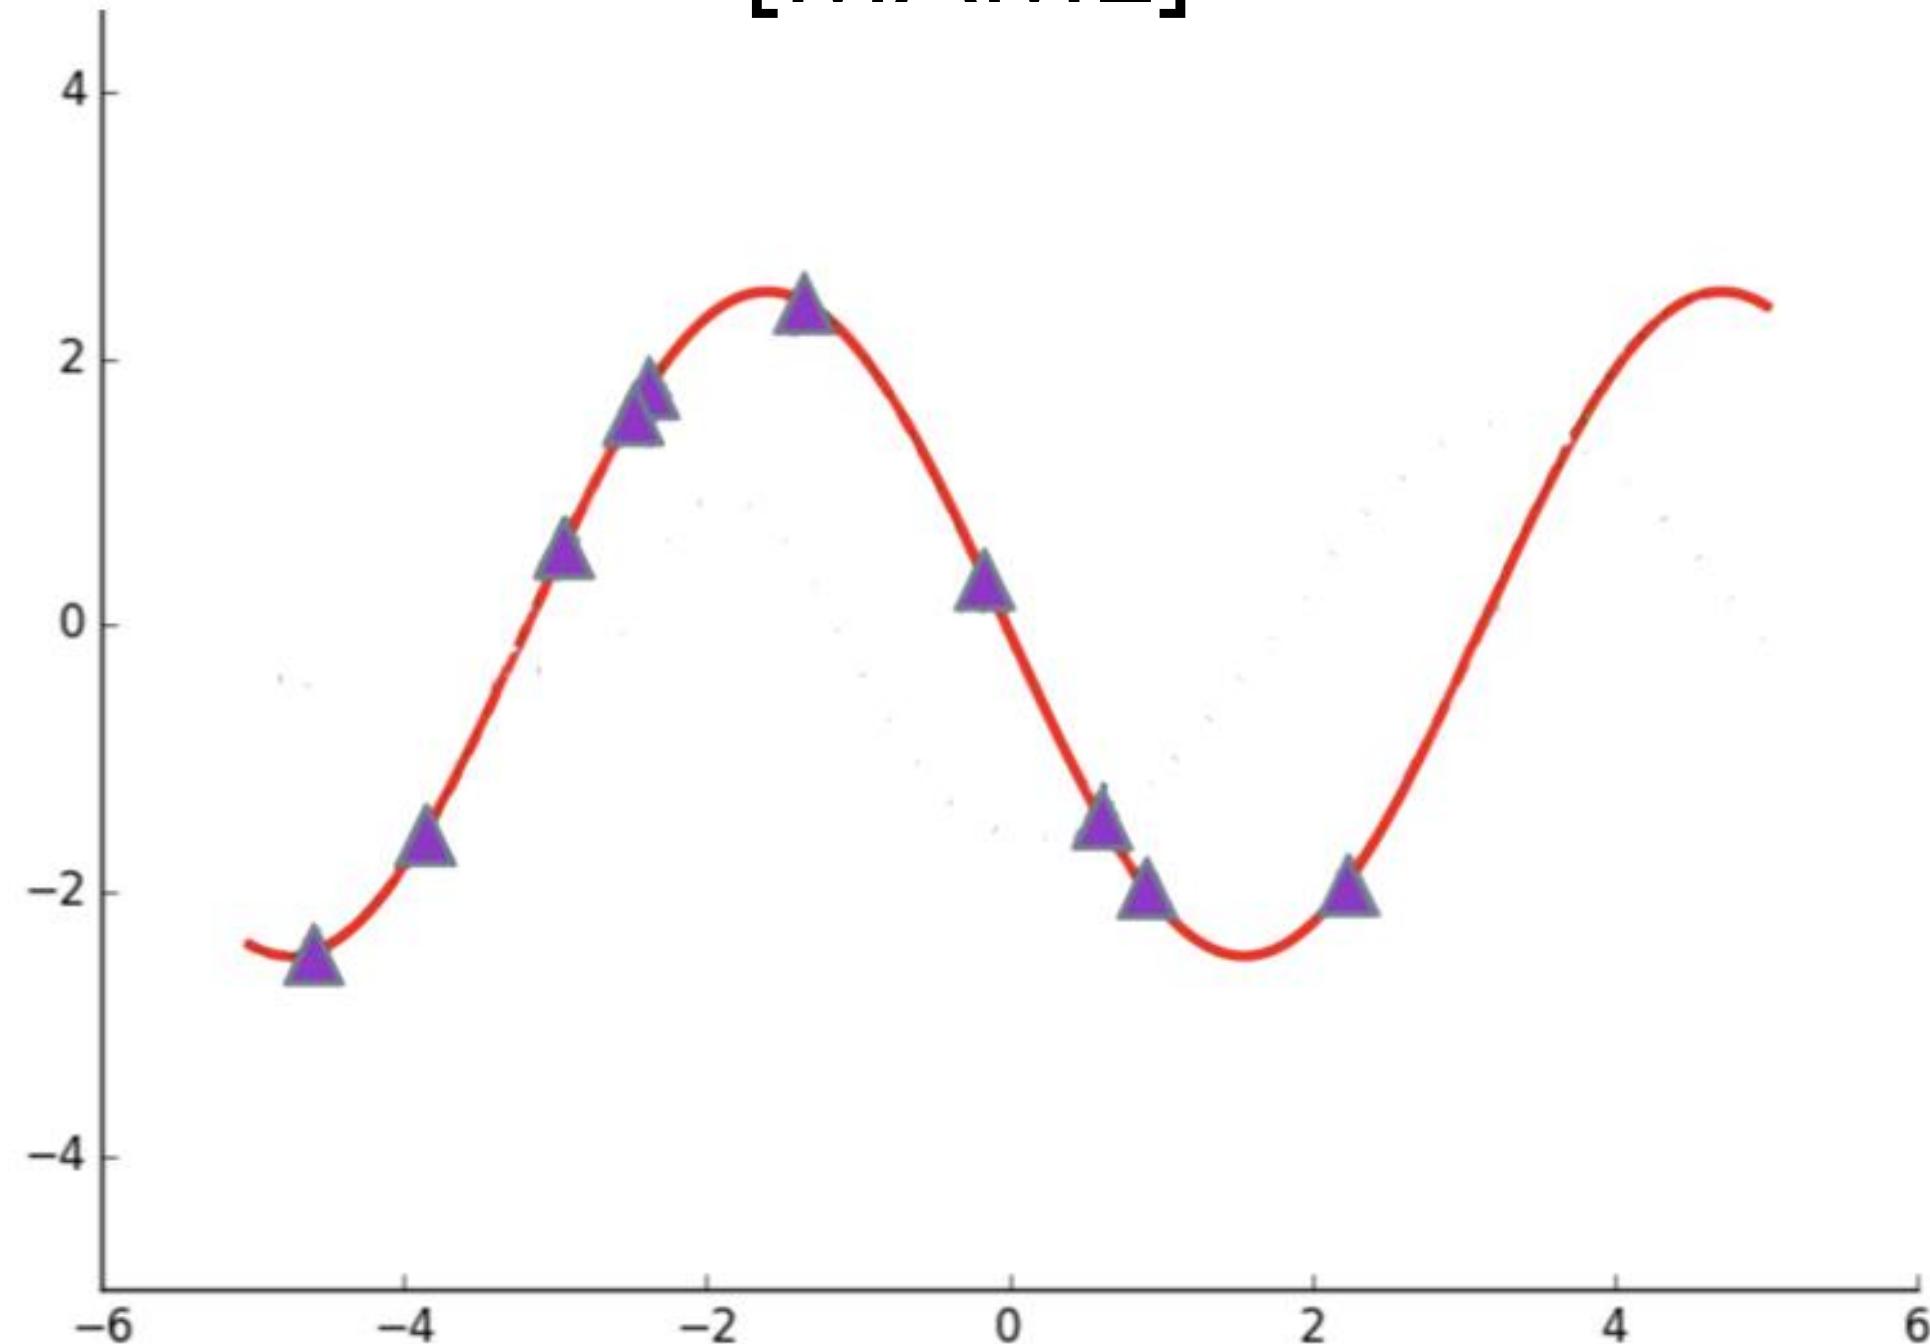
Experiments on Regression

- Sinusoid function:
 - Amplitude (A) and phase (ϕ) are varied between tasks
 - A in $[0.1, 0.5]$
 - ϕ in $[0, \pi]$
 - x in $[-5.0, 5.0]$
- Loss function: Mean Squared Error (MSE)
- Regressor: 2 hidden layers with 40 units and ReLU
- Training
 - Use **only 1 gradient step** for learner
 - $K = 5$ or 10 example (5-shot learning or 10-shot learning)
 - Fixed step size ($\alpha=0.01$) for Adam optimizer.

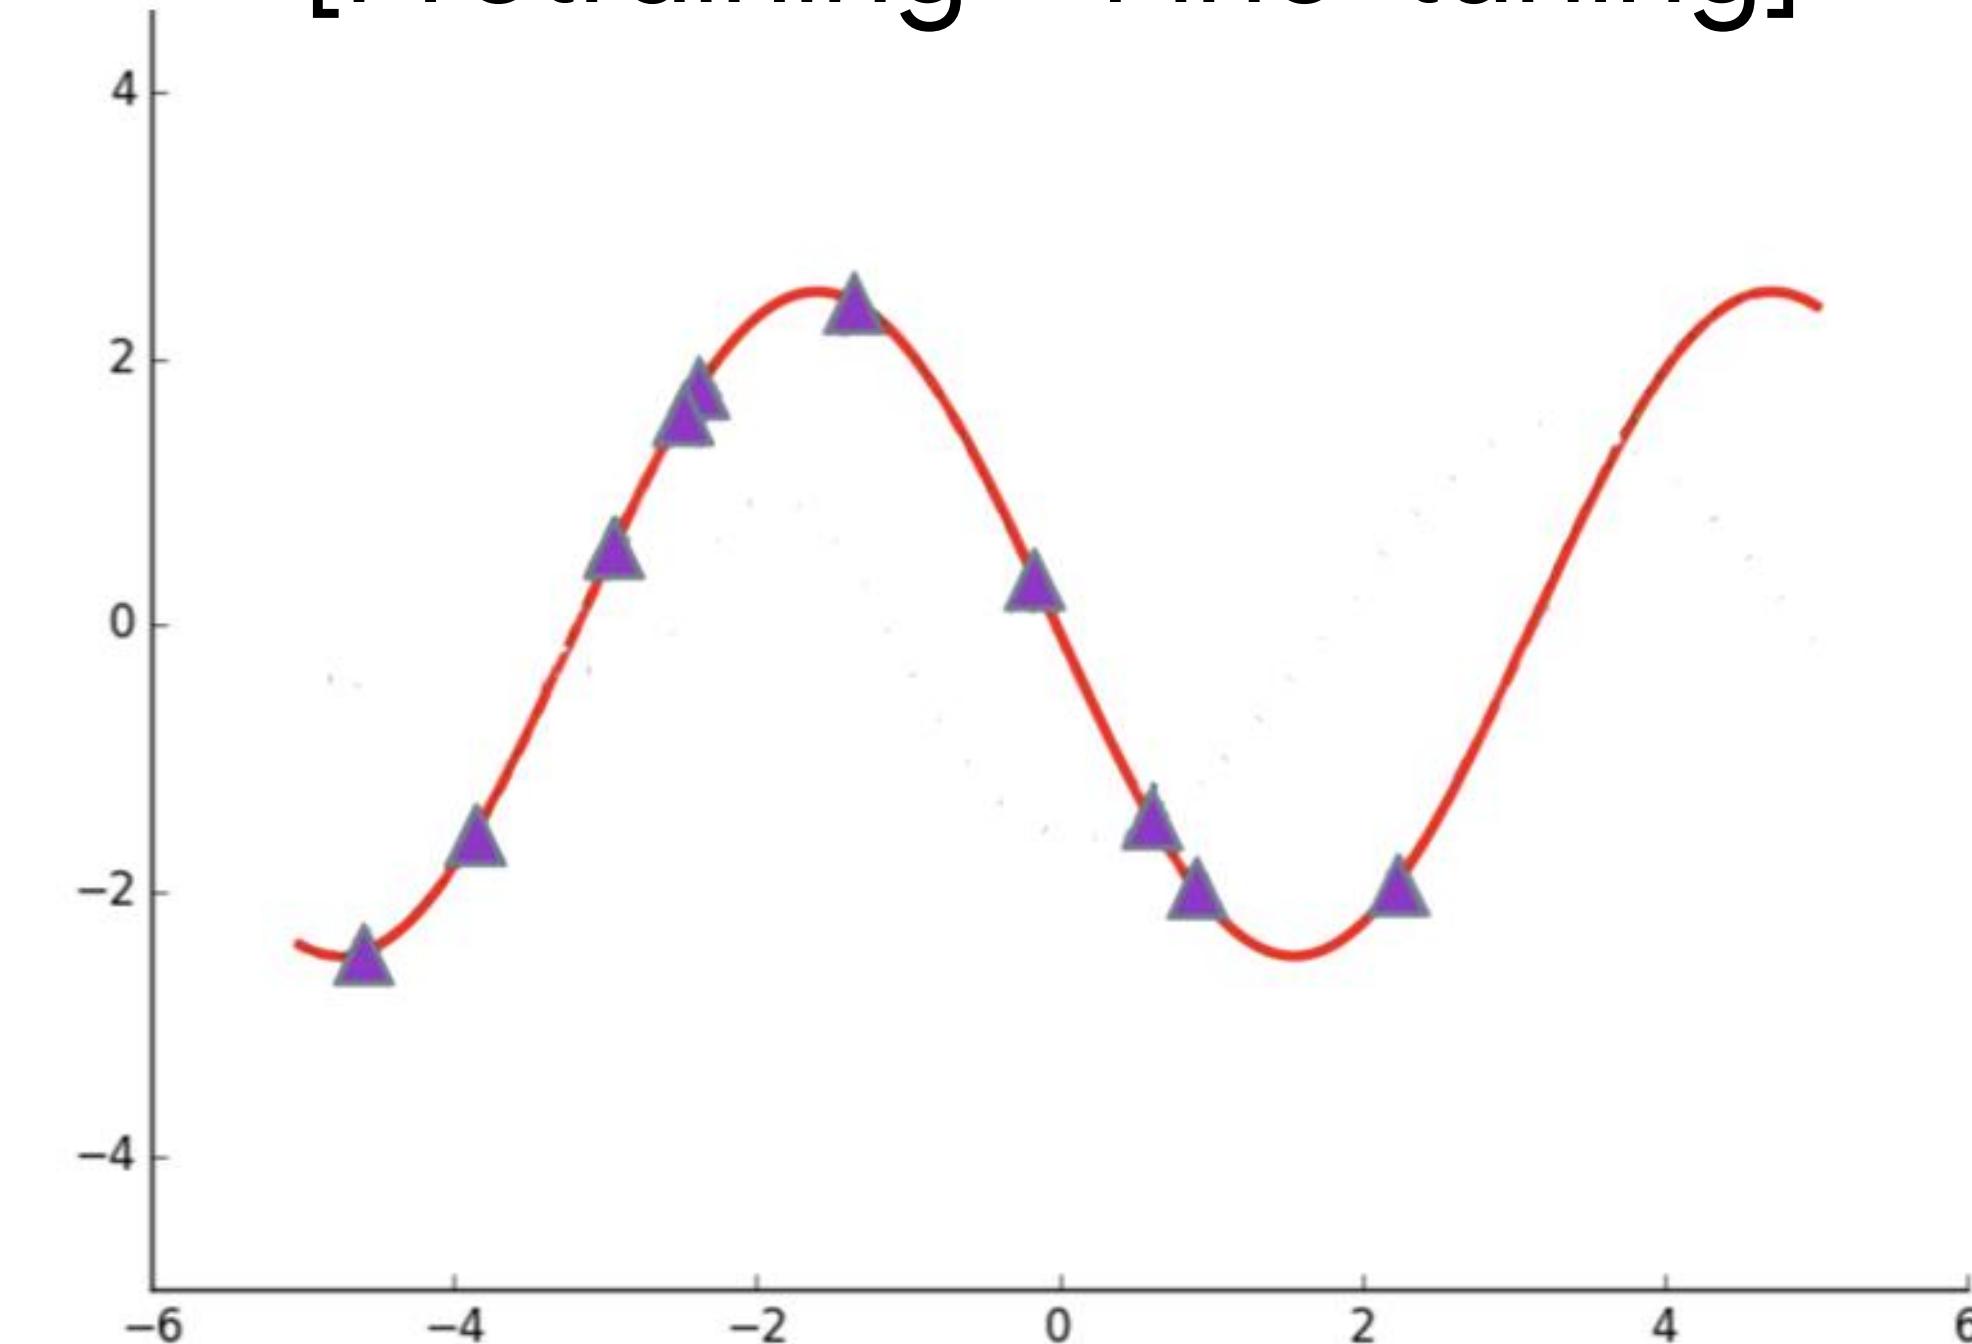
5. Experiment

Results of 10-Shot Learning Regression

[MAML]



[Pretraining + Fine-tuning]

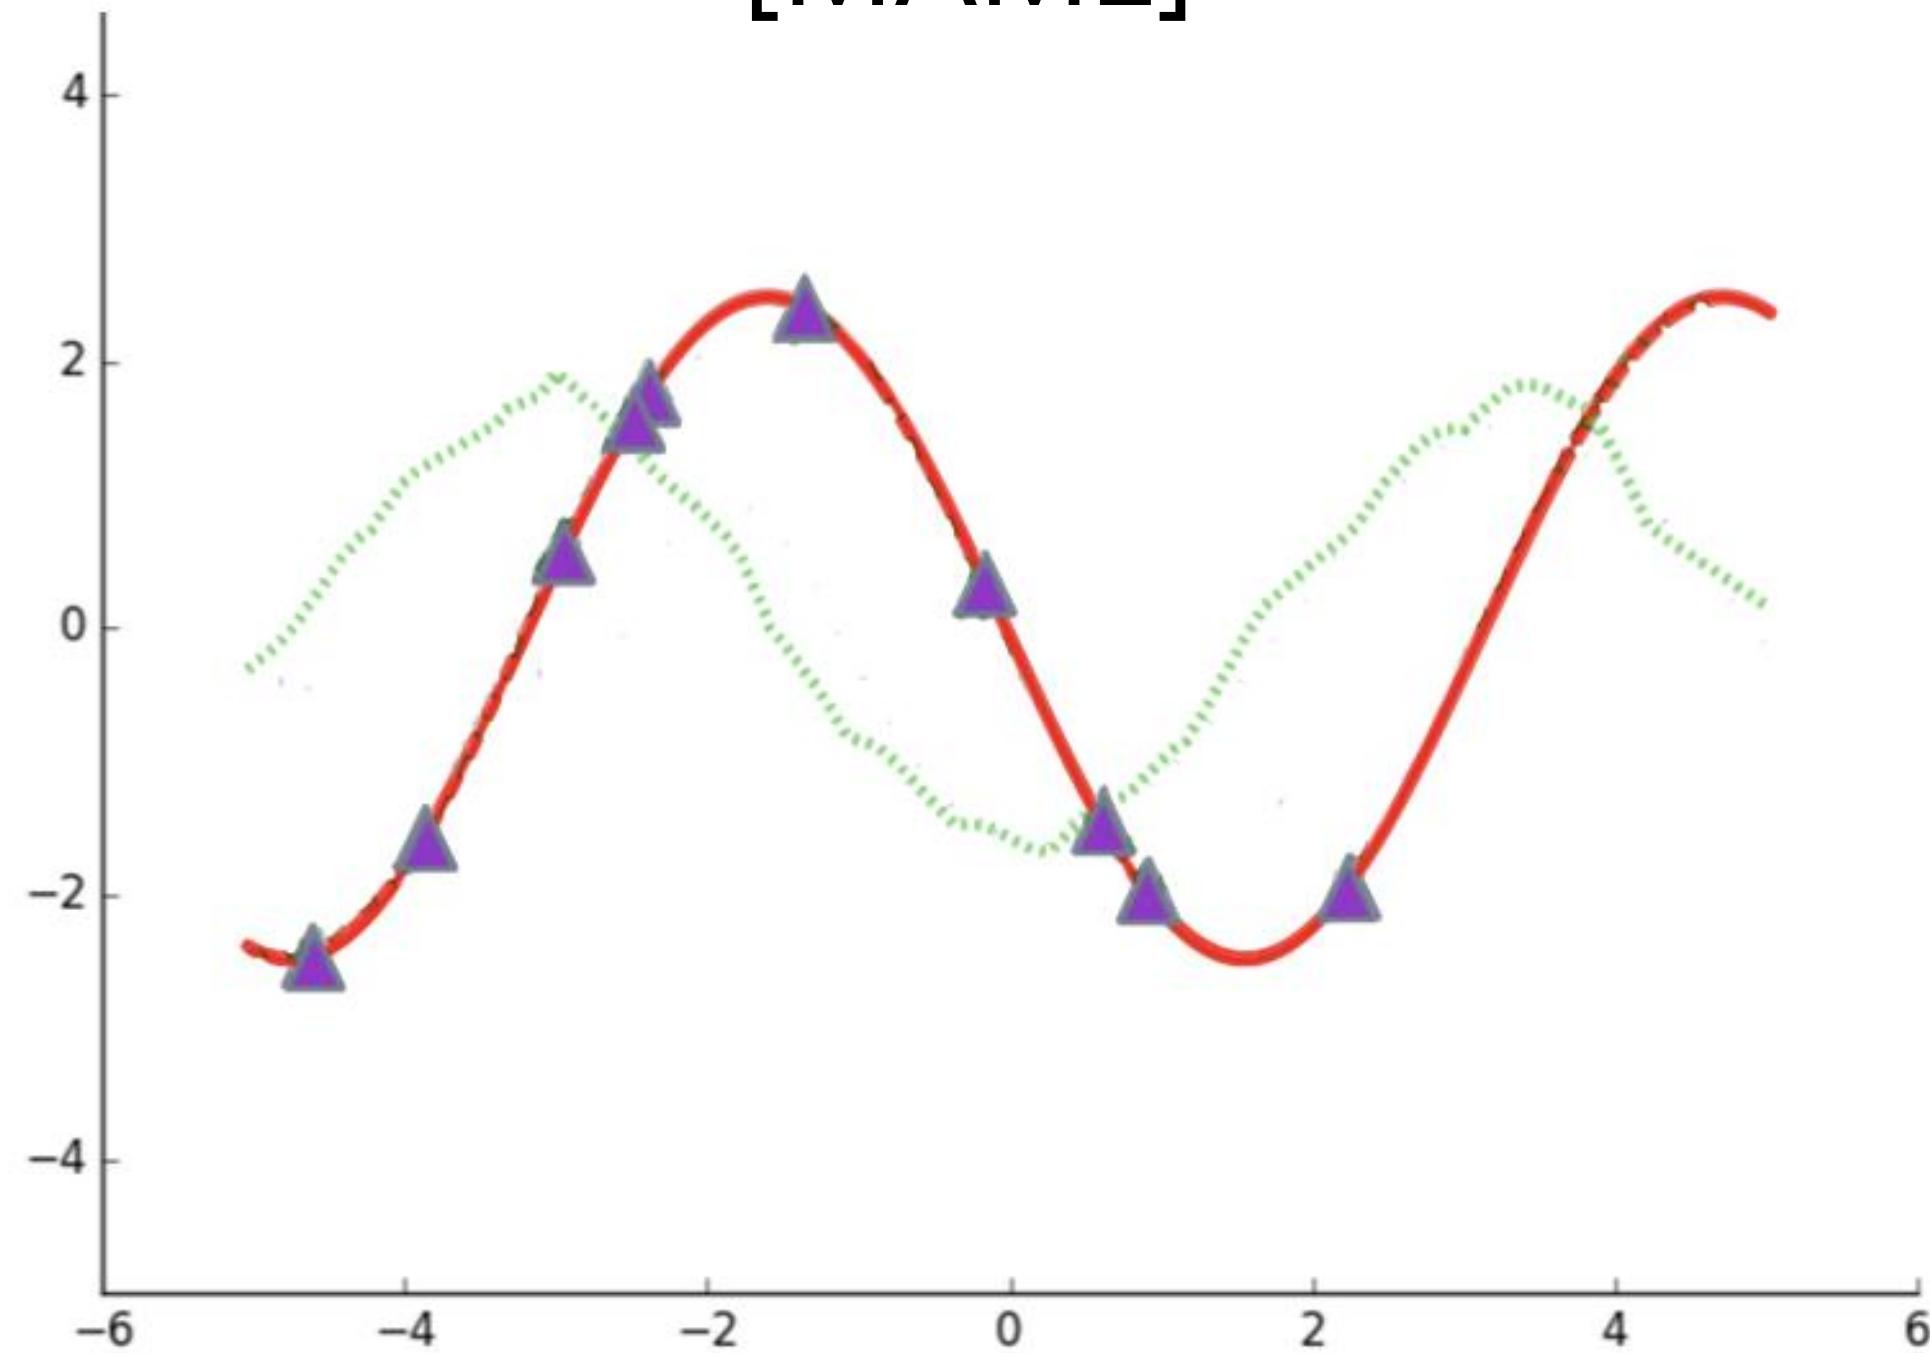


The red line is ground truth.
Fit this sine function with only few (10) samples.

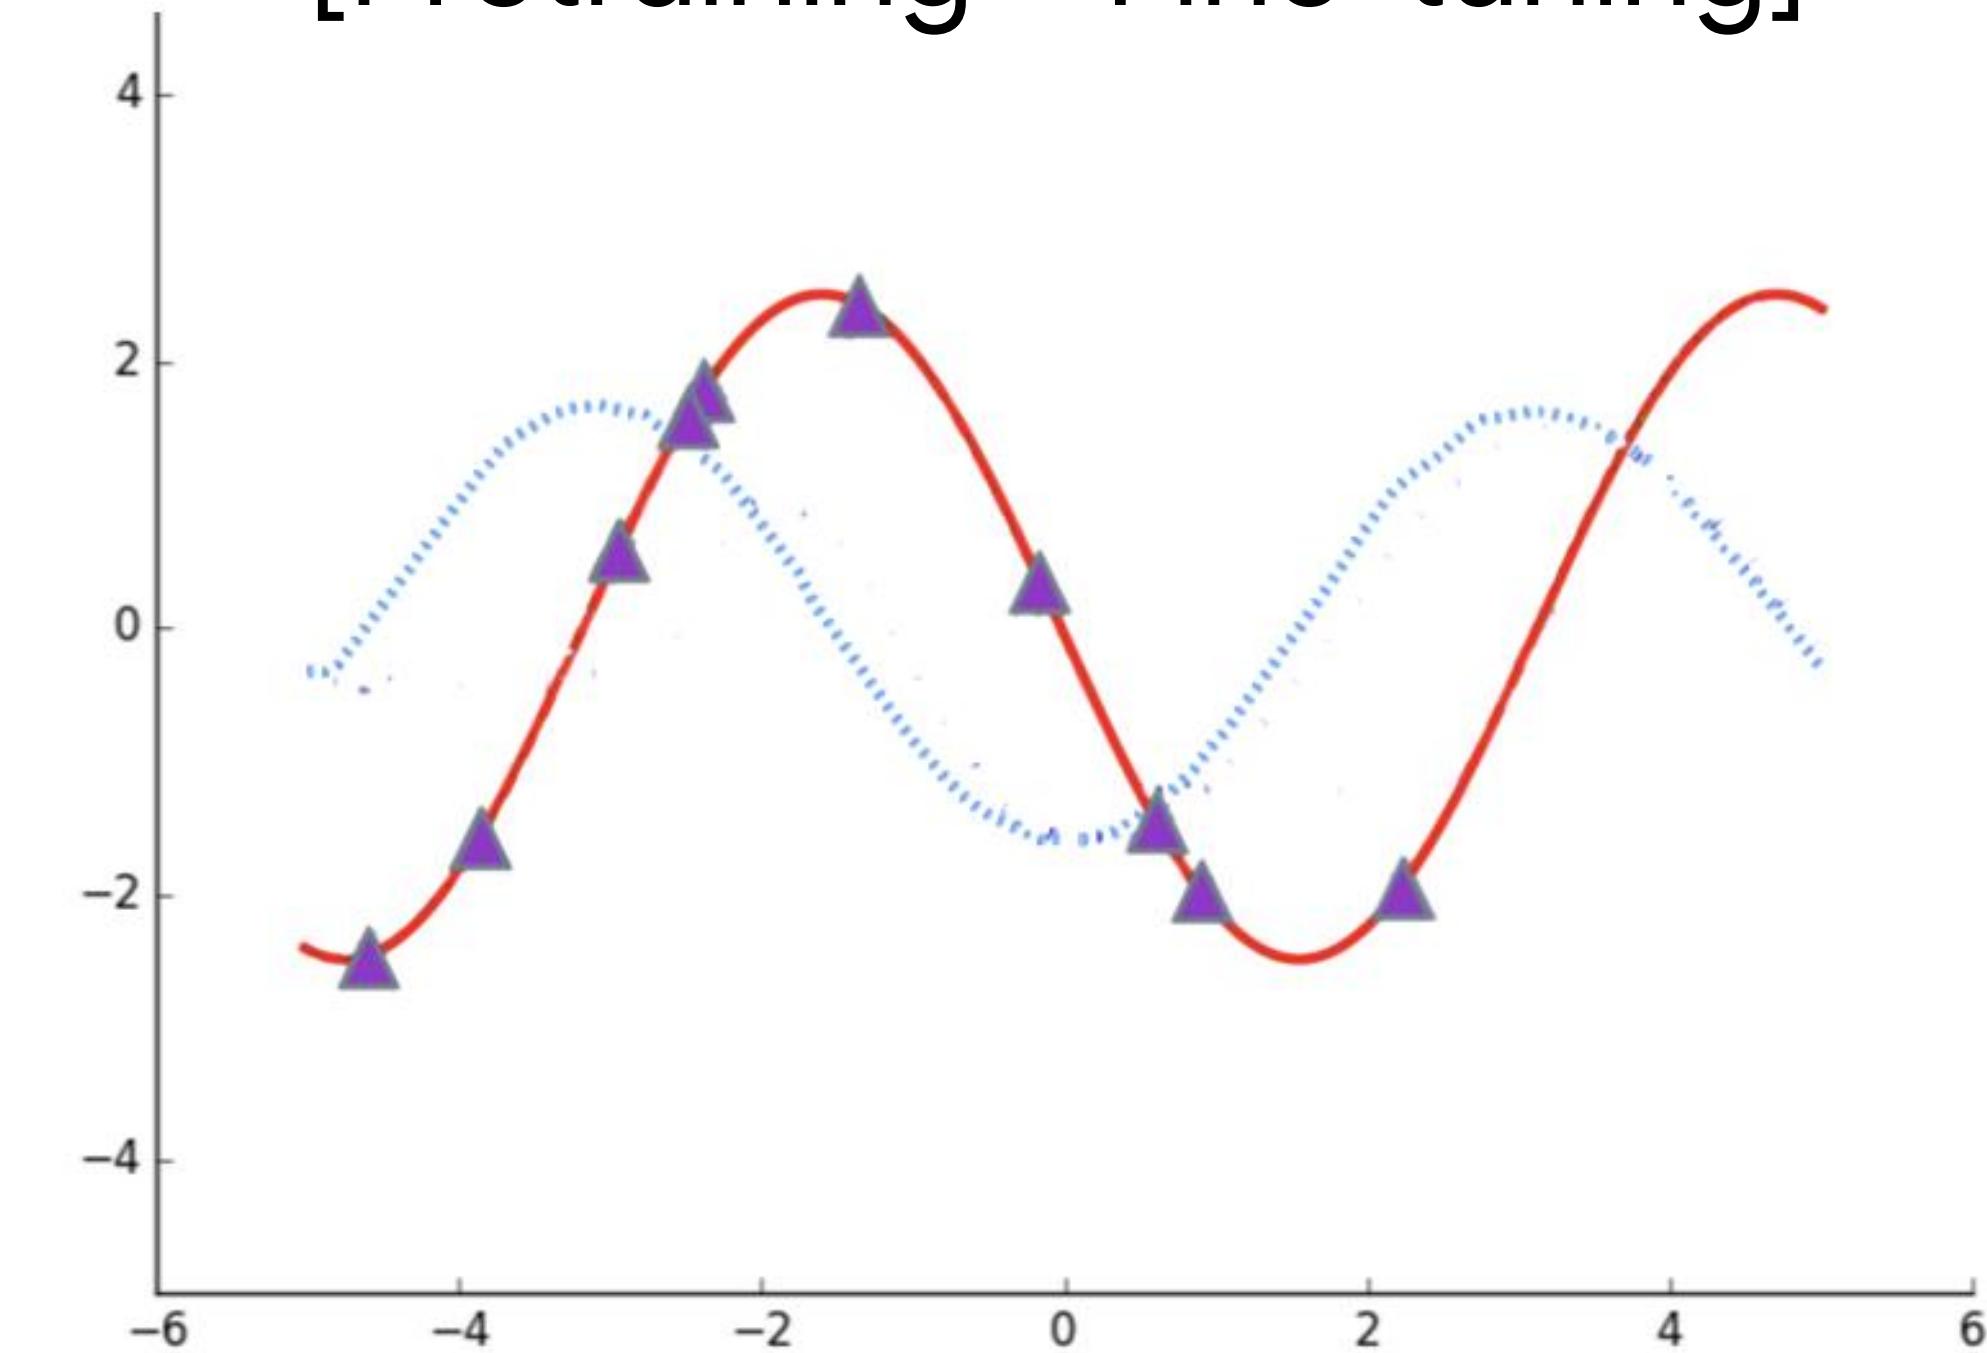
5. Experiment

Results of 10-Shot Learning Regression

[MAML]



[Pretraining + Fine-tuning]



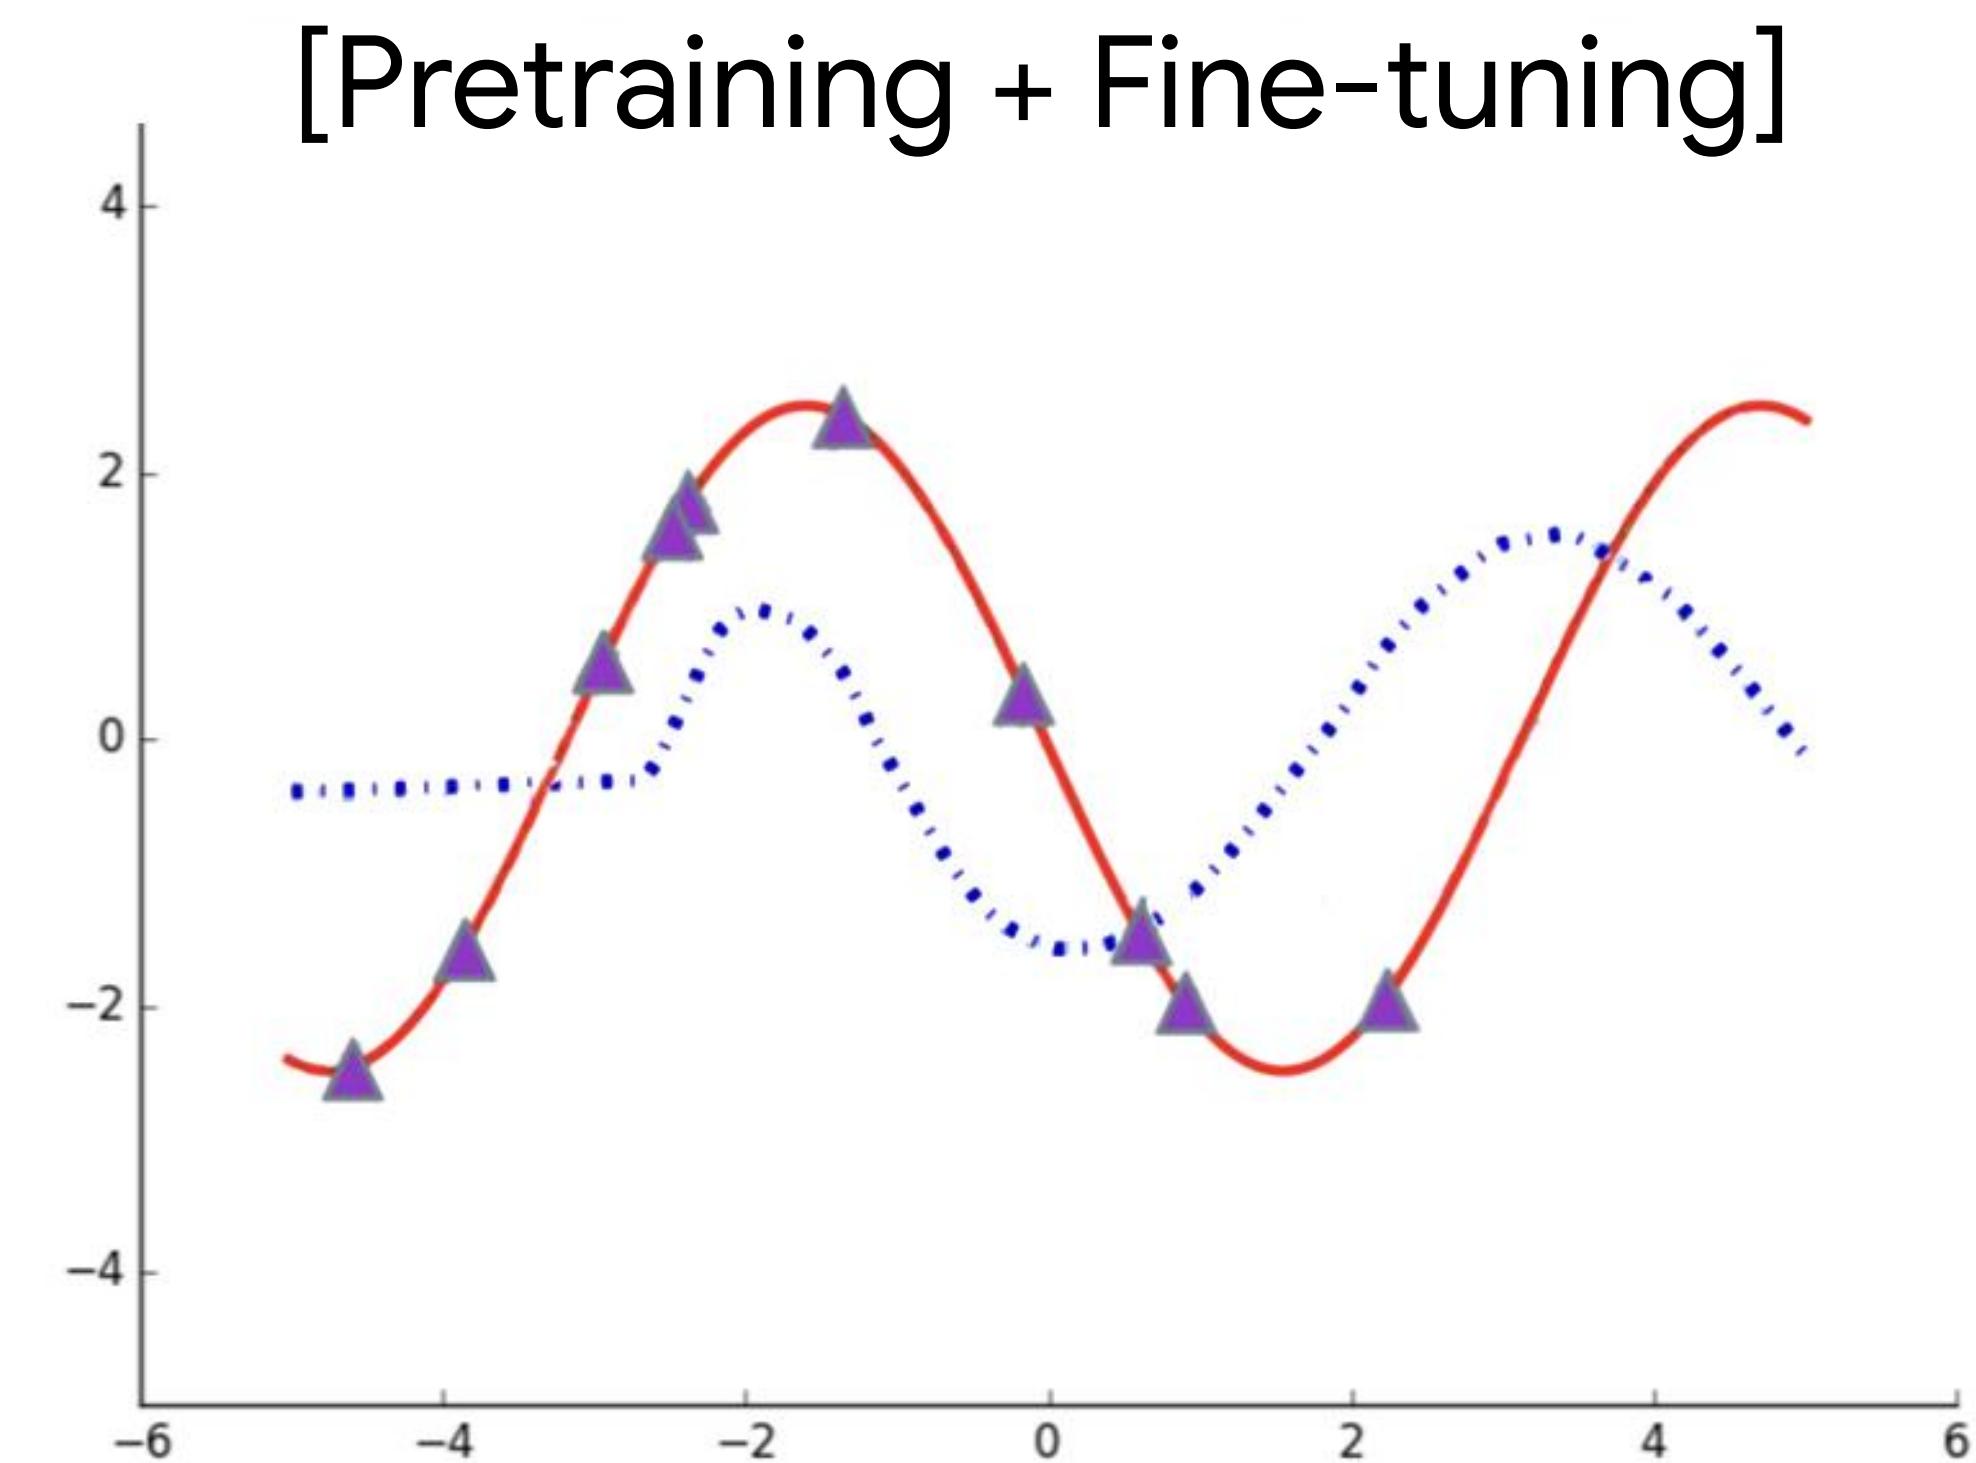
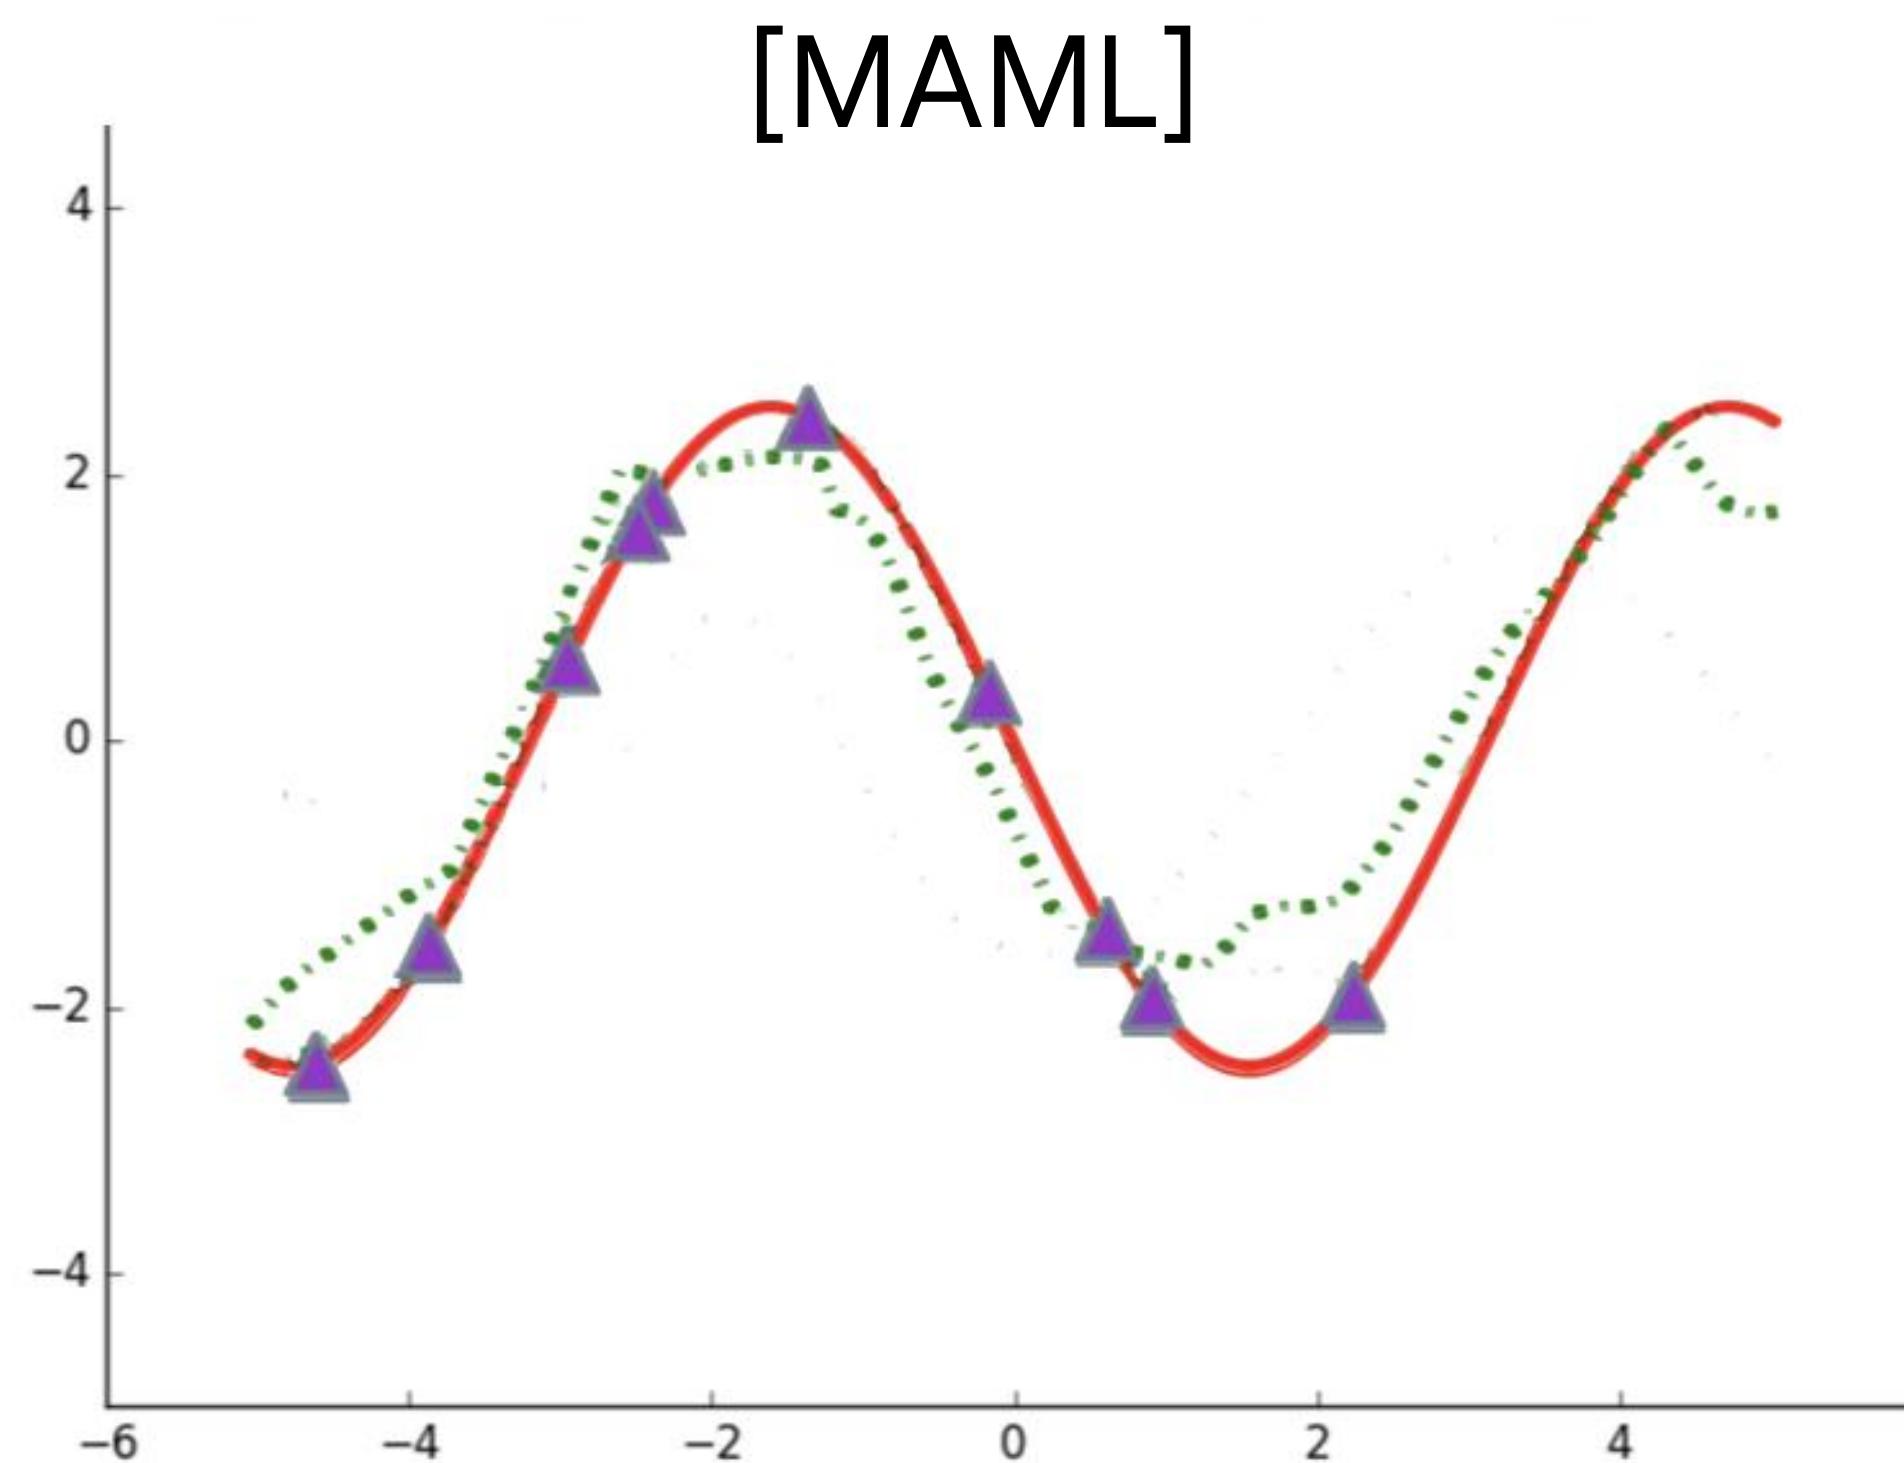
Above plots are the pre-trained function of two models.

(The prediction of meta-parameter of MAML,

The prediction of co-learned parameter of vanilla multi-task learning)

5. Experiment

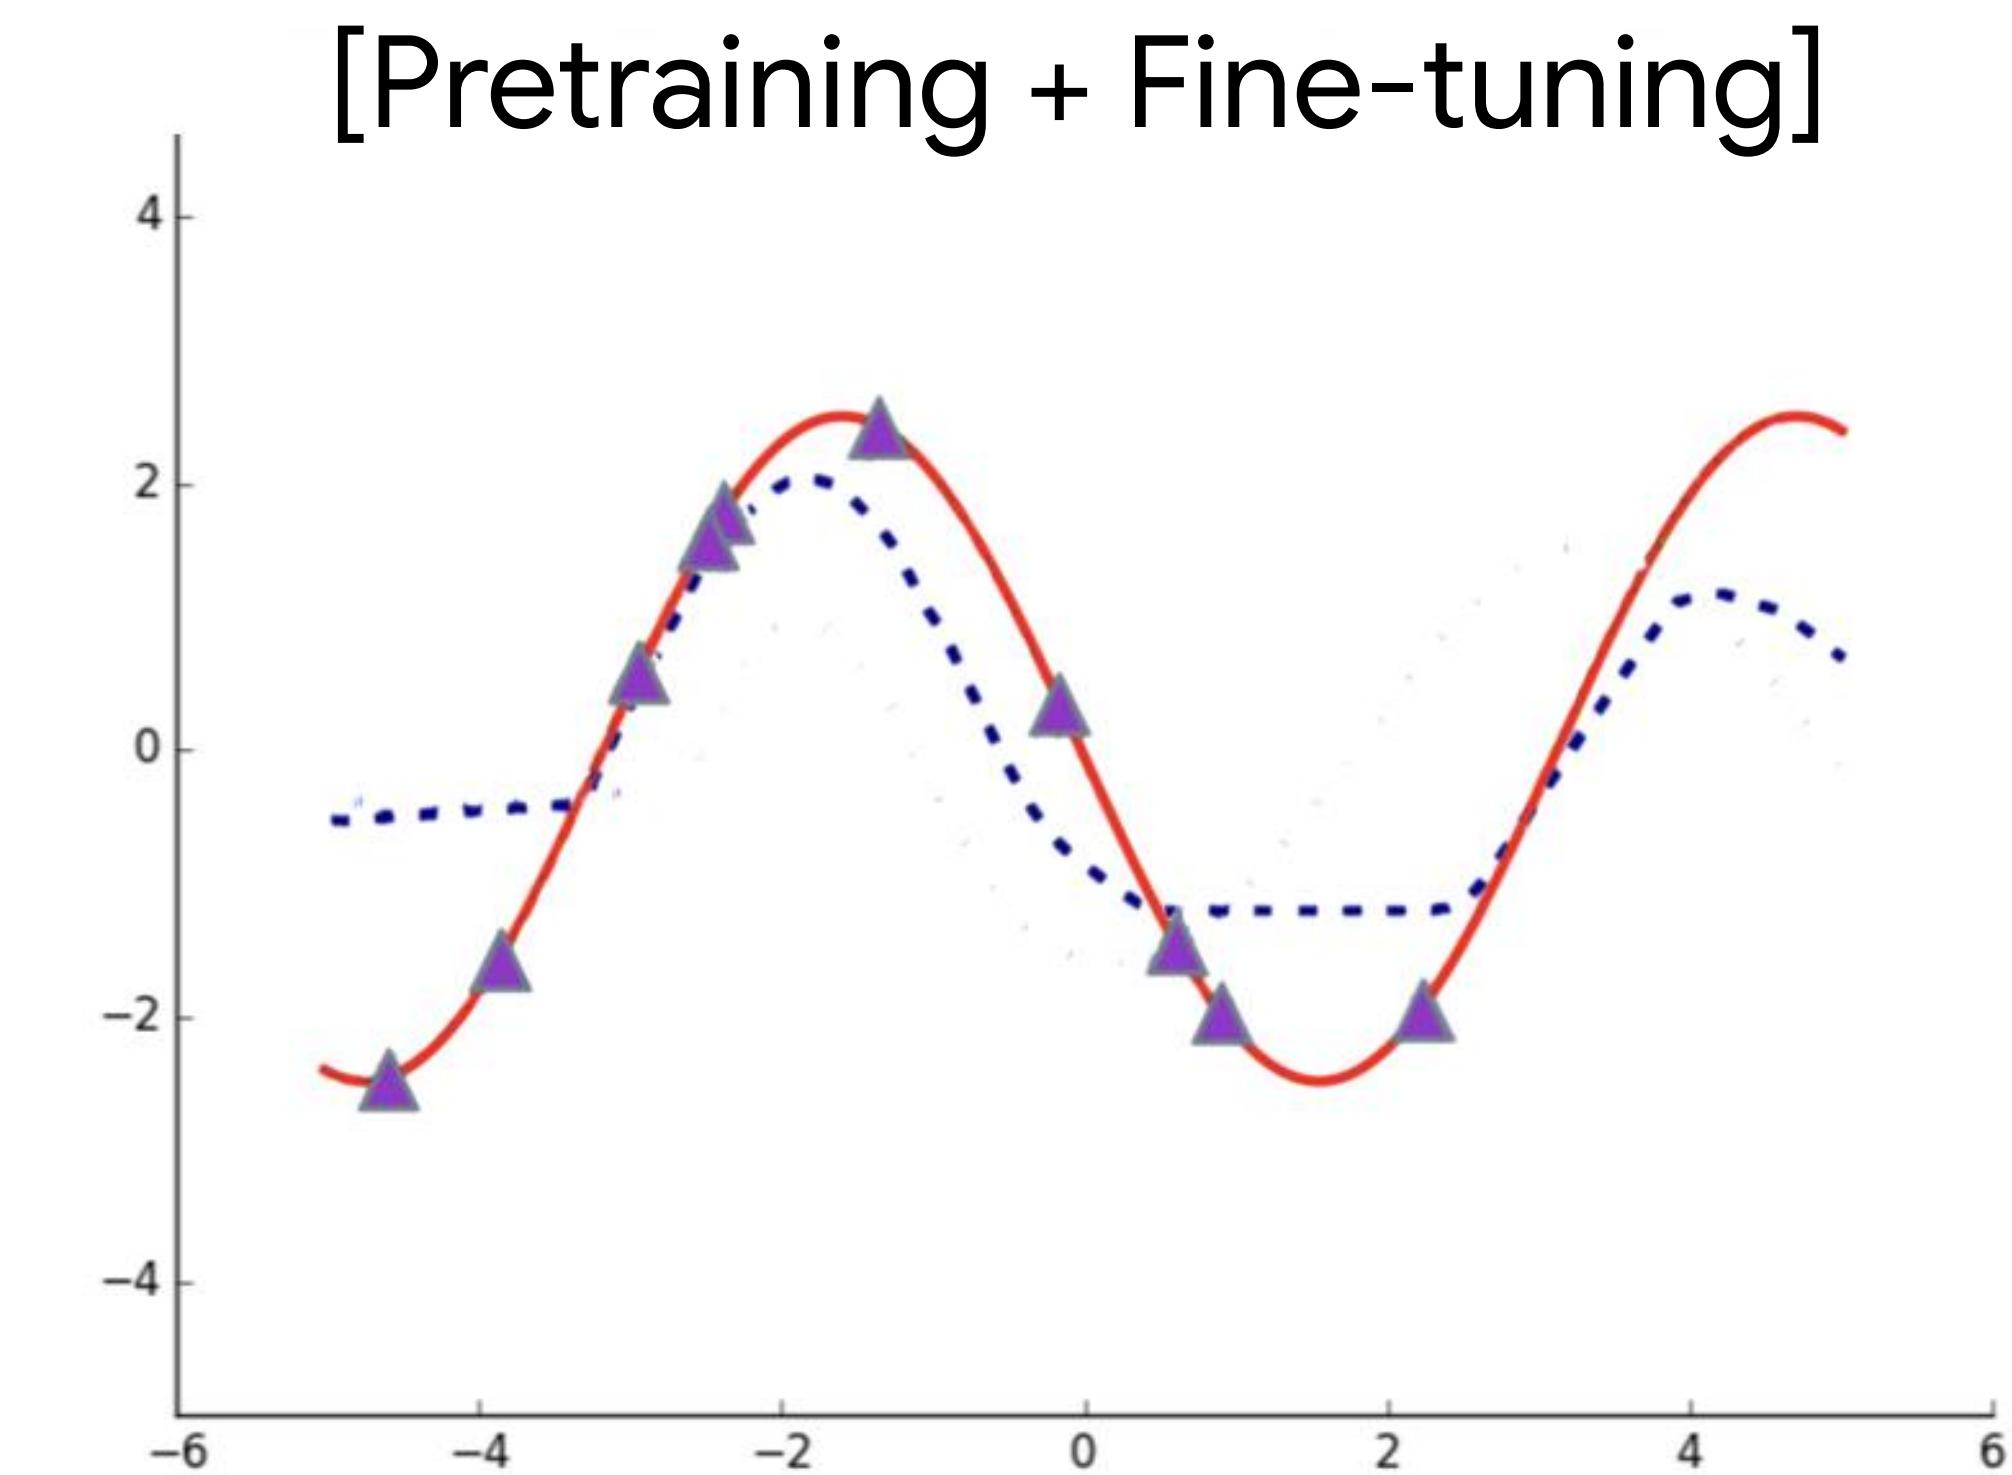
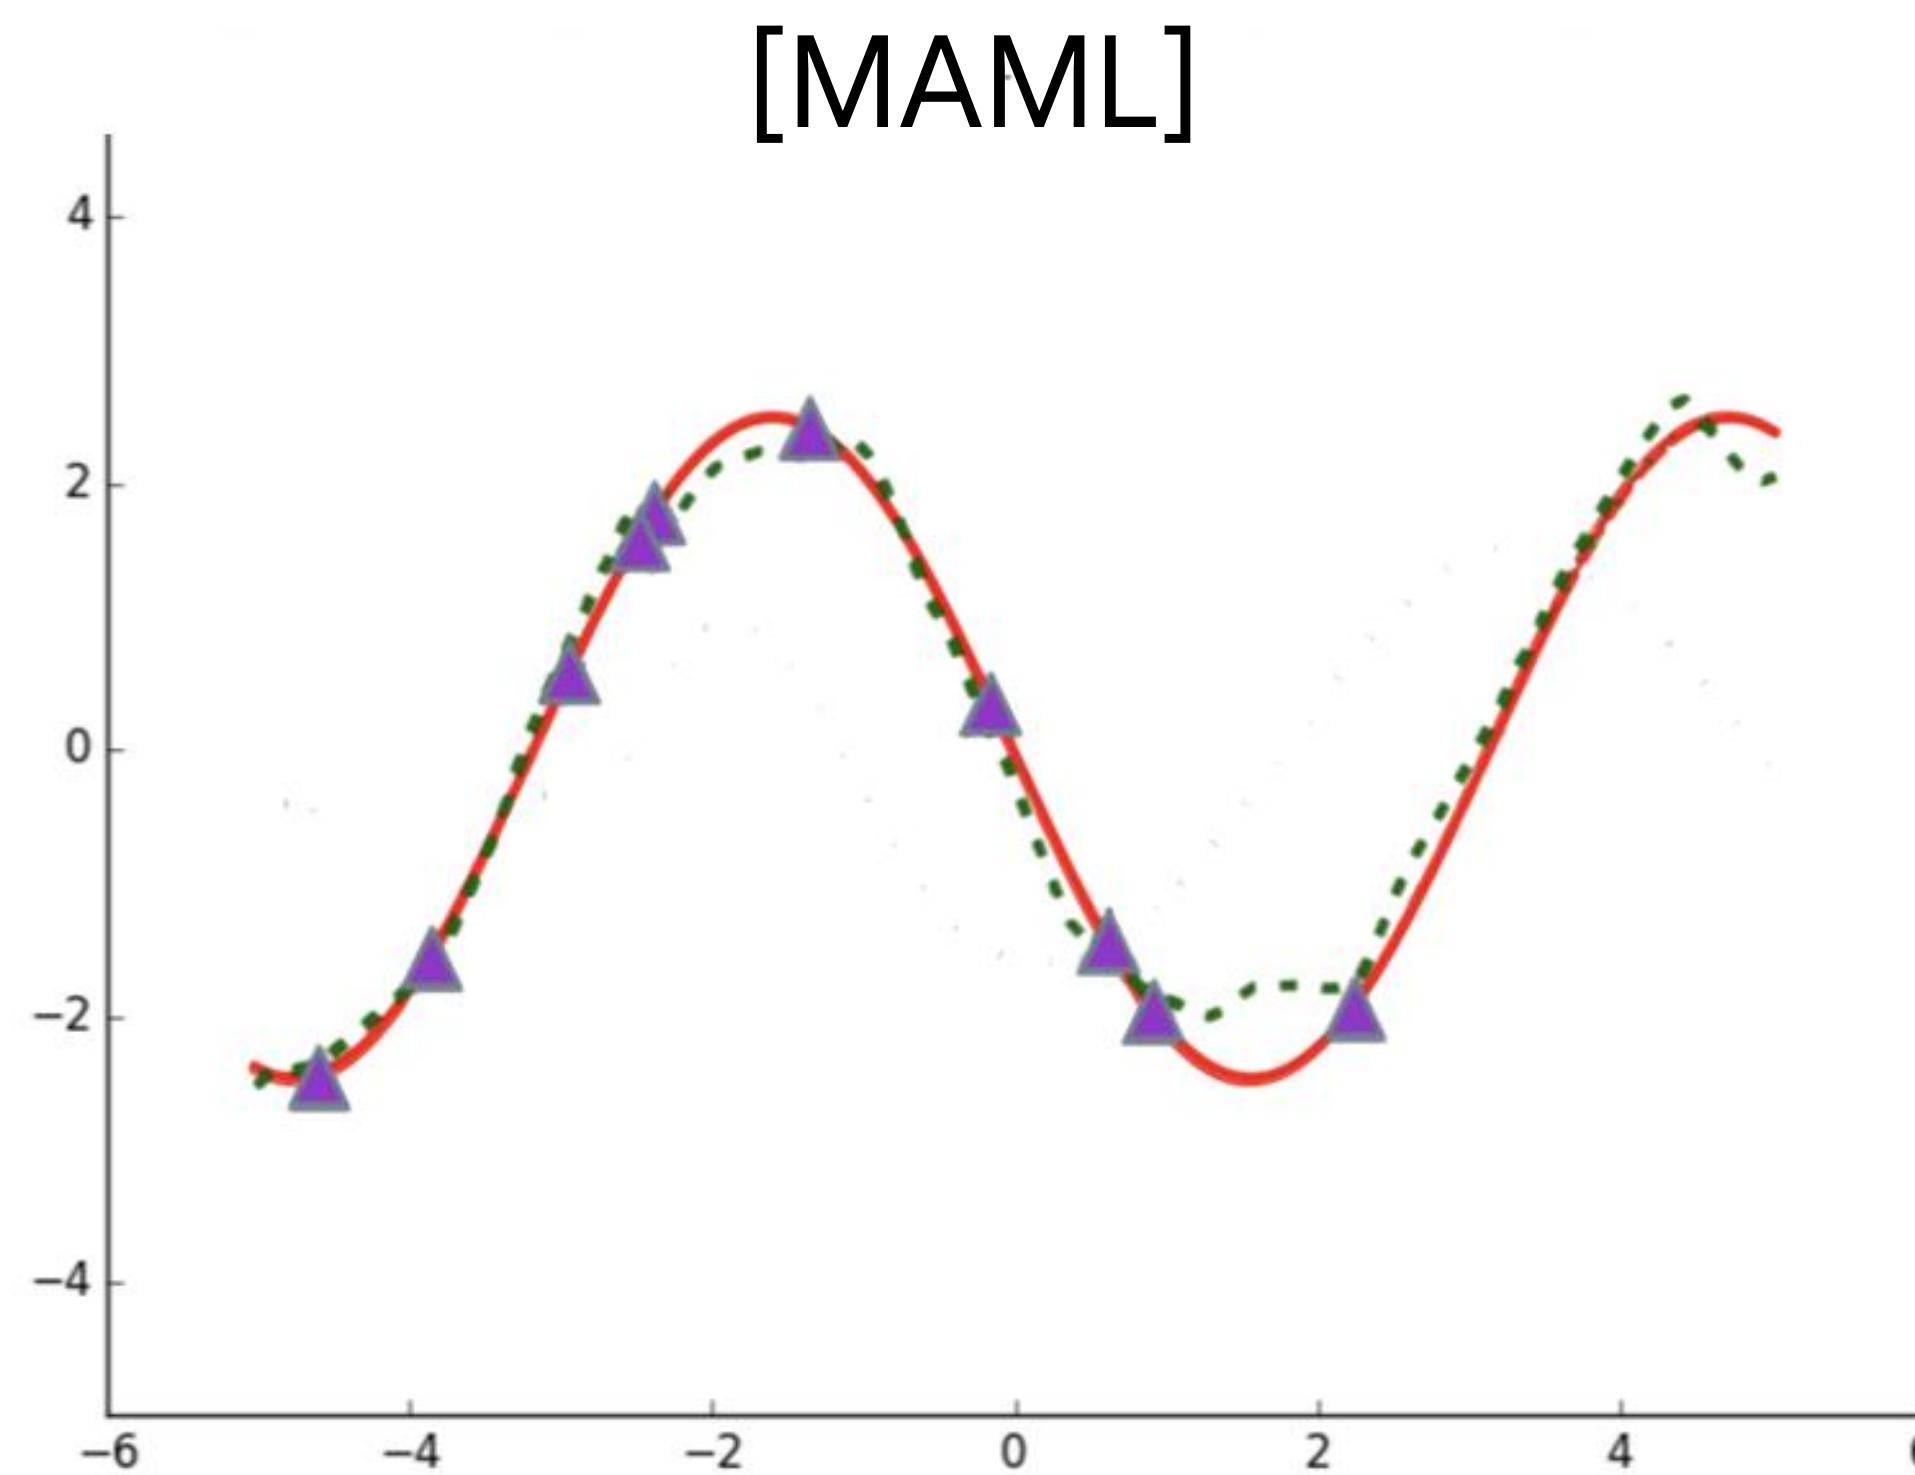
Results of 10-Shot Learning Regression



After 1 gradient step update.

5. Experiment

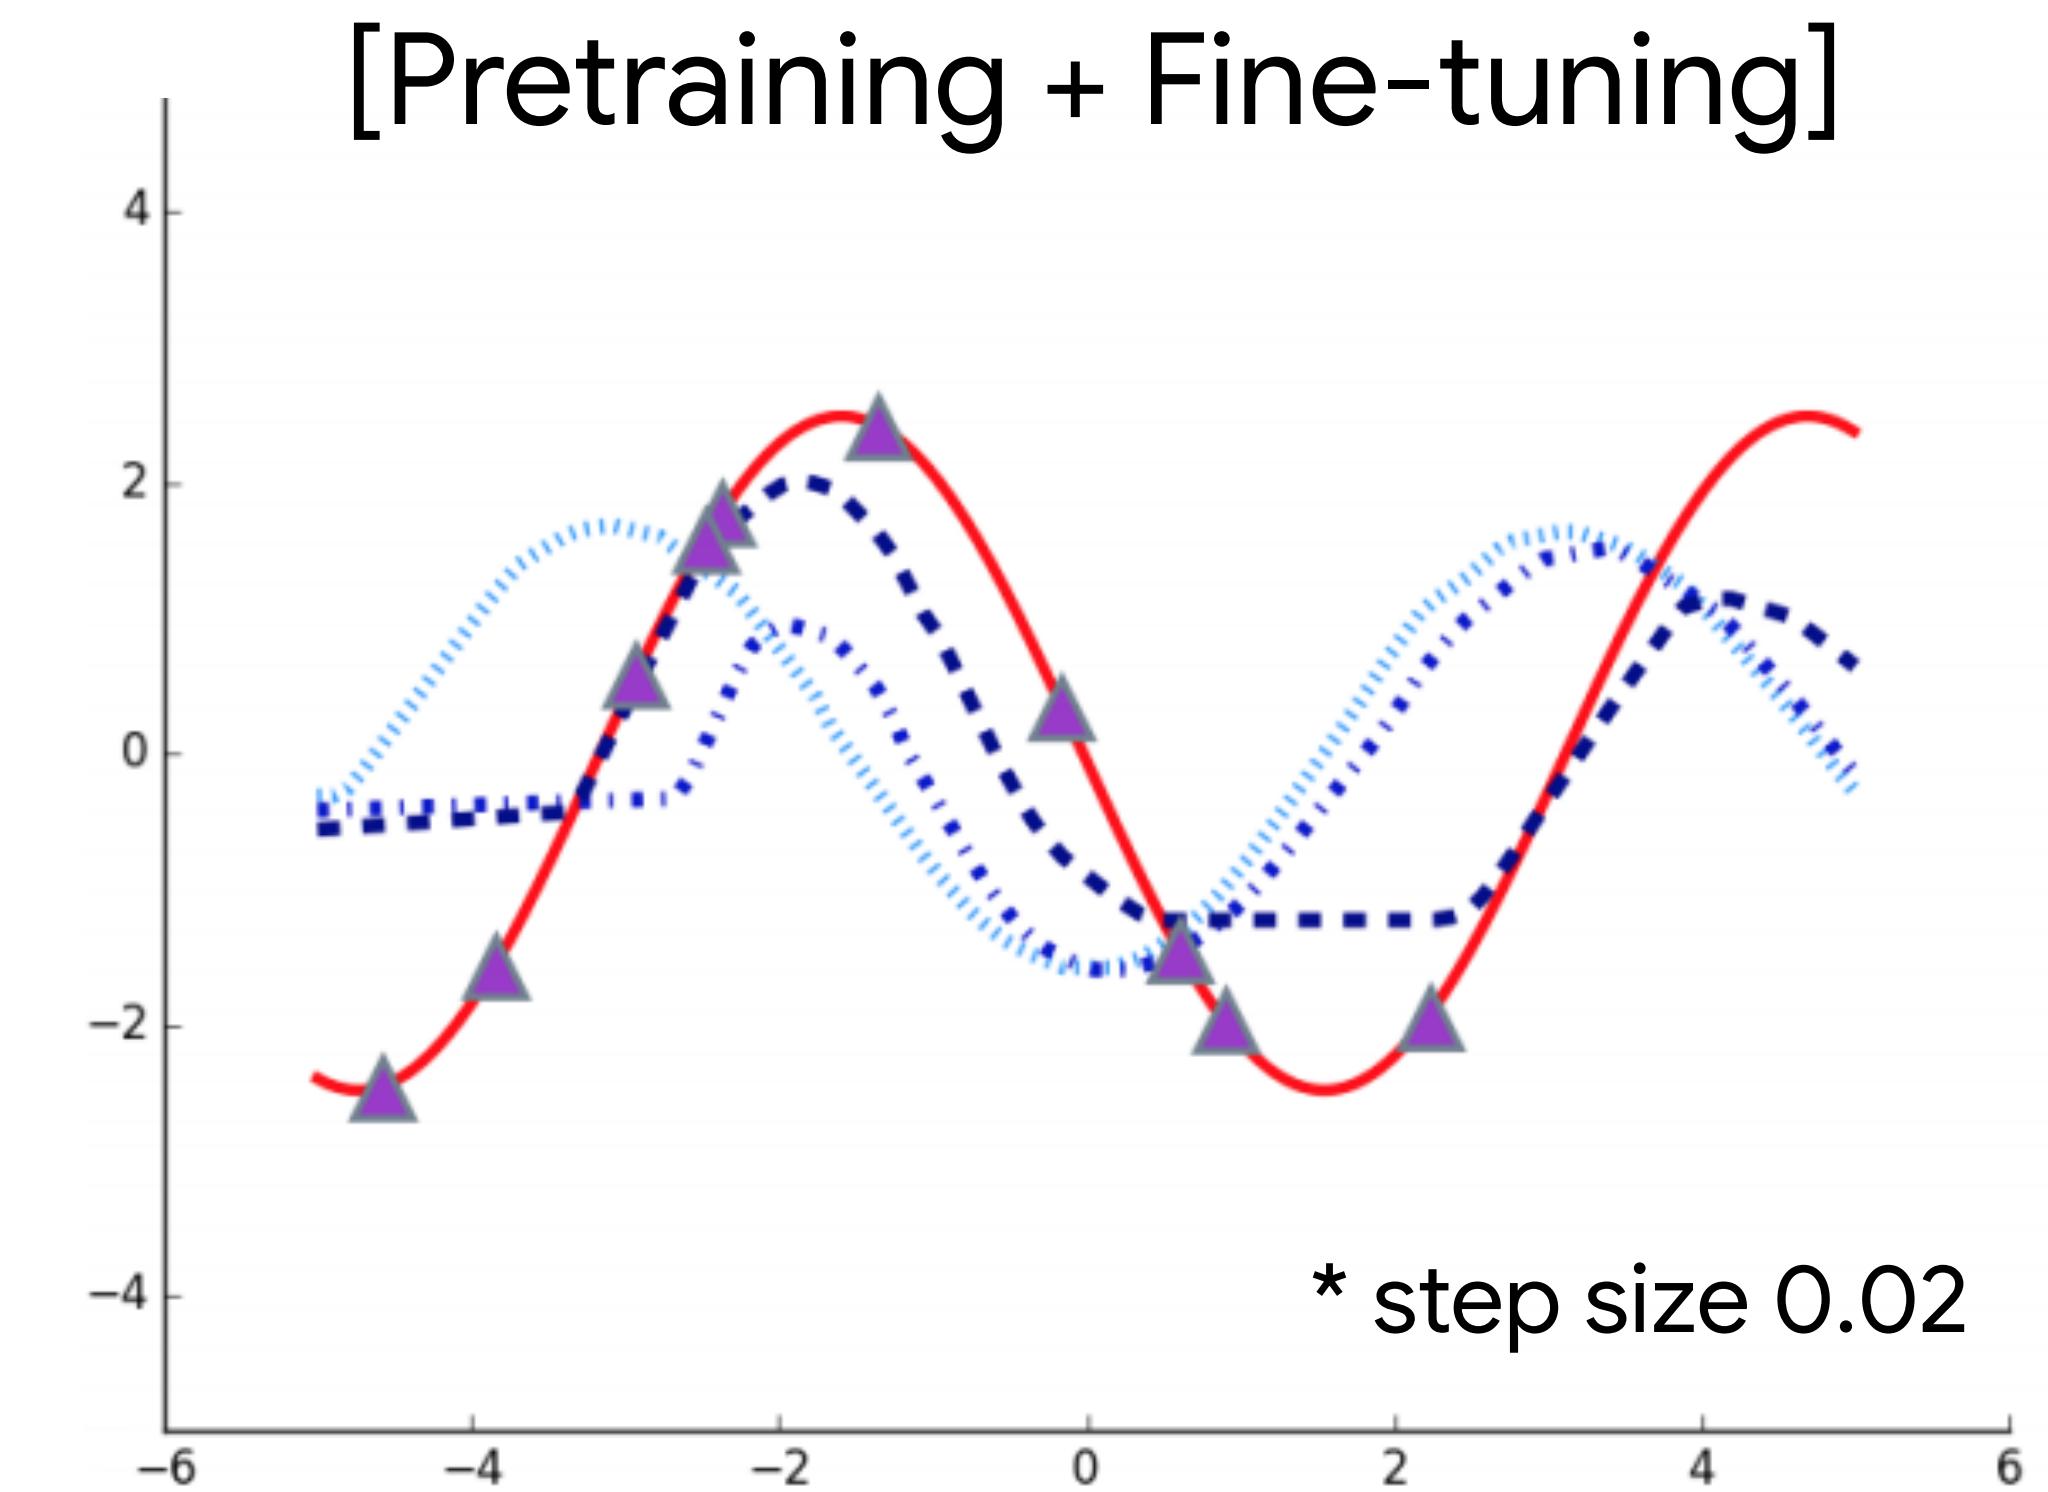
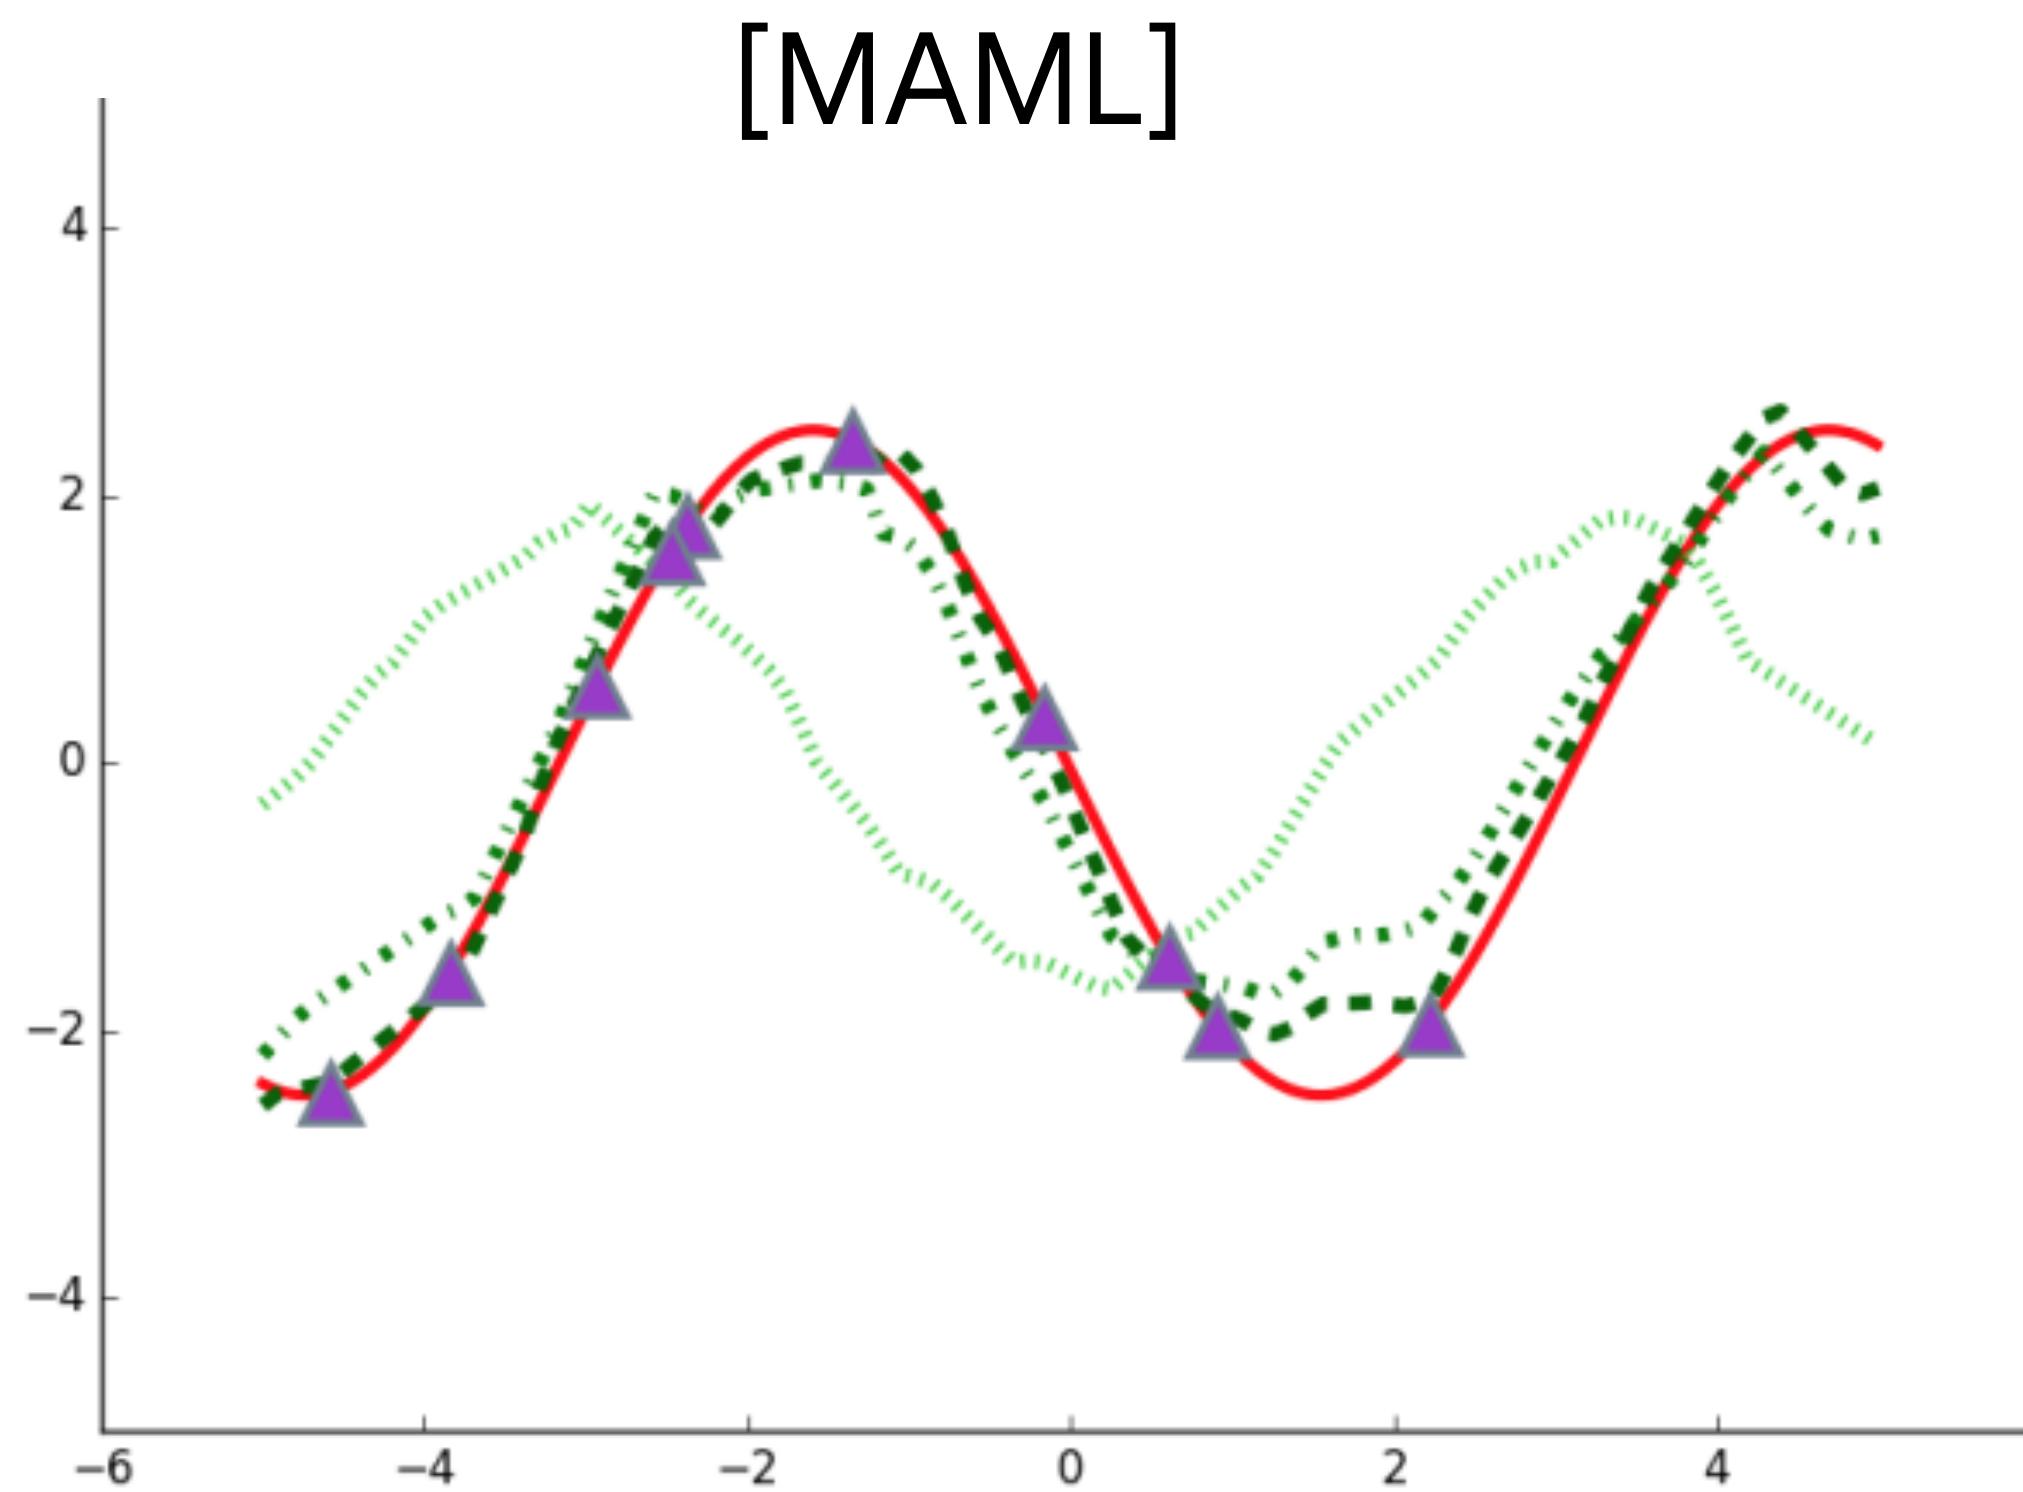
Results of 10-Shot Learning Regression



After 10 gradient step update.

5. Experiment

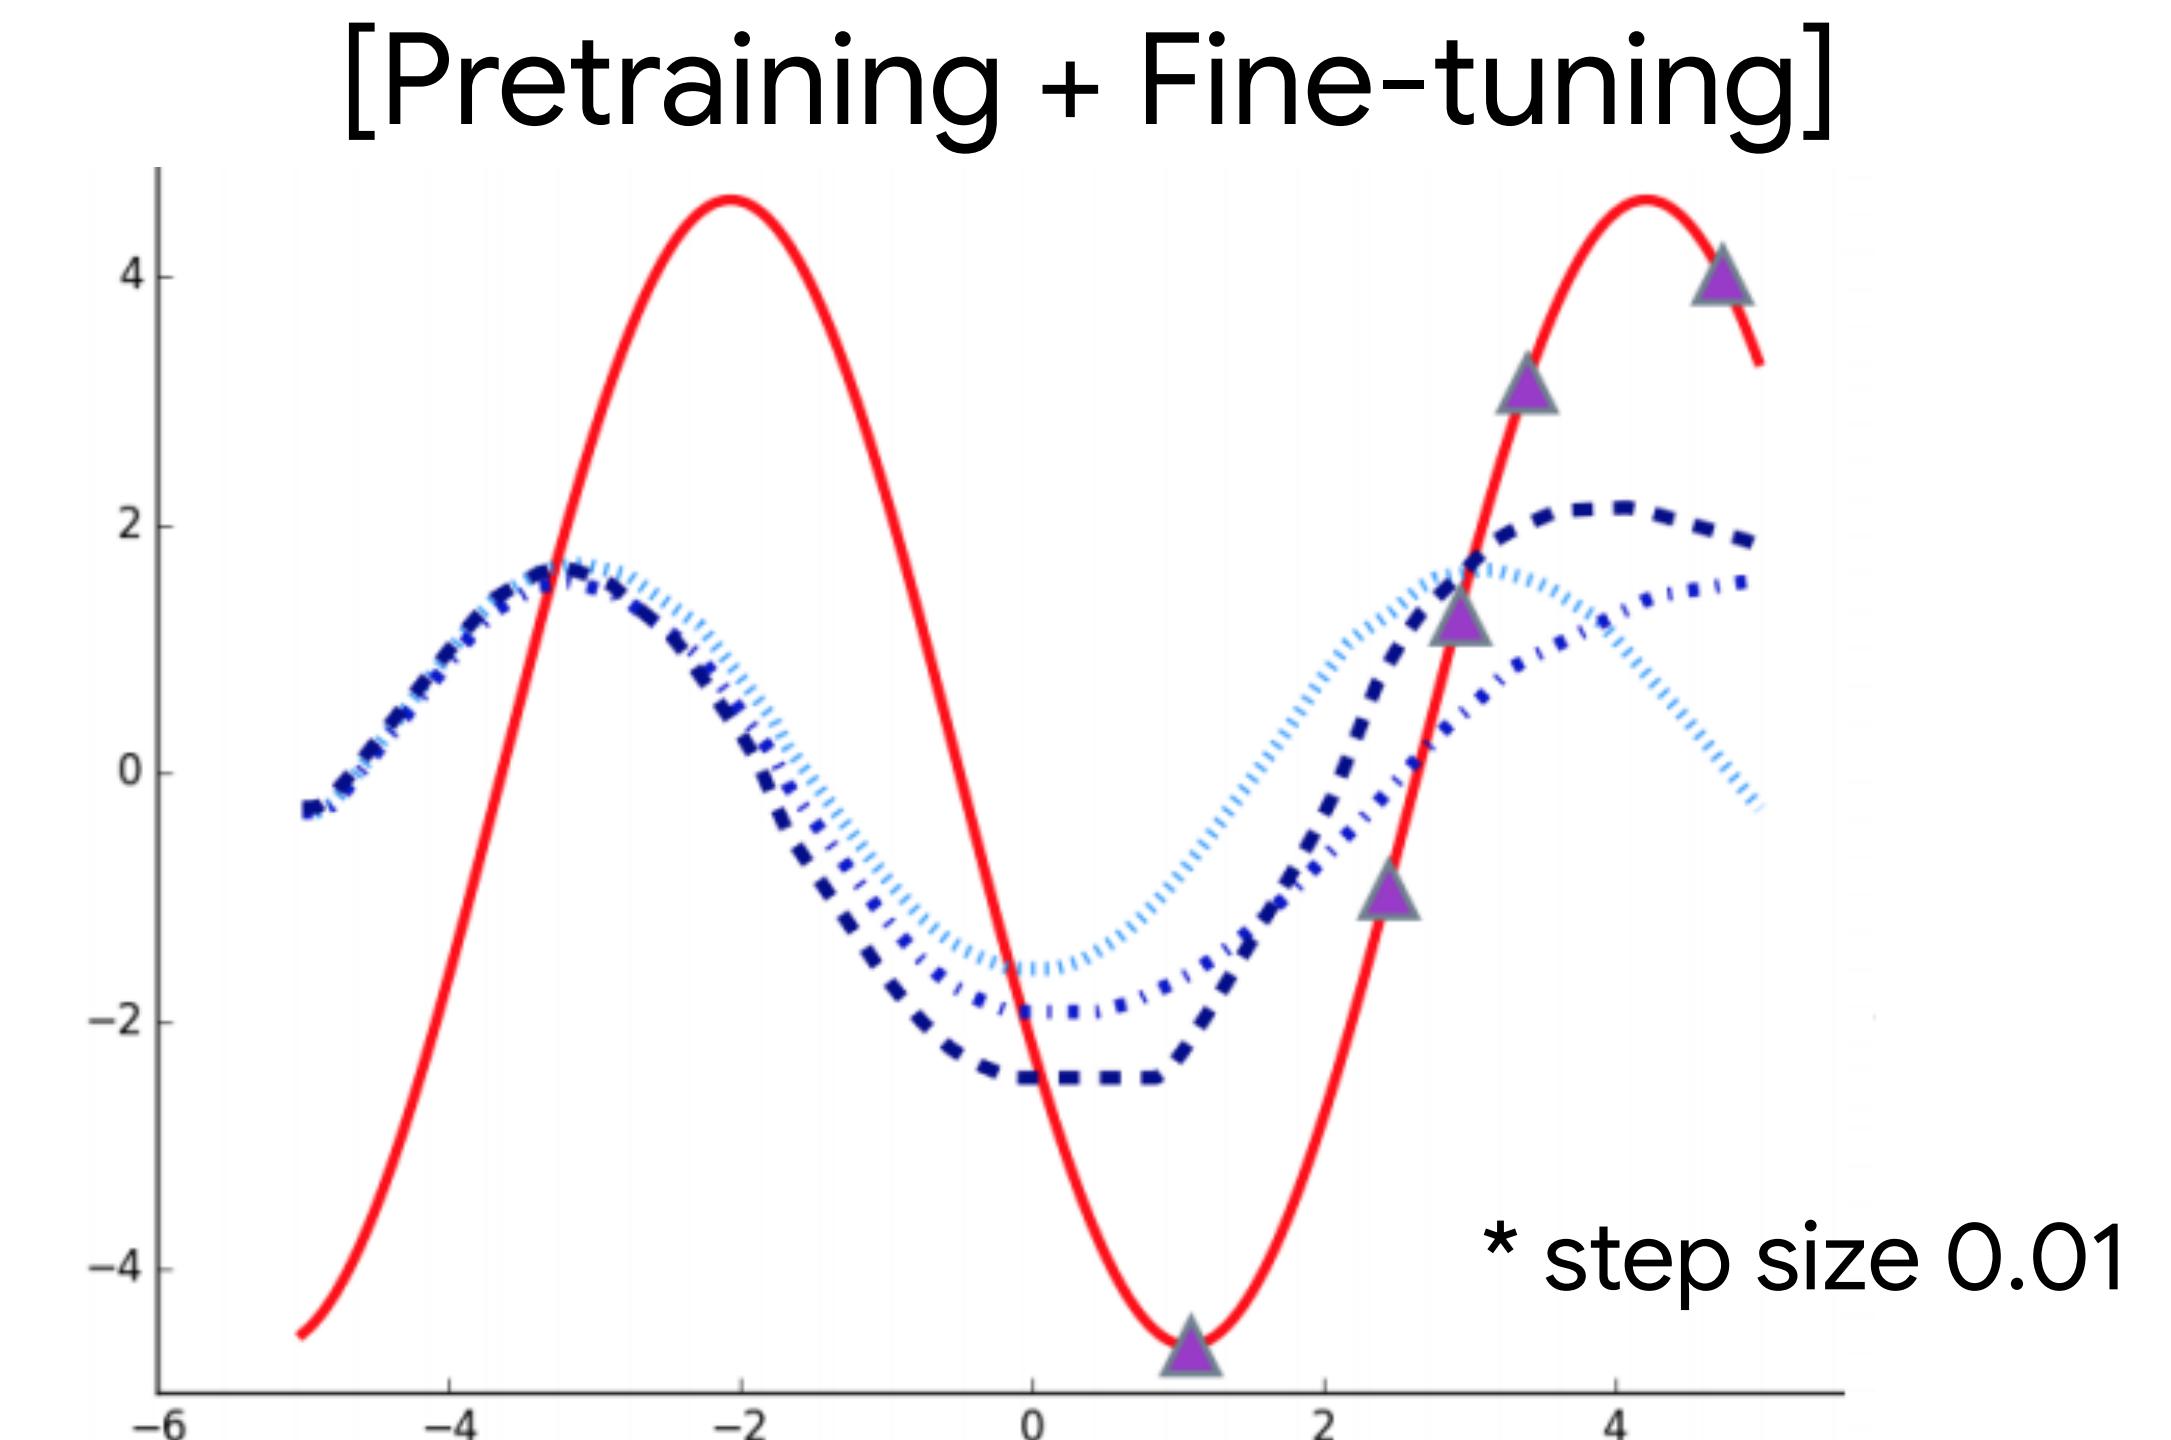
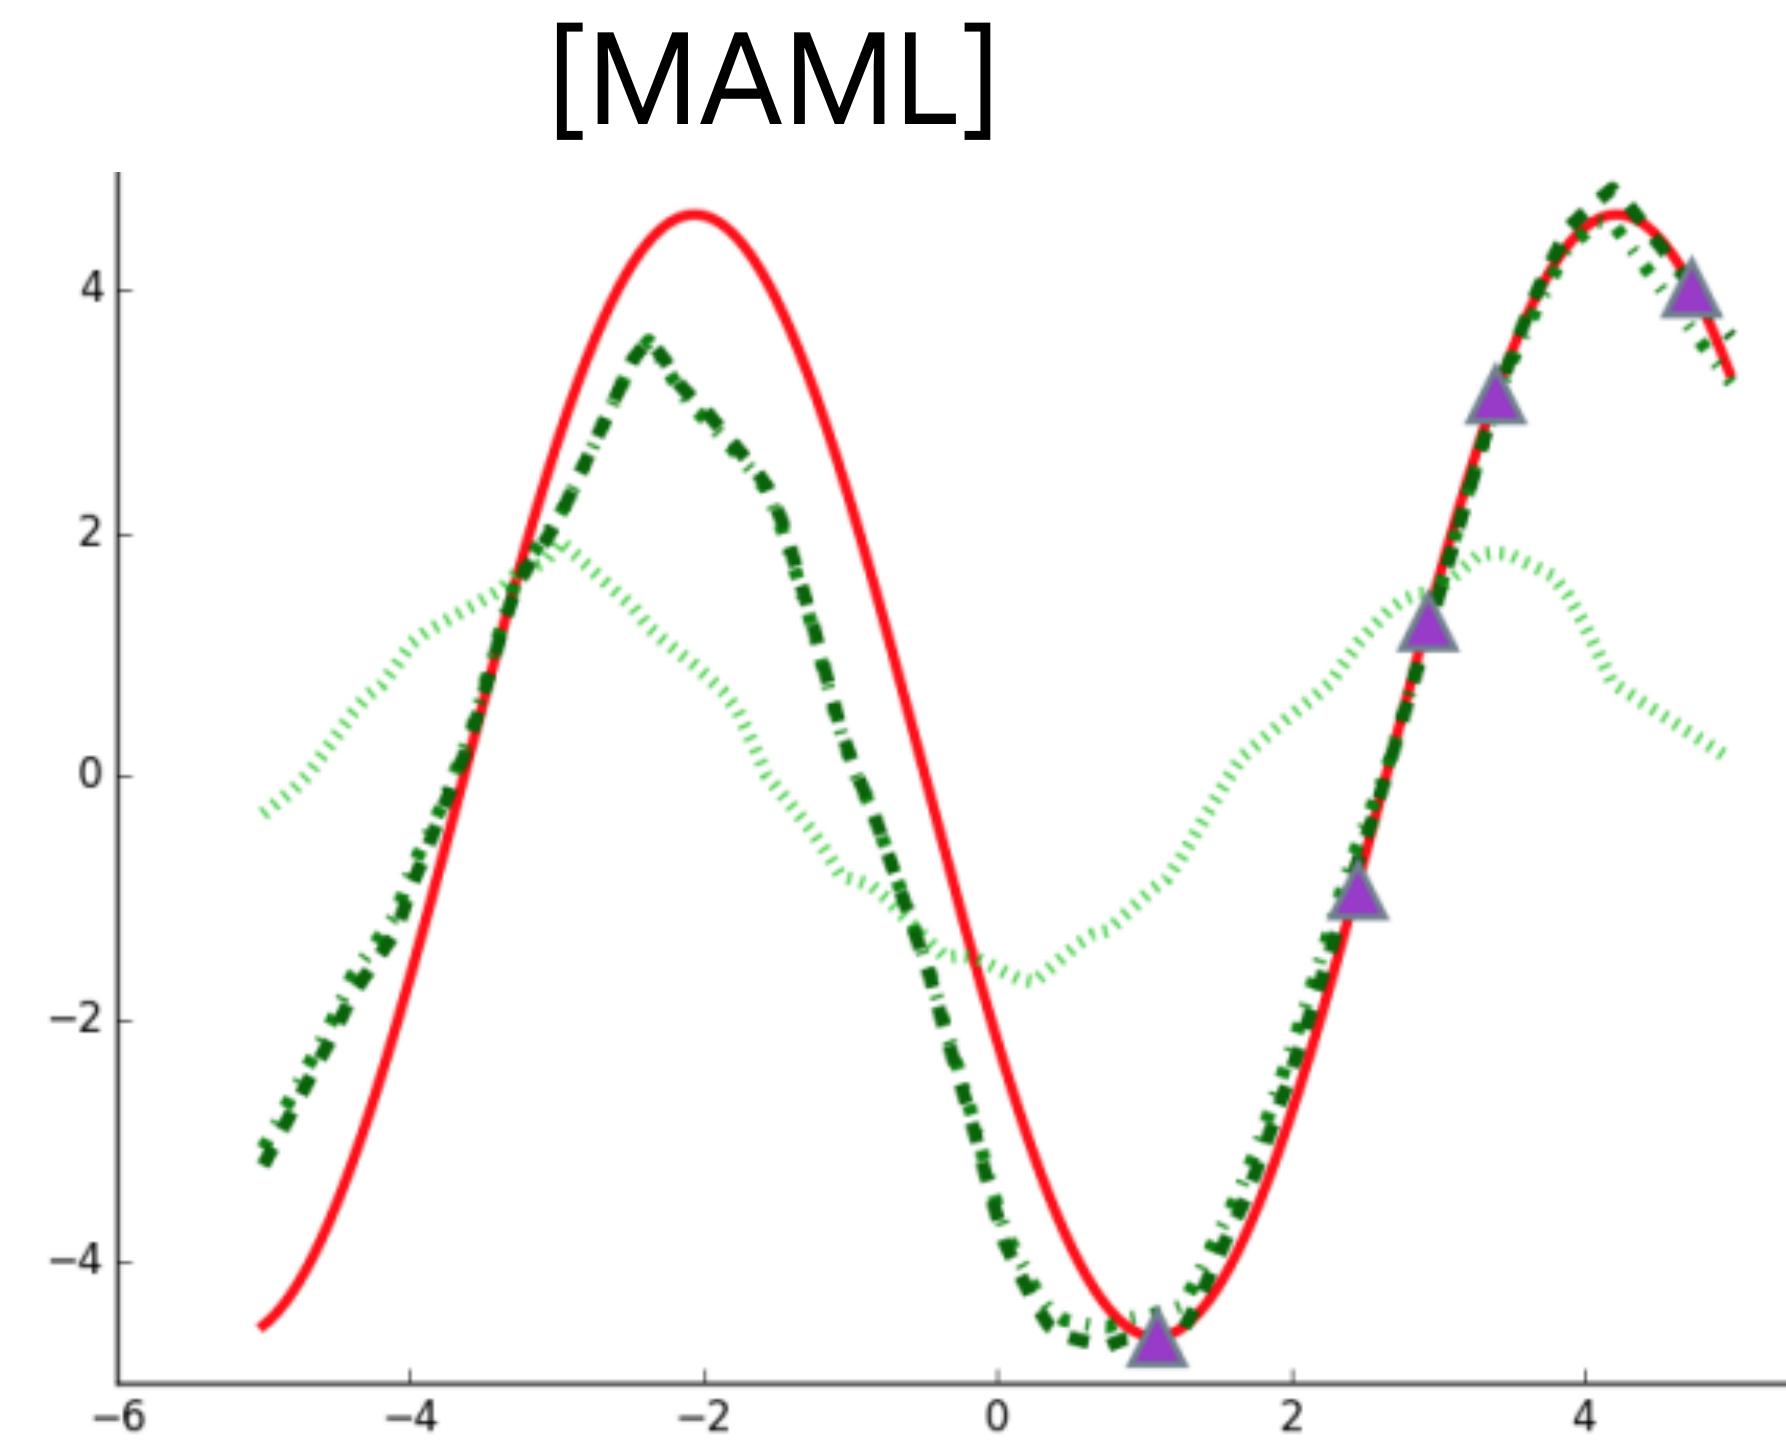
Results of 10-Shot Learning Regression



In the case of MAML, they quickly adapted to new samples.
But, adaptation failed in the case of pretraining model.

5. Experiment

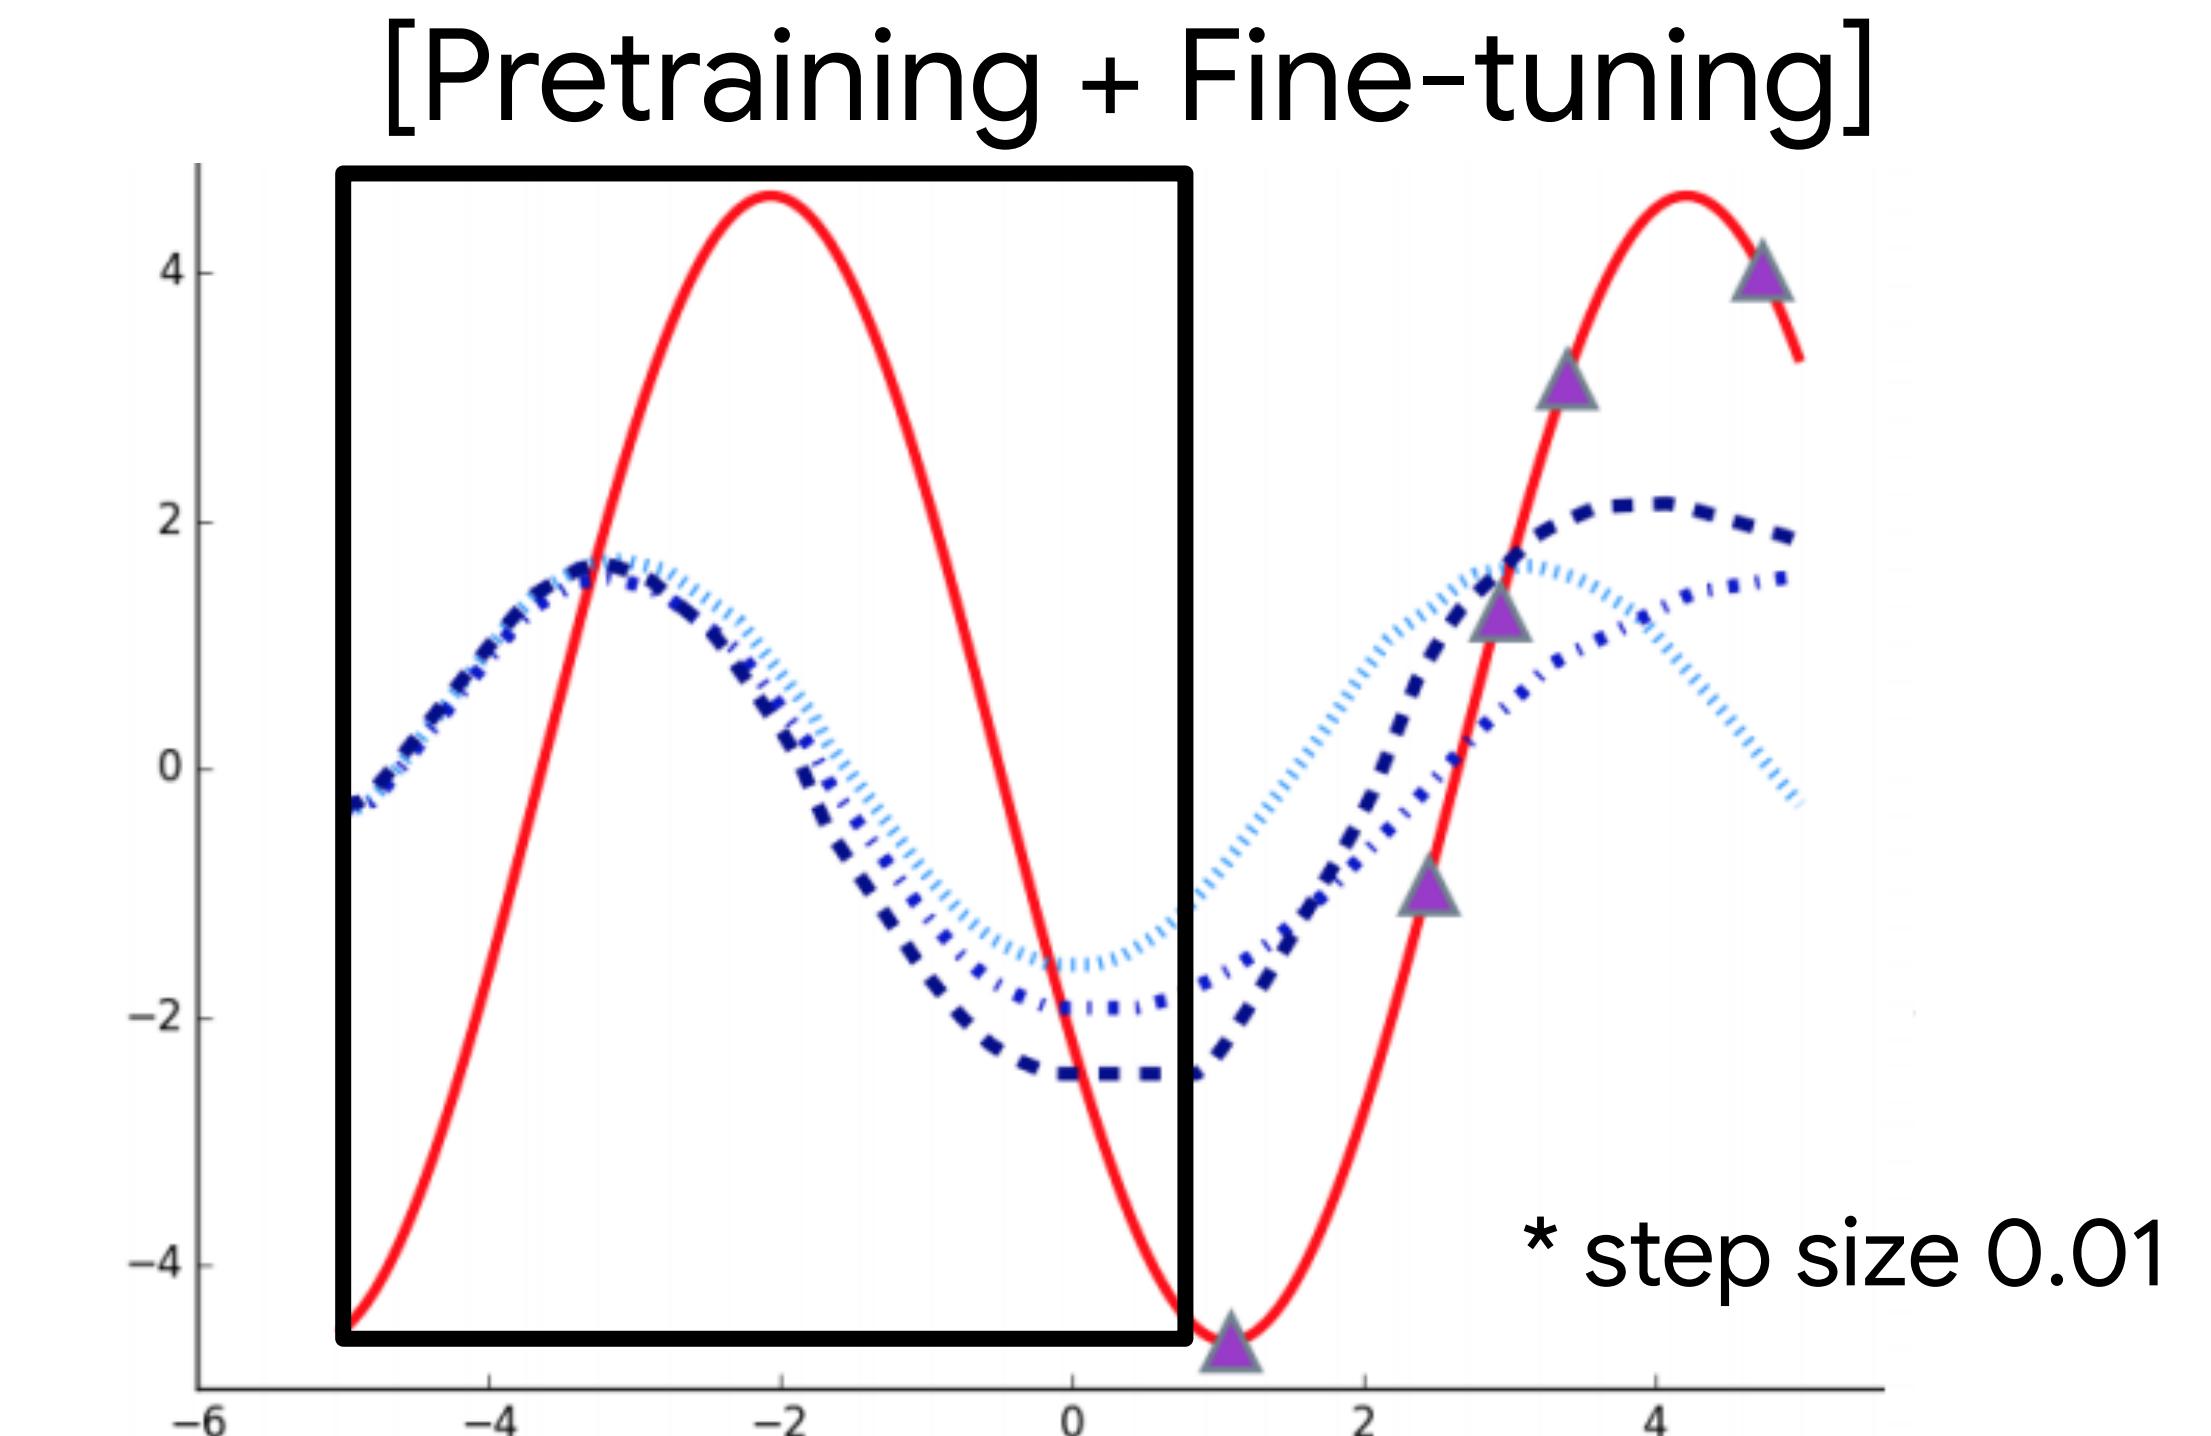
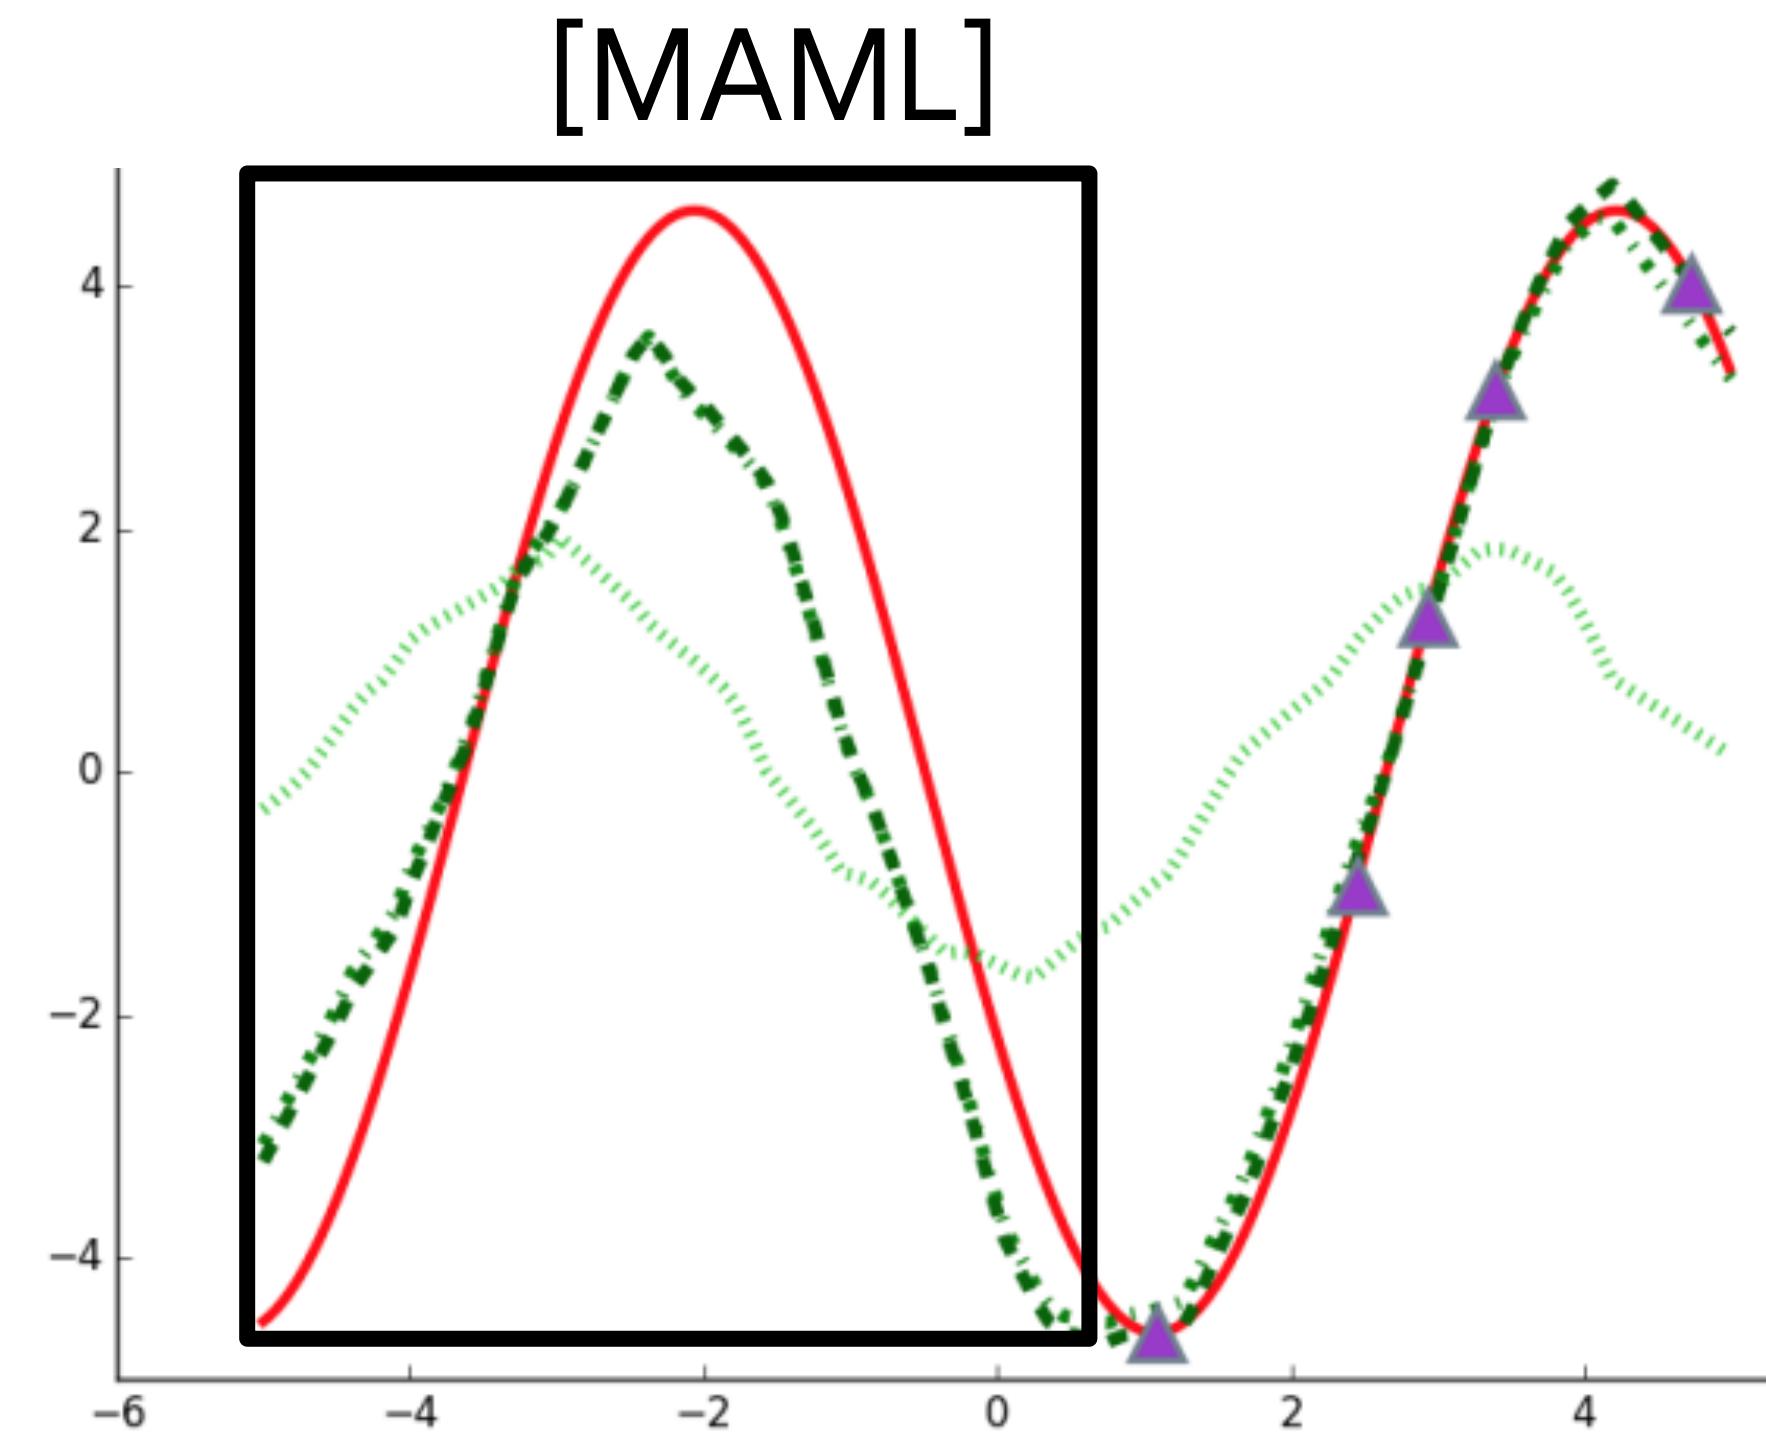
Results of 5-Shot Learning Regression



In the 5-shot learning, the difference is pervasive.

5. Experiment

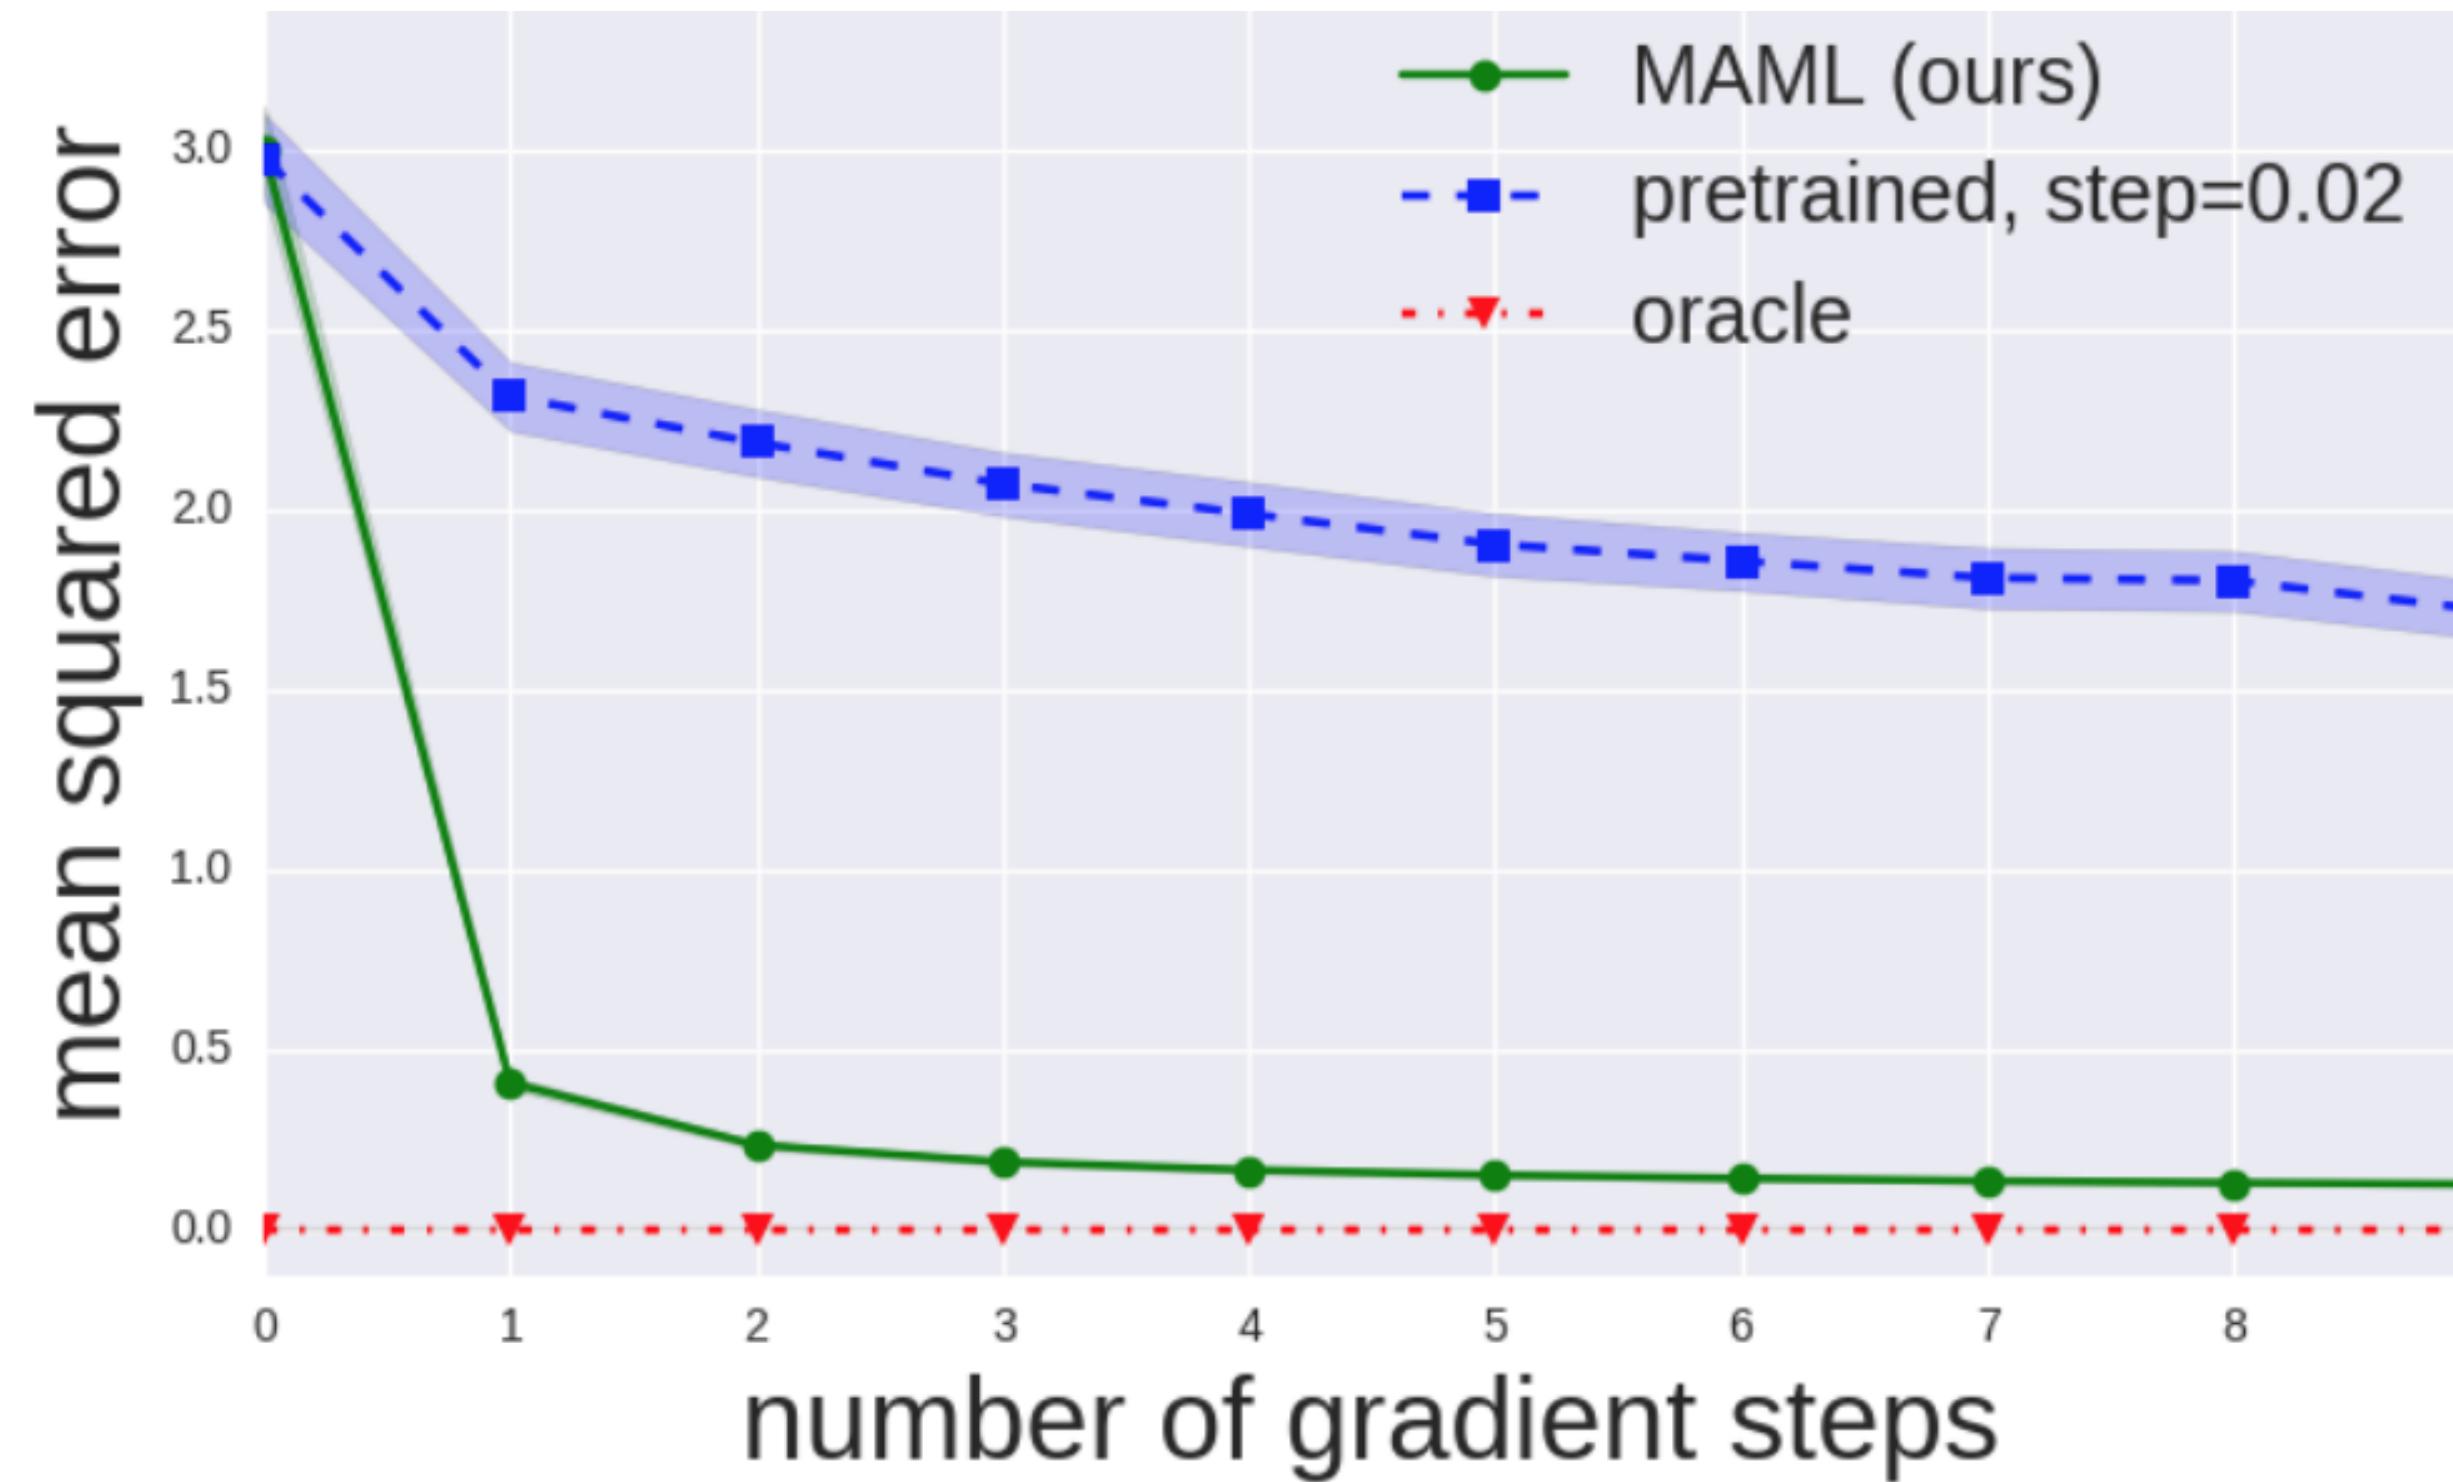
Results of 5-Shot Learning Regression



In the 5-shot learning, the difference is pervasive.
Good predictions are also made for ranges not particularly seen.

5. Experiment

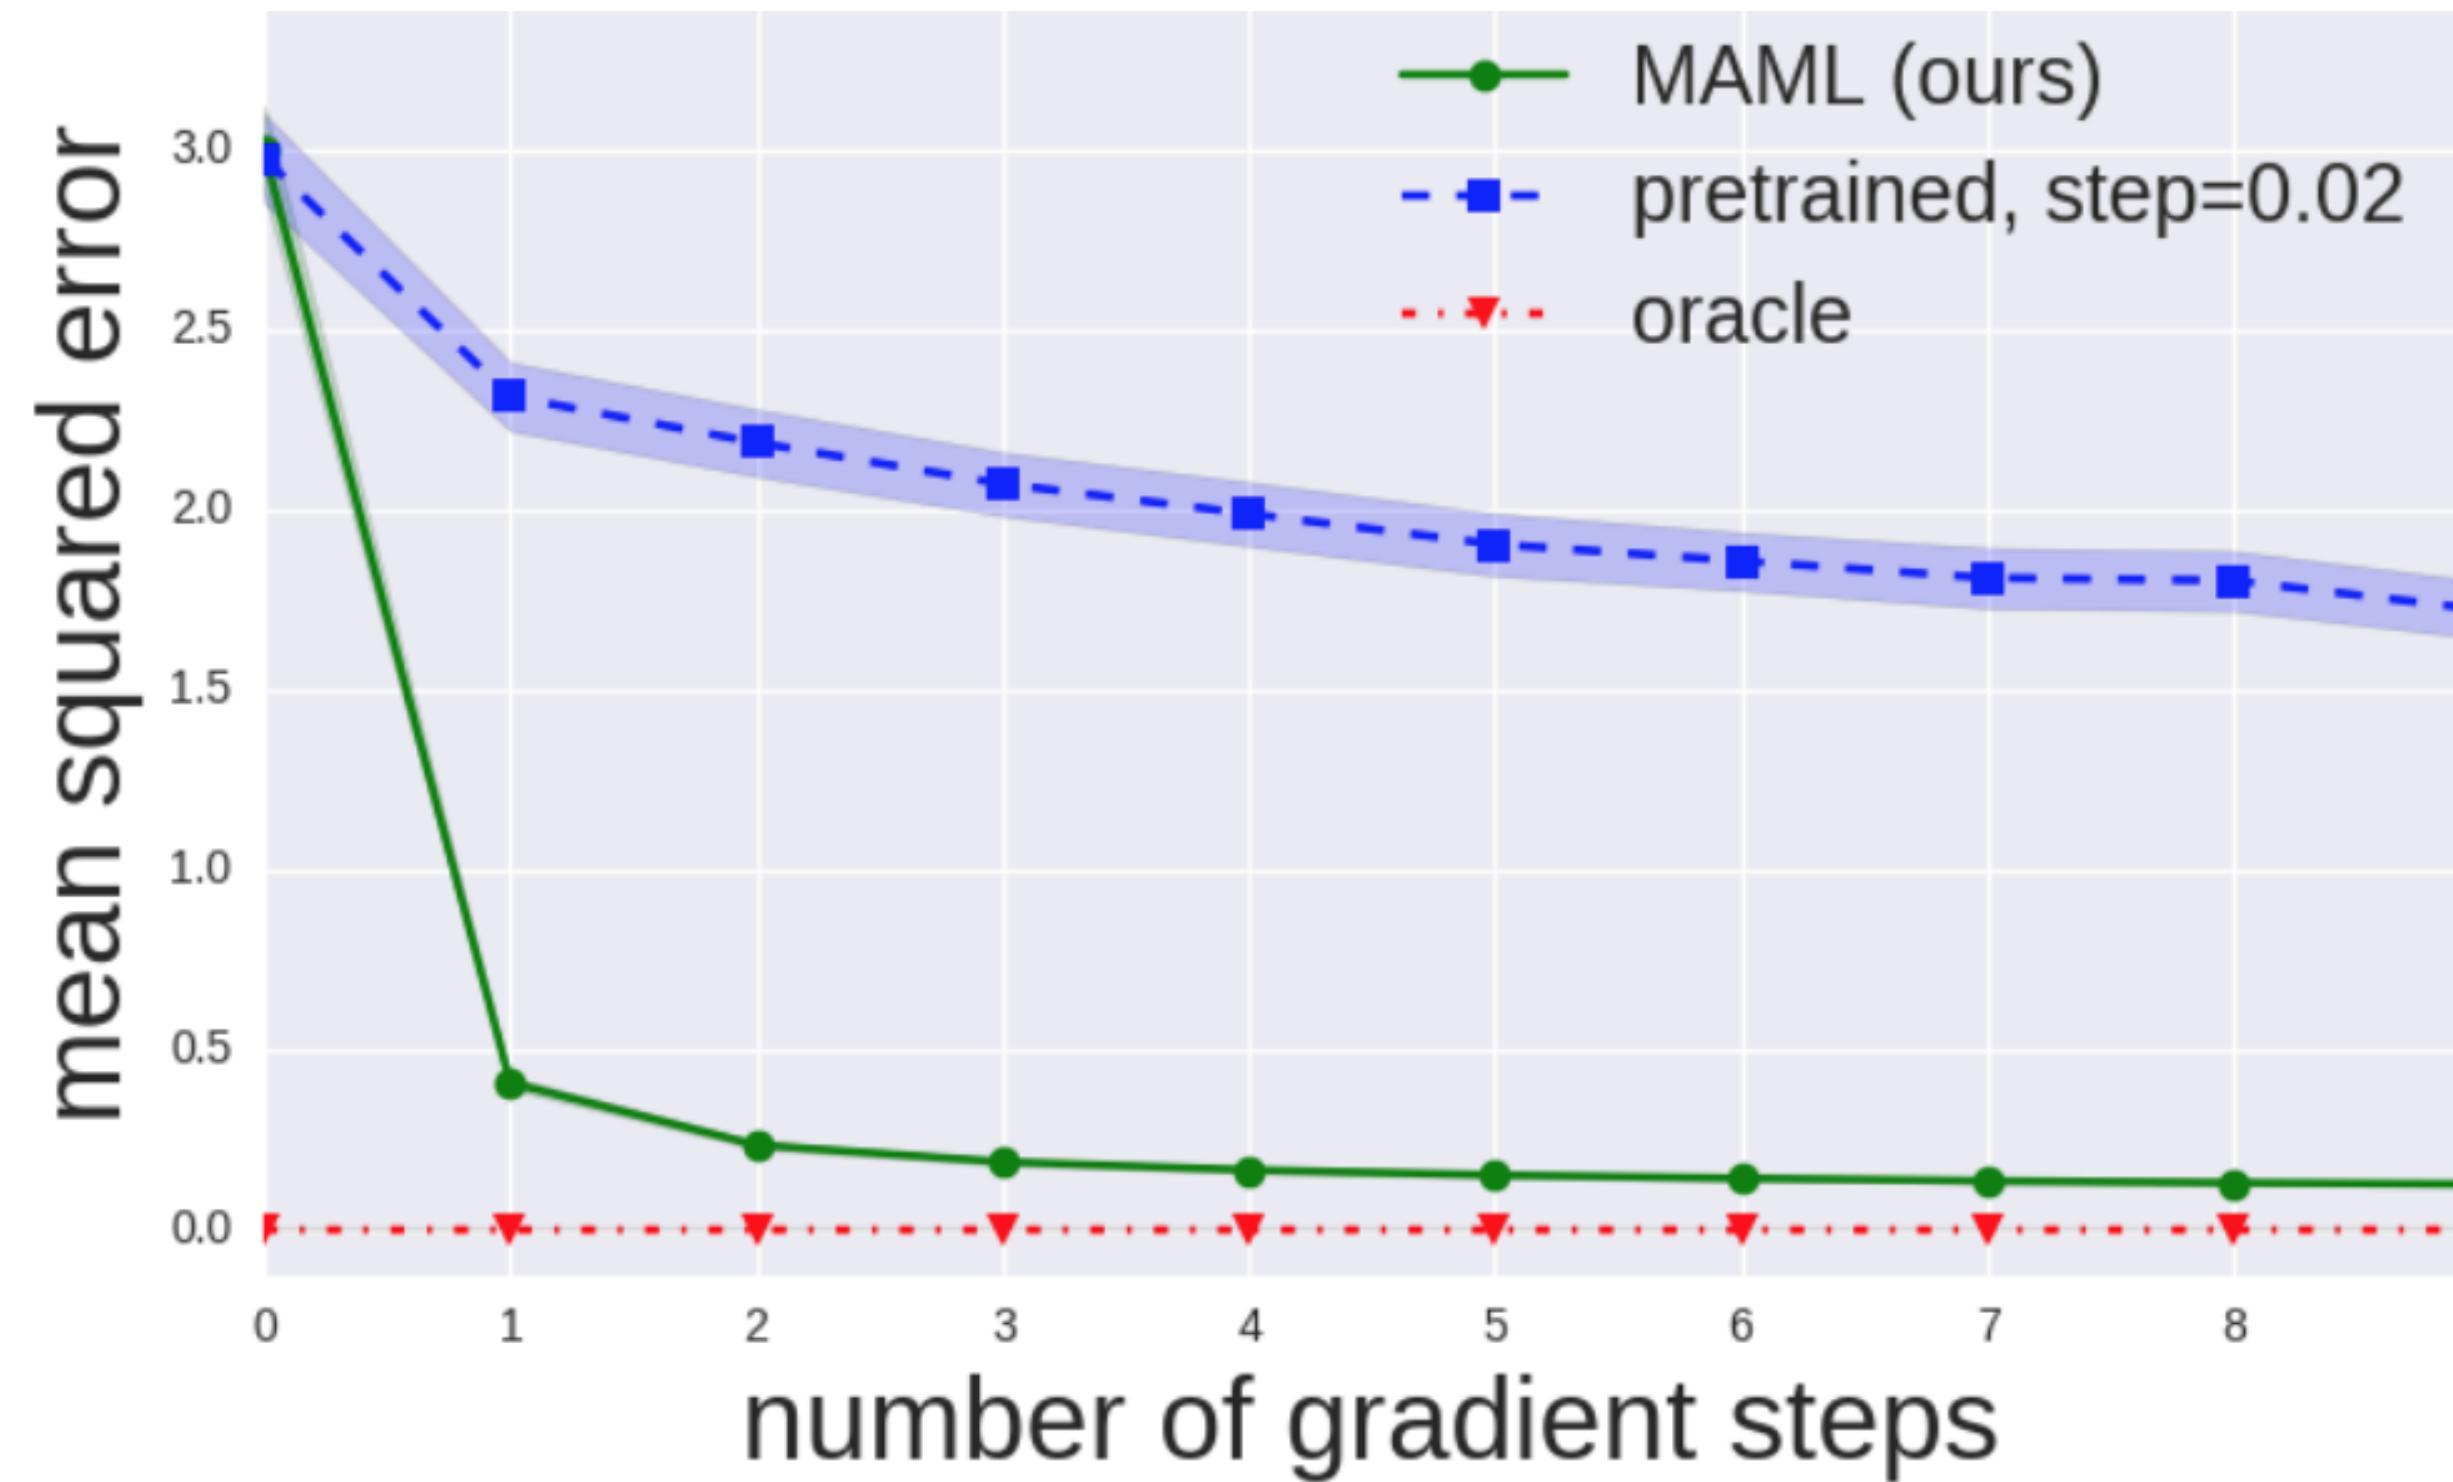
MAML Needs Only One Gradient Step



Vanilla pretrained model adapted slowly,
but, the MAML method quickly adapted **even in one gradient step**.

5. Experiment

MAML Needs Only One Gradient Step

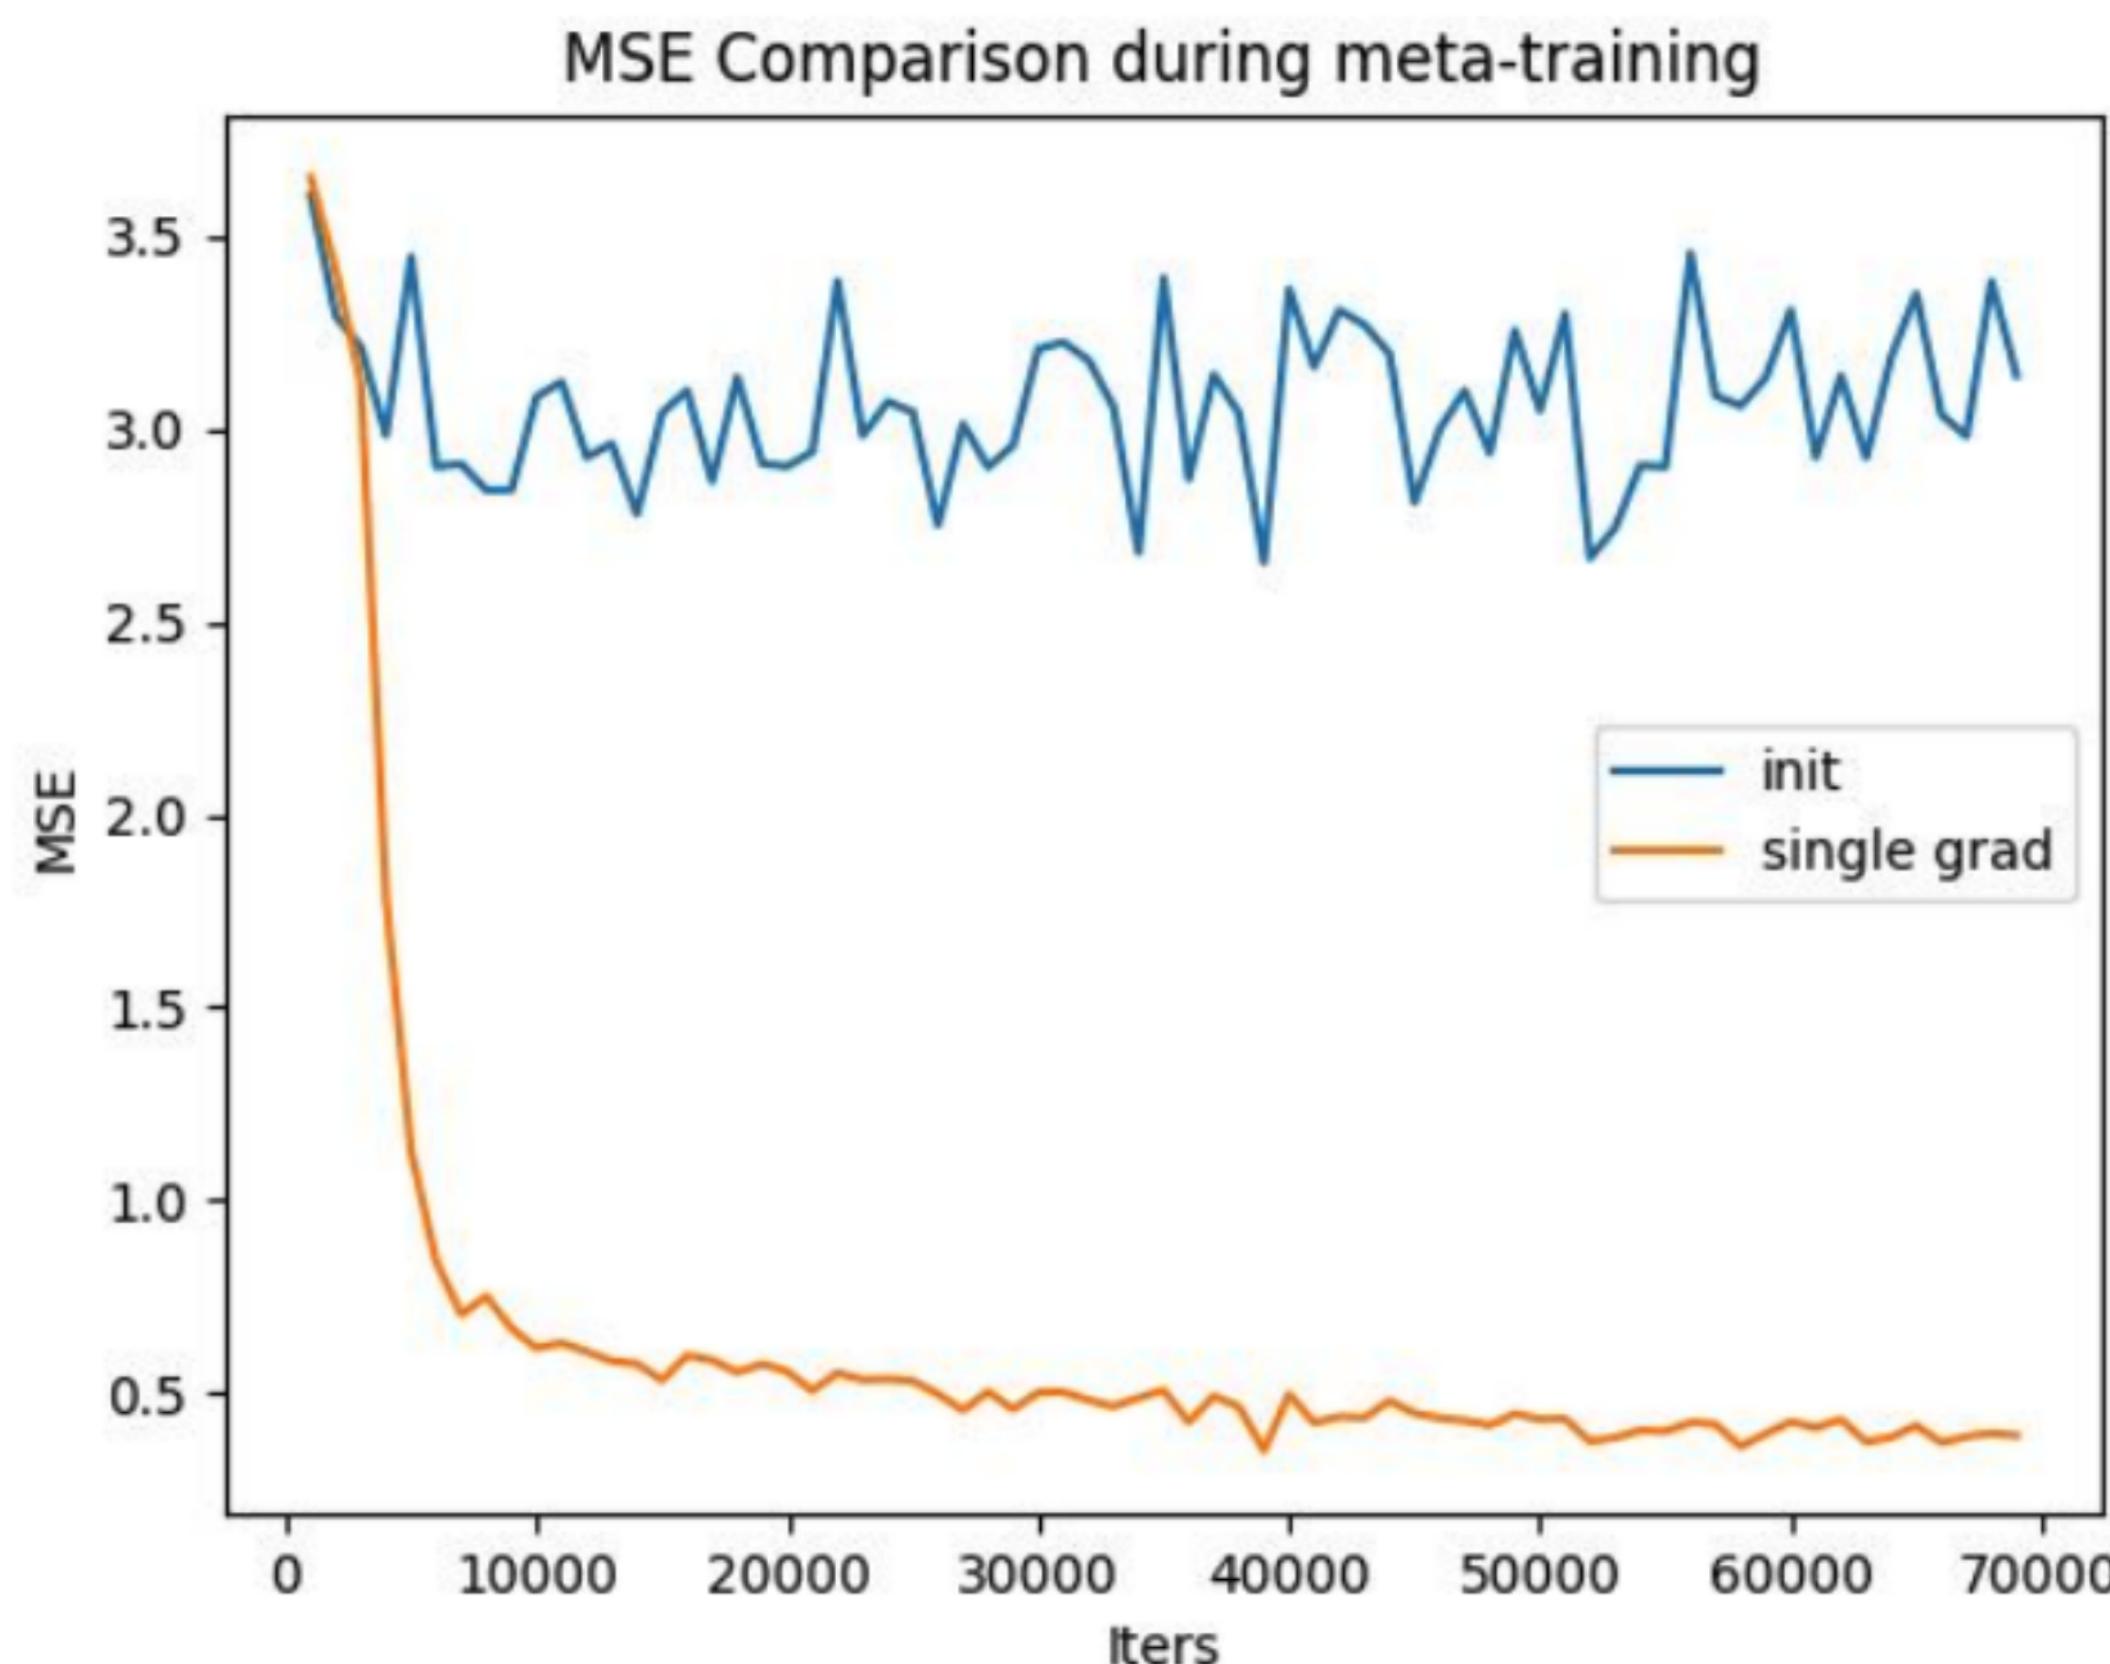


Vanilla pretrained model adapted slowly,
but, the MAML method quickly adapted **even in one gradient step**.

Meta parameter theta of MAML is **more sensitive**
than pre-updated theta of vanilla pretrained model.

5. Experiment

MSE of the Meta-Parameters



- The performance of the meta-parameters was not improved much in training.
- However, the performance of the single gradient updated parameters started on meta-parameters improved as training progressed.

5. Experiment

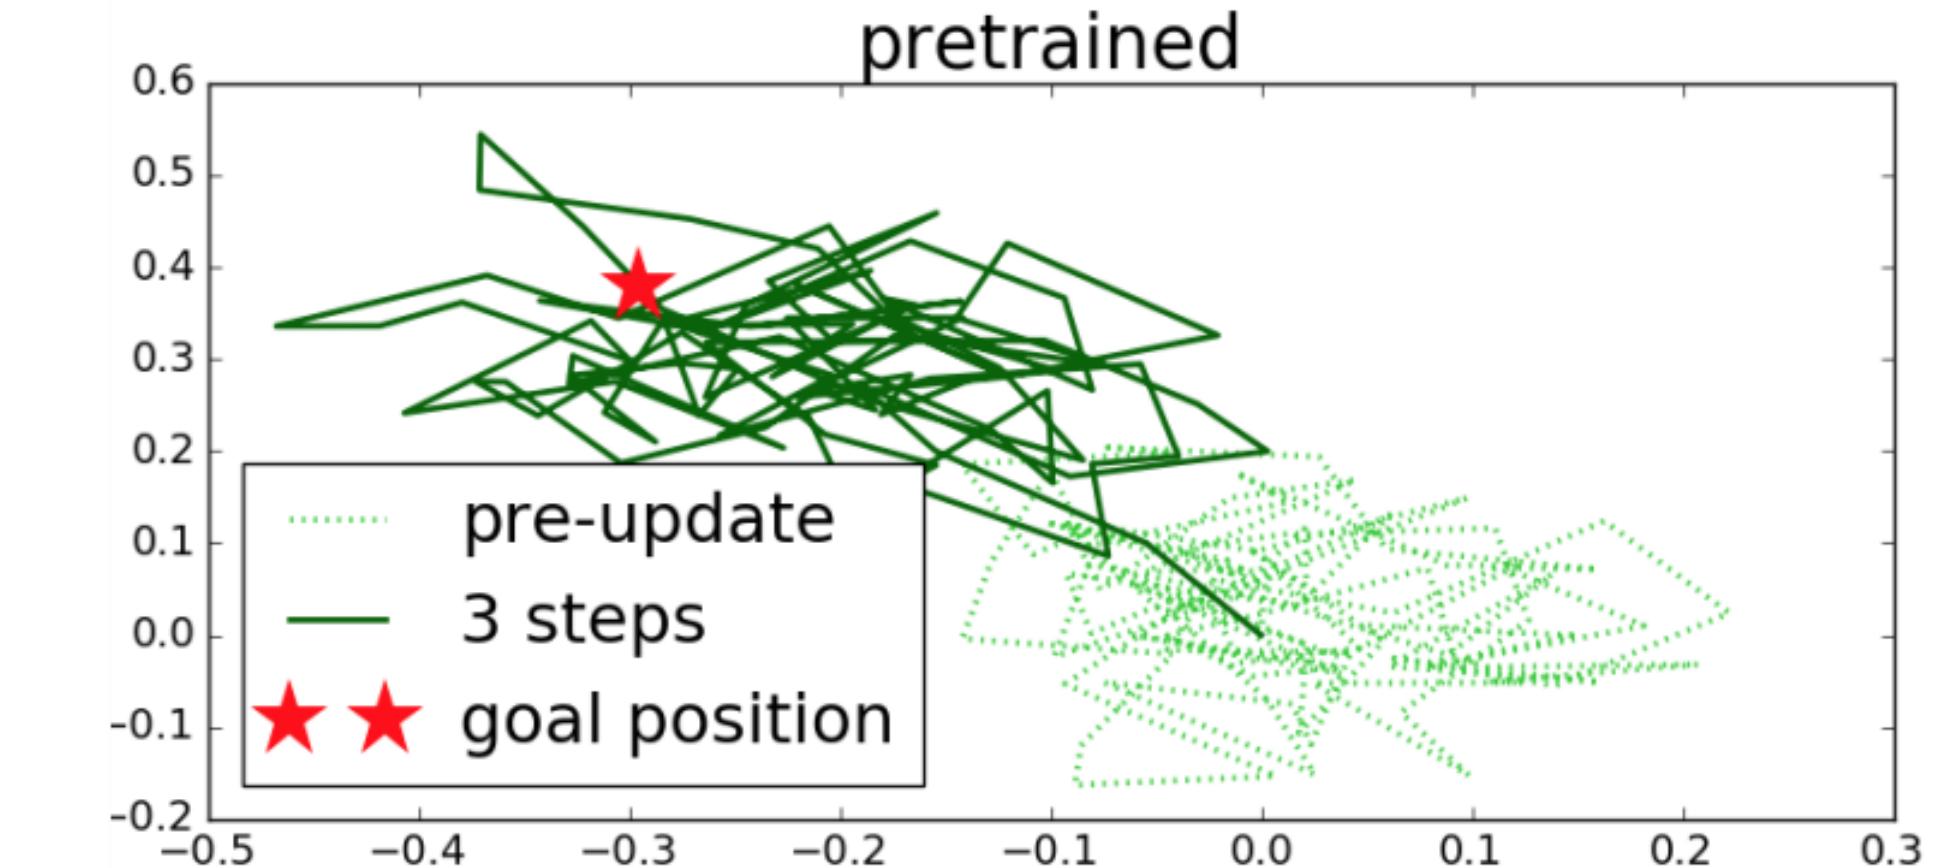
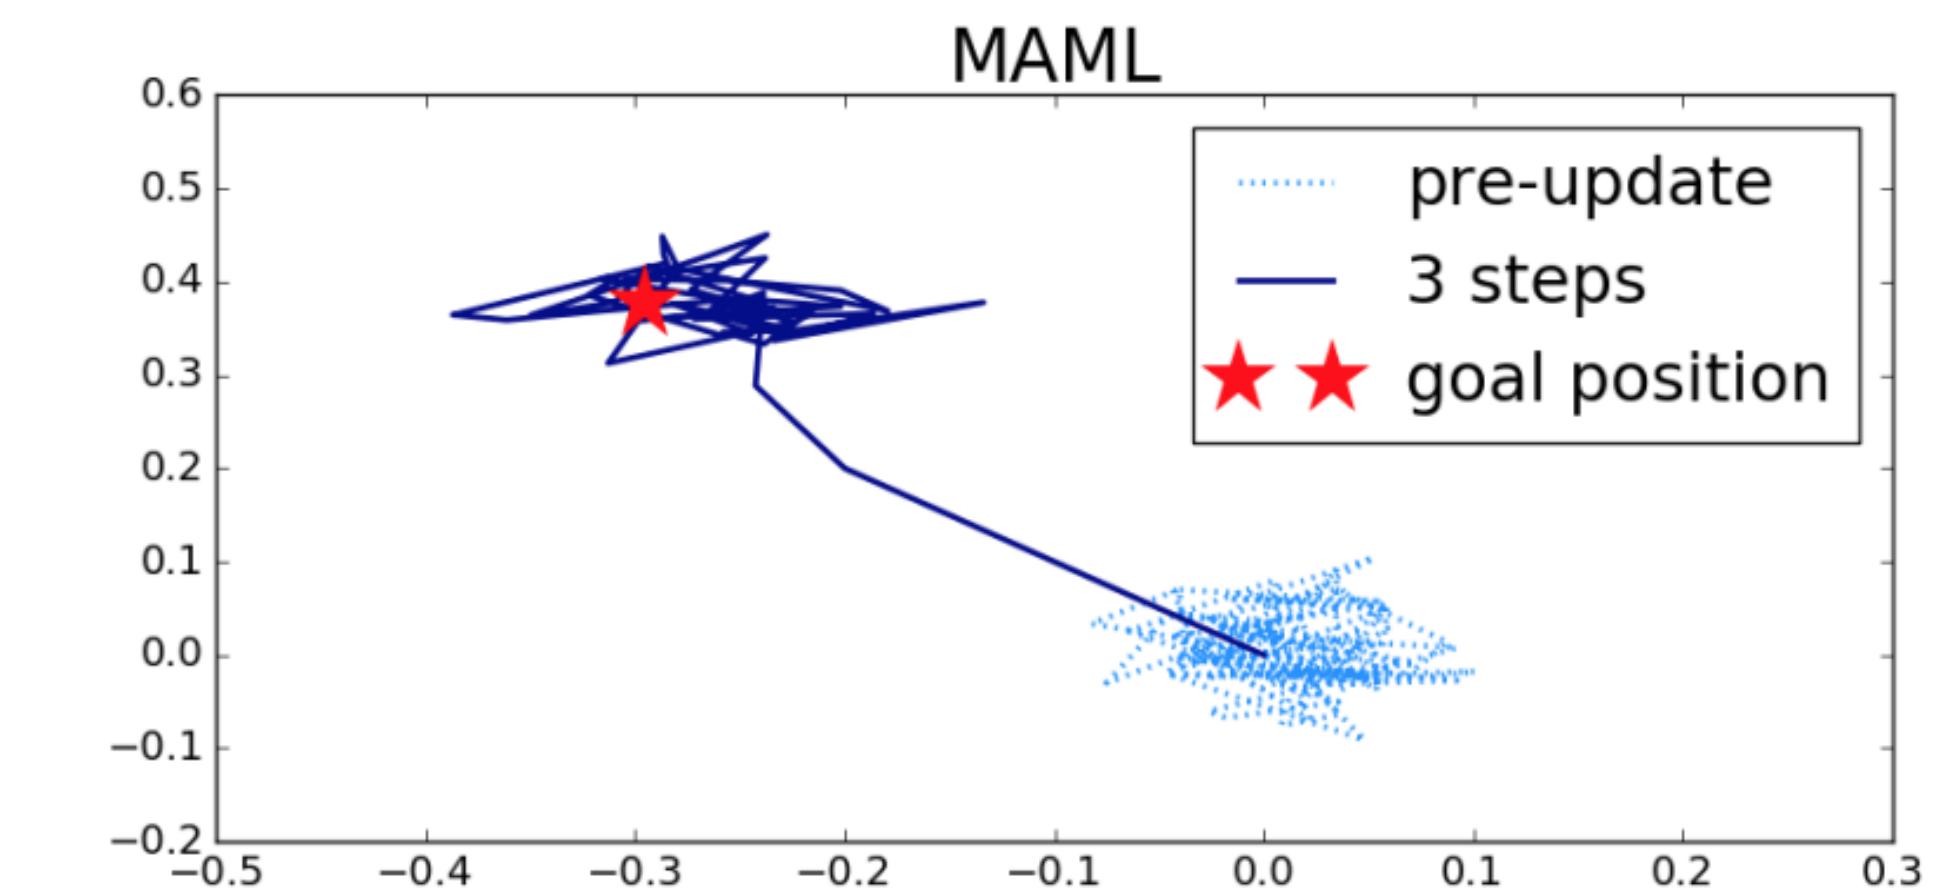
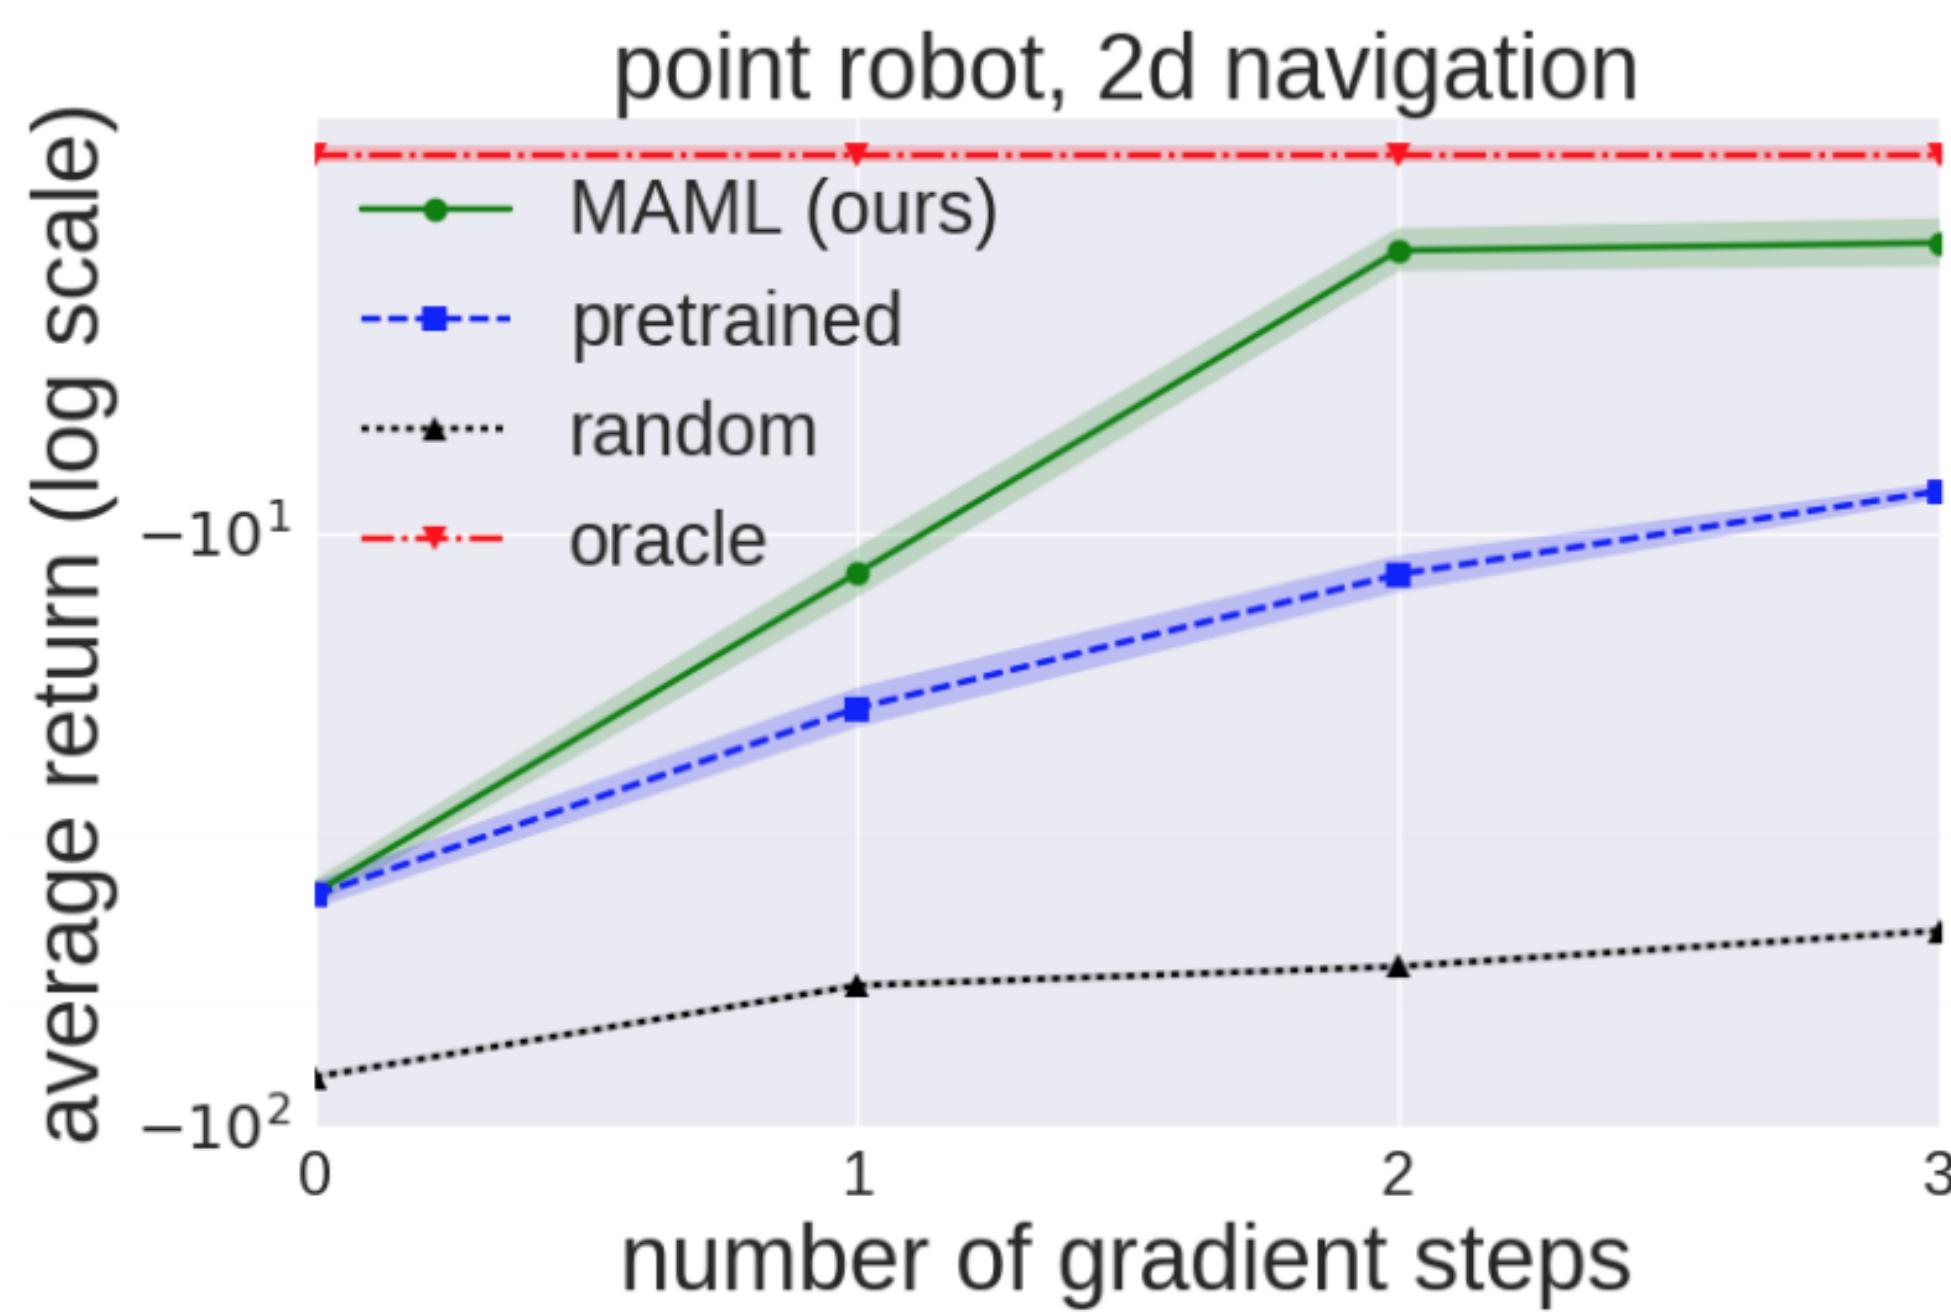
Experiments on Reinforcement Learning

- rllab benchmark suite
- Neural network policy with two hidden layers of size 100 with ReLU
- Gradients updates are computed using vanilla policy gradient (REINFORCE) and trust-region policy (TRPO) optimization as meta-optimizer.
- Comparison
 - Pretraining one policy on all of the tasks and fine-tuning
 - Training a policy from randomly initialized weights
 - Oracle policy

5. Experiment

Results on Reinforcement Learning

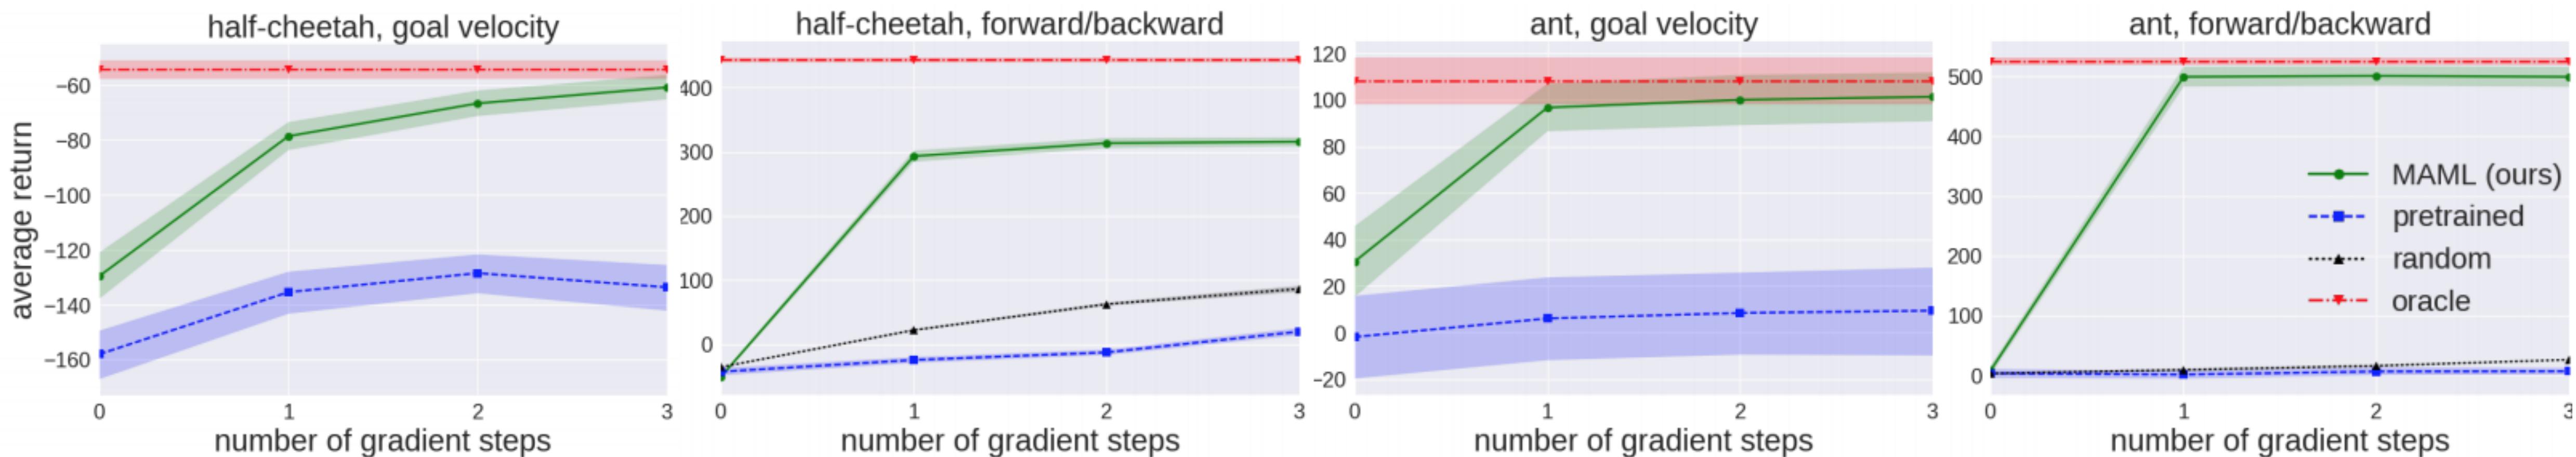
- 2d navigation



5. Experiment

Results on Reinforcement Learning

- Locomotion
 - High-dimensional locomotion tasks with the MuJoCo simulator



num. grad steps	0	1	2	3
context vector	-40.49	-44.08	-38.27	-42.50
MAML (ours)	-50.69	293.19	313.48	315.65

Thank you!

Reference

- Finn, Chelsea, Pieter Abbeel, and Sergey Levine. "Model-agnostic meta-learning for fast adaptation of deep networks." arXiv preprint arXiv:1703.03400 (2017).
- Montufar, Guido F., et al. "On the number of linear regions of deep neural networks." Advances in neural information processing systems. 2014.
- Baldi, Pierre. "Autoencoders, unsupervised learning, and deep architectures." Proceedings of ICML workshop on unsupervised and transfer learning. 2012.
- Li, Zhenguo, et al. "Meta-sgd: Learning to learn quickly for few shot learning." arXiv preprint arXiv:1707.09835 (2017)

Reference

- <https://www.slideshare.net/TaesuKim3/pr12094-modelagnostic-metalearning-for-fast-adaptation-of-deep-networks>
- https://people.eecs.berkeley.edu/~cbfinn/_files/nips2017_metaworkshop.pdf
- <https://github.com/cbfinn/maml>
- Lab seminar slide by Young-ki Hong.